

**LARGE LANGUAGE MODELS IN SOFTWARE
ENGINEERING:
TRANSFORMING CODE GENERATION, TESTING,
DEBUGGING,
AND DEVOPS WORKFLOWS**

An Independent Research Study

by

N Bhouvana Viswakarma

June 2026

DECLARATION

I, N Bhouvana Viswakarma, hereby declare that this research study is my own original independent work. The work presented here represents my independent synthesis and analysis of the literature cited throughout this study. This research has not previously been submitted for the award of any degree or diploma. All sources consulted have been duly acknowledged and cited in accordance with IEEE citation guidelines, and no part of this work has been plagiarized, fabricated, or misrepresented.

N Bhouvana Viswakarma

Date: June 2026

ACKNOWLEDGEMENTS

The completion of this thesis was made possible through the support and guidance of numerous individuals and institutions. I am grateful to the educators and mentors whose courses in algorithms, machine learning, software engineering, and systems design equipped me with the conceptual tools necessary to undertake research at the intersection of these disciplines. Access to open academic databases and publicly available research infrastructure provided the scholarly resources essential for a systematic literature review of this scope.

The research community whose work forms the foundation of this thesis deserves acknowledgment. Researchers across academia and industry have collectively advanced the field through open publication of findings, model weights, datasets, and benchmarks — a commitment to cumulative scientific progress that benefits the broader community and made this study possible.

Most profoundly, I express my deepest gratitude to my family for their unconditional love, patience, and support throughout this endeavour. Their confidence in my abilities provided the foundation without which this work could not have been completed. This thesis is dedicated to them.

N Bhouvana Viswakarma

June 2026, Bengaluru

ABSTRACT

Large language models (LLMs) have emerged as one of the more consequential technological developments of the early twenty-first century, with wide-ranging implications for the discipline of software engineering. Originating in the Transformer architecture of Vaswani et al. (2017) and accelerating with the introduction of GPT-3 (2020), Codex (2021), and successive generations of increasingly capable models, LLMs have shown a growing ability to generate, explain, refactor, test, debug, and document software across dozens of programming languages and paradigms. The deployment of GitHub Copilot in June 2022 marked the transition of LLM-assisted programming from academic curiosity to commercial reality, beginning a period of rapid tool proliferation and enterprise adoption that continues today.

The present study examines the role of large language models across core software engineering activities — code generation, automated testing, debugging, documentation generation, code review, infrastructure-as-code generation, and DevOps workflow automation — through a comparative, lifecycle-spanning analysis. The work is motivated by a persistent gap in the scholarly literature: while individual LLM applications in software engineering have received growing attention, a unified treatment that addresses technical capabilities, empirical evidence, limitations, ethical implications, and future research directions across the full lifecycle has remained absent.

The gap is addressed through a systematic literature review (SLR) augmented by secondary comparative empirical analysis. Following the protocol of Kitchenham and Charters (2007), the SLR was applied to a corpus of over 220 peer-reviewed publications, technical reports from major research laboratories, and empirical deployment studies spanning 2017 to 2025. The search encompassed five major databases — ACM Digital Library, IEEE Xplore, SpringerLink, arXiv, and Google Scholar — with an initial corpus of approximately 1,800 candidate papers reduced to 220 primary sources through systematic screening against predefined inclusion and exclusion criteria.

The comparative empirical analysis evaluates twelve leading models and tools including GPT-4o, Claude 3.5 Sonnet, Gemini 1.5 Pro, GitHub Copilot, Amazon Q Developer, Cursor, Windsurf, Claude Code, Devin, OpenHands, DeepSeek-Coder-V2, and Qwen2.5-Coder across ten performance and utility dimensions. Evaluation dimensions encompass functional correctness on HumanEval, SWE-bench Verified, LiveCodeBench, and MultiPL-E benchmarks; multi-step reasoning capability; context window utilization; hallucination rates; inference cost; response latency; and enterprise readiness.

Key empirical findings include the following. Contemporary frontier LLMs achieve pass@1 rates exceeding 90% on HumanEval — up from 28.8% for the original Codex in 2021. Agentic systems show growing but still limited capacity to resolve real-

world tasks, with the best achieving 37–49% resolution on SWE-bench Verified. GitHub Copilot has been shown to increase developer productivity by 35–55% on routine coding tasks, and LLM-assisted documentation generation reduces production time by 50–67% across documented enterprise deployments. Significant challenges remain: hallucination rates vary substantially across models and task types, with particular severity for proprietary API invocations; approximately 40% of LLM-generated code in security-sensitive contexts contains vulnerabilities; and unresolved questions of intellectual property, developer overreliance, environmental cost, and governance require continued scholarly and policy attention.

Five key contributions of this study are made to the scholarly literature: (1) a unified taxonomy of LLM applications across eight major software engineering lifecycle activities; (2) a comparative analysis of twelve leading models and agentic systems across ten evaluation dimensions, synthesizing benchmark results as of mid-2025; (3) six case studies of enterprise deployments, extracting lessons regarding architecture, workflow integration, productivity measurement, and business impact; (4) an analysis of ten challenge and risk categories — hallucination, security, intellectual property, overreliance, environmental impact, bias, governance, prompt injection, context window limitations, and benchmark contamination; and (5) a forward-looking research agenda identifying thirteen priority directions with estimated time horizons, key challenges, and current readiness assessments.

The evidence suggests that large language models have moved from experimental productivity supplements to nascent infrastructure within software engineering practice, altering how code is written, reviewed, tested, documented, and deployed. Responsible integration into professional workflows, however, requires the research community, industry, policymakers, and educators to jointly address unresolved challenges in reliability, security, ethics, environmental sustainability, and governance.

Keywords: Large Language Models, Software Engineering, Code Generation, Automated Testing, Automated Debugging, DevOps Automation, GitHub Copilot, Transformer Architecture, HumanEval, SWE-bench, Agentic AI, Retrieval-Augmented Generation, RLHF, Constitutional AI, Code Intelligence, AI-Assisted Development, Hallucination, Software Security, Benchmark Evaluation, Enterprise Adoption

TABLE OF CONTENTS

Declaration

Acknowledgements

Abstract

Table of Contents

List of Tables

List of Figures

List of Abbreviations

CHAPTER 1: INTRODUCTION

1.1 Background and Historical Context

1.2 The Emergence of AI in Software Engineering

1.3 The LLM Revolution: A Chronological Survey

1.4 Research Motivation and Problem Statement

1.5 Research Objectives and Questions

1.6 Significance of the Study

1.7 Scope, Delimitations, and Limitations

1.8 Research Methodology Overview

1.9 Principal Contributions

1.10 Organization of the Thesis

CHAPTER 2: LITERATURE REVIEW

2.1 Artificial Intelligence: Foundations

2.2 Machine Learning Paradigms

2.3 Deep Learning Architectures

2.4 The Transformer Architecture

2.5 BERT and Bidirectional Pre-training

2.6 CodeBERT and Graph-Enhanced Models

2.7 GPT Series and Autoregressive Generation

2.8 T5 and Encoder-Decoder Approaches

2.9 Codex: Foundation of AI Coding Assistants

2.10 StarCoder and Open-Source Code Models

2.11 Code Llama and Meta's Open Models

2.12 DeepSeek-Coder: Efficiency at Scale

2.13 Qwen2.5-Coder: Multilingual Support

2.14 Claude: Safety-Aligned Models

2.15 Gemini: Natively Multimodal Models

2.16 GitHub Copilot in the Developer Workflow

2.17 Cursor and AI-Native IDEs

- 2.18 Windsurf, Claude Code, and Agentic Coding
- 2.19 SWE-agent, Devin, and OpenHands
- 2.20 Retrieval-Augmented Generation for Code
- 2.21 RLHF and Alignment Techniques
- 2.22 MCP, Tool Use, and Agentic Frameworks
- 2.23 Evaluation Benchmarks
- 2.24 Hallucination Research and Mitigation
- 2.25 Cross-Lingual Code Support
- 2.26 Economic Impact and Labor Market Effects
- 2.27 Synthesis and Research Gap

CHAPTER 3: FOUNDATIONS OF LARGE LANGUAGE MODELS

- 3.1 Tokenization and Vocabulary
- 3.2 Embeddings and Positional Encoding
- 3.3 Self-Attention and Multi-Head Attention
- 3.4 Feed-Forward Networks and Normalization
- 3.5 Decoder-Only vs Encoder-Decoder Architectures
- 3.6 Mixture of Experts
- 3.7 Training Pipeline and Scaling Laws
- 3.8 Inference Optimization Techniques
- 3.9 Alignment: RLHF, CAI, and DPO
- 3.10 In-Context Learning and Prompting Theory
- 3.11 Long-Context Architectures
- 3.12 Emergent Abilities in Large Models

CHAPTER 4: LLM APPLICATIONS IN SOFTWARE ENGINEERING

- 4.1 Taxonomy of LLM Applications
- 4.2 Code Generation
- 4.3 Code Explanation and Summarization
- 4.4 Code Refactoring
- 4.5 Bug Detection and Static Analysis
- 4.6 Debugging and Fault Localization
- 4.7 Code Review Automation
- 4.8 Test Generation
- 4.9 Documentation Generation
- 4.10 DevOps and Infrastructure Automation
- 4.11 SQL and Database Code Generation
- 4.12 API Generation and Specification
- 4.13 Code Migration and Modernization
- 4.14 Accessibility and Inclusive Development

4.15 Natural Language Requirements to Code

CHAPTER 5: COMPARATIVE STUDY OF LLM-BASED SE TOOLS

5.1 Evaluation Framework and Methodology

5.2 GPT-4o Analysis

5.3 Claude 3.5 Sonnet Analysis

5.4 Gemini 1.5 Pro Analysis

5.5 GitHub Copilot Analysis

5.6 Amazon Q Developer Analysis

5.7 Cursor and Windsurf Analysis

5.8 Claude Code and Devin Analysis

5.9 OpenHands Platform Analysis

5.10 Comprehensive Comparative Table

5.11 Prompt Engineering Strategies

5.12 Cost-Performance Analysis

5.13 Open-Source vs Proprietary Trade-offs

CHAPTER 6: CASE STUDIES

6.1 Introduction and Methodology

6.2 Microsoft GitHub Copilot

6.3 Google Gemini Code Assist

6.4 Amazon Q Developer

6.5 Anthropic Claude Code

6.6 Sourcegraph Cody

6.7 OpenHands Agentic Platform

6.8 Meta's Internal AI Coding Deployment

6.9 Emerging Patterns in Enterprise Adoption

CHAPTER 7: CHALLENGES AND RISKS

7.1 Hallucinations and Factual Unreliability

7.2 Security Vulnerabilities in Generated Code

7.3 Intellectual Property and Copyright

7.4 Developer Overreliance and Skill Atrophy

7.5 Environmental and Computational Cost

7.6 Bias, Fairness, and Representation

7.7 Governance, Accountability, and Compliance

7.8 Prompt Injection and Adversarial Attacks

7.9 Context Window Limitations

7.10 Evaluation Validity and Benchmark Contamination

CHAPTER 8: FUTURE RESEARCH DIRECTIONS

8.1 Agentic Software Engineers

- 8.2 Self-Improving Software Systems
- 8.3 Autonomous Debugging and Testing
- 8.4 Autonomous DevOps
- 8.5 Neuro-Symbolic Integration
- 8.6 Smaller Local Models and Edge AI
- 8.7 Human-AI Collaboration Models
- 8.8 Multimodal Software Engineering
- 8.9 Evaluation Methodology Advances
- 8.10 Enterprise Governance and Policy
- 8.11 Continuous Code Quality Management
- 8.12 Cross-Repository and Ecosystem Analysis
- 8.13 Personalized AI Development Assistance

CHAPTER 9: CONCLUSION

- 9.1 Summary of Research Findings
- 9.2 Answers to Research Questions
- 9.3 Theoretical and Practical Implications
- 9.4 Limitations of this Study
- 9.5 Recommendations for Practitioners
- 9.6 Closing Remarks

REFERENCES

APPENDIX A: Benchmark Descriptions and Datasets

APPENDIX B: Prompt Engineering Templates

APPENDIX C: Glossary of Terms

APPENDIX D: Research Methodology Details

LIST OF TABLES

Table 2.1: Comprehensive Model and Tool Comparison

Table 3.1: Summary of Positional Encoding Methods

Table 3.2: Scaling Laws: Key Studies and Findings

Table 3.3: Alignment Techniques Comparison

Table 4.1: Code Generation Benchmark Performance (mid-2025)

Table 4.2: LLM Capabilities in DevOps Domains

Table 4.3: Test Generation: LLM vs Traditional Tools

Table 4.4: Prompt Engineering Strategy Comparison

Table 5.1: Comprehensive Tool Comparison: 10 Metrics

Table 5.2: Cost Structure Comparison for Enterprise Deployment

Table 5.3: Open-Source vs Proprietary Model Trade-offs

Table 6.1: Case Study Overview: Deployment Characteristics

Table 6.2: GitHub Copilot: Productivity Measurement Studies

Table 6.3: Enterprise Adoption Patterns Summary

Table 7.1: Security Vulnerability Categories in LLM-Generated Code

Table 7.2: Environmental Cost Estimates for LLM Training

Table 7.3: Challenge Categories: Severity and Mitigation Status

Table 8.1: Research Agenda: Directions, Horizons, and Priorities

Table 9.1: Research Questions and Summary Answers

Table A.1: Benchmark Dataset Characteristics

LIST OF ABBREVIATIONS

Abbreviation	Expansion
AGI	Artificial General Intelligence
AI	Artificial Intelligence
ACI	Agent-Computer Interface
API	Application Programming Interface
APR	Automated Program Repair
AST	Abstract Syntax Tree
BPE	Byte-Pair Encoding
CAI	Constitutional AI
CI/CD	Continuous Integration / Continuous Delivery
CLI	Command Line Interface
CNN	Convolutional Neural Network
CoT	Chain of Thought
DPO	Direct Preference Optimization
E2E	End-to-End
FFN	Feed-Forward Network
FIM	Fill-in-the-Middle
FP16	16-bit Floating Point
GQA	Grouped Query Attention
GPU	Graphics Processing Unit
HumanEval	Human Evaluation Benchmark (OpenAI)
IaC	Infrastructure as Code
IDE	Integrated Development Environment
LLM	Large Language Model
LoRA	Low-Rank Adaptation
MBPP	Mostly Basic Python Problems Benchmark
MCP	Model Context Protocol
MLM	Masked Language Modeling
MoE	Mixture of Experts
MTTR	Mean Time to Resolve
NLP	Natural Language Processing
NSP	Next Sentence Prediction
OWASP	Open Web Application Security Project
PPO	Proximal Policy Optimization
PR	Pull Request
RAG	Retrieval-Augmented Generation
ReLU	Rectified Linear Unit
RLHF	Reinforcement Learning from Human Feedback
RNN	Recurrent Neural Network
RoPE	Rotary Position Embedding
SE	Software Engineering
SLR	Systematic Literature Review
SQL	Structured Query Language
SAST	Static Application Security Testing
SWE-bench	Software Engineering Benchmark
TPU	Tensor Processing Unit
T2S	Text-to-SQL
YAML	YAML Ain't Markup Language

CHAPTER 1: INTRODUCTION

1.1 Background and Historical Context

Software engineering, as a formal discipline, emerged from the crucible of what Dijkstra and Naur famously characterized as the 'software crisis' at the NATO Software Engineering Conferences of 1968 and 1969 [1]. The precipitating observation was stark: the complexity of software systems being demanded by government, military, and commercial clients had grown to a scale that far outpaced the capacity of individual programmers and informal development teams to deliver them reliably, on schedule, and within budget. The Apollo Guidance Computer, developed at MIT under the direction of Margaret Hamilton, offered a prescient illustration of both the heights that disciplined software engineering could reach and the extraordinary cognitive demands that large-scale software placed on its developers [2]. The response to the software crisis was the emergence of software engineering as a systematic, disciplined, and quantifiable approach to software development — an engineering discipline in the fullest sense of the term.

Over the half-century following those seminal NATO conferences, software engineering evolved through successive paradigms of practice and tooling, each addressing limitations exposed by the preceding generation. Structured programming, championed by Dijkstra, Hoare, and Wirth through the 1970s, provided the theoretical and practical foundations for replacing undisciplined control flow with hierarchically organized program structures built from sequence, selection, and iteration [3]. The concept of modular design, elaborated by Parnas through his canonical papers on information hiding and decomposition [181], established the principle that managing software complexity required disciplined partitioning of system concerns into modules with narrow, well-defined interfaces. Object-oriented programming, emergent from Simula and Smalltalk and popularized through C++ and Java in the 1980s and 1990s, provided abstraction mechanisms enabling larger teams to manage greater complexity by modelling domain concepts as software objects [4]. The 1990s witnessed the rise of agile methodologies — eXtreme Programming, Scrum, and Lean development — which replaced heavyweight waterfall processes with iterative, customer-centric development cycles capable of accommodating the reality that software requirements evolve [5].

Despite this extraordinary disciplinary evolution, software engineering remained — through the first decade of the twenty-first century — fundamentally a human cognitive activity. The translation of requirements into executable specifications, the diagnosis of elusive defects in complex system interactions, the design of scalable and maintainable architectures, and the evaluation of competing implementation strategies all required the kind of specialized expertise cultivated over years of intensive practice.

Automation tools existed — compilers, static analysis tools, automated test runners, build systems — but these tools automated well-defined, mechanically specifiable sub-tasks. The higher-level cognitive activities at the core of software engineering remained stubbornly resistant to automation. The prospect of automating these activities at a level approaching human proficiency was, until recently, largely the province of science fiction and long-horizon speculation [6].

The emergence of machine learning as a practical engineering discipline in the early 2010s, catalysed by the availability of large datasets, GPU computing, and the AlexNet breakthrough in 2012, began to shift this landscape. Neural approaches to code completion, code search, and bug detection demonstrated that statistical learning from large code corpora could capture useful patterns in software artifacts [7]. Program synthesis research explored formal methods for generating programs from logical specifications, revealing both the theoretical possibility and practical difficulty of automated code generation [8]. Recurrent neural networks, particularly LSTM architectures, were applied to code modeling tasks and demonstrated ability to learn syntactic patterns and short-range semantic relationships in code sequences [9].

The watershed transformation arrived with the introduction of the Transformer architecture by Vaswani et al. in 2017 [10]. By replacing sequential recurrent processing with parallel self-attention computation — enabling each element of a sequence to relate directly to every other element regardless of distance — the Transformer architecture unlocked the ability to train models at scales and on datasets previously impractical. The immediate applications were in natural language processing, where pre-trained Transformer models achieved rapid state-of-the-art improvements across benchmarks. The translation of this architectural revolution to code-related tasks, through models including CodeBERT, GPT-3, and ultimately Codex, established the template for the current generation of LLM-based software engineering tools [10, 11, 12, 13, 14, 15, 16, 17, 18].

1.2 The Emergence of Artificial Intelligence in Software Engineering

The application of artificial intelligence to software engineering tasks has a history that predates the current LLM era by several decades, beginning with program synthesis research in the 1970s and progressing through successive generations of increasingly capable techniques. Understanding this historical trajectory illuminates both the novelty of current LLM-based approaches and the extent to which they build on and supersede prior work.

Early program synthesis research pursued a vision of automated programming grounded in formal logic: given a specification expressed in a formal language, construct a program guaranteed to satisfy that specification by systematic search through the space of possible programs [182]. This approach demonstrated theoretical feasibility for

restricted domains but proved computationally intractable for the complexity levels required by real-world software systems. The fundamental challenge — that the space of possible programs grows exponentially with program size, making exhaustive search infeasible — was addressed through inductive synthesis, constraint-based synthesis, and example-based synthesis, each of which narrowed the search space through additional assumptions but retained significant practical limitations [183, 184].

The resurgence of neural network research in the 2000s and the deep learning revolution of the early 2010s opened qualitatively different avenues for statistical approaches to code understanding and generation. Rather than searching through a formally defined program space, neural approaches learned distributed representations of code semantics from large corpora of existing software, capturing statistical regularities that could inform code completion, bug detection, and code search. Bug2Fix [185], CCT (Code Clone Tool), and related work demonstrated that sequence-to-sequence neural models could learn to fix simple bugs when trained on large datasets of before-and-after code pairs. DeepFix [186] applied recurrent neural networks to fix common compilation errors in student-written C programs, demonstrating practical utility for a restricted but important class of software defects.

The natural language processing community's parallel development of pre-trained language models provided the architectural and methodological foundation that would prove transformative for code. The success of ELMo [187], followed rapidly by GPT [188] and BERT [2], established the pre-training and fine-tuning paradigm as dramatically superior to task-specific model training for most NLP tasks. The insight — that models pre-trained on vast quantities of unlabeled text developed rich representations transferable to diverse downstream tasks — translated directly to the code domain. IntelliCode, introduced by Microsoft Research in 2018 [97], provided the first production-scale demonstration that pre-trained language models could meaningfully improve developer productivity through contextually ranked code completions. Though IntelliCode's underlying models were small by today's standards, its commercial deployment validated the fundamental premise that neural language model-based code assistance was practically useful.

The pivotal commercial inflection point arrived with the demonstration of Codex by Chen et al. at OpenAI in June 2021 [18]. Derived from GPT-3 through fine-tuning on approximately 159 gigabytes of Python code from public GitHub repositories, Codex demonstrated that a language model at the scale of GPT-3 — 12 billion parameters in the version underlying early Copilot deployments — could generate functionally correct solutions to a significant and practically meaningful fraction of programming challenge problems when given natural language specifications. The HumanEval benchmark introduced alongside Codex provided the first standardized evaluation for code generation, enabling systematic comparison of subsequent models. Codex's immediate

deployment as the foundation for GitHub Copilot, which entered technical preview in June 2021 and general availability in June 2022, marked the transition of LLM-based coding assistance from academic research to commercial reality at scale.

1.3 The LLM Revolution: A Chronological Survey, 2017–2025

The period from 2017 to 2025 witnessed a transformation in AI capability for software engineering widely regarded as one of the most significant in the history of the discipline. Understanding this transformation in chronological terms provides essential context for evaluating the state of the field and assessing the credibility of claims about future progress.

The Transformer architecture, introduced in the seminal 'Attention Is All You Need' paper by Vaswani et al. in 2017 [10], provided the architectural foundation for all subsequent LLM developments. Though the original paper's focus was machine translation, the architecture's generality quickly became apparent as researchers applied it across NLP tasks. The key innovations — multi-head self-attention enabling parallel computation over full sequences, positional encodings providing sequence order information, and residual connections enabling training of deep stacks — proved robust across scales from millions to hundreds of billions of parameters.

The year 2018 saw two parallel developments that would shape the subsequent trajectory. BERT, introduced by Devlin et al. at Google [32], demonstrated that bidirectional pre-training through masked language modeling yielded representations dramatically superior to those of unidirectional models for understanding tasks. GPT-1, introduced by Radford et al. at OpenAI [11], demonstrated that unidirectional autoregressive pre-training, while less powerful for understanding tasks, enabled fluent and contextually appropriate text generation. The complementarity of these approaches reflected a genuine architectural trade-off that would drive subsequent specialization: encoder-only models for understanding, decoder-only models for generation, and encoder-decoder models for translation.

The scaling era began in earnest with GPT-2 (1.5B parameters, 2019) [12] and accelerated dramatically with GPT-3 (175B parameters, 2020) [13]. GPT-3's demonstration of few-shot in-context learning — competent task performance from only a handful of in-context examples without gradient updates — represented a qualitative capability shift that established the 'emergent abilities' paradigm: capabilities that appear suddenly at sufficient scale rather than improving smoothly with model size [90]. For software engineering specifically, GPT-3 demonstrated surprisingly capable code generation in some scenarios, motivating targeted code-domain pre-training that would produce Codex. Simultaneously, the open-source community's translation of transformer architectures for code began producing specialized models: CodeBERT [33],

GraphCodeBERT [34], and CodeT5 [38] demonstrated strong performance on understanding and translation tasks.

The period 2021–2022 was defined by commercialization and capability consolidation. Codex [18] and GitHub Copilot [19] validated at production scale the hypothesis that LLM code generation could provide practically useful assistance to professional developers. InstructGPT [60] demonstrated that RLHF-aligned models were substantially more useful than base models for real-world applications, initiating the alignment era of model development. AlphaCode [189] demonstrated competitive-programming-level code generation capability, showing that LLMs could approach human-level performance on algorithmic problem-solving at the high end of the difficulty spectrum.

The capability frontier expanded dramatically in 2023 and 2024 through a combination of model scaling, training methodology improvements, and architectural innovations. GPT-4 [36] introduced multimodal understanding and professional-level performance across diverse cognitive tasks. Claude [16] established that Constitutional AI training yielded safer and more reliably aligned models without sacrificing capability. Gemini [17] demonstrated natively multimodal architectures trained jointly on diverse modalities from the ground up. The open-source ecosystem accelerated with LLaMA [12], Code Llama [44], StarCoder2 [43], DeepSeek-Coder [45], and Qwen2.5-Coder [47] providing open-weight frontier-quality models. The emergence of agentic systems — SWE-agent [54], Devin [55], OpenHands [56], and Claude Code [49] — extended LLM capability from single-turn code generation to multi-step autonomous software engineering.

By mid-2025, the state of the field is characterized by: frontier model performance exceeding 90% on HumanEval; agentic systems achieving 37–49% resolution on the challenging SWE-bench Verified benchmark; widespread enterprise adoption of AI-assisted development tools; growing evidence of meaningful productivity improvements; and a rich ecosystem of open-source alternatives providing frontier-quality capability to organizations prioritizing privacy, cost, or customization. The trajectory of progress over the preceding eight years provides compelling grounds for believing that the capabilities of mid-2025 represent only an intermediate point in an ongoing transformation whose ultimate scope remains difficult to assess.

1.4 Research Motivation and Problem Statement

The proliferation of LLMs in software engineering contexts has generated an extensive and rapidly growing body of empirical literature. Individual papers report benchmark results, controlled productivity studies, deployment case studies, and analyses of specific challenges including hallucination, security vulnerabilities, and copyright concerns. The volume of this literature — well over one thousand papers published on

LLMs and software engineering between 2020 and 2025 — reflects the field's extraordinary dynamism and the breadth of research interest it has generated.

Yet despite this richness, the literature exhibits a persistent structural gap: the absence of a comprehensive, systematic, and comparative treatment that integrates findings across the full software development lifecycle. Existing surveys and reviews tend to focus on specific sub-domains — code generation [6], automated testing [130], or bug fixing [49] — without providing the cross-cutting synthesis necessary to understand the collective impact of LLMs on software engineering as an integrated practice. Individual benchmark papers report capability comparisons but do not integrate productivity evidence, challenge analyses, and enterprise deployment lessons. Case study reports from deploying organizations provide practical insights but lack the systematic comparative framework necessary to extract generalizable lessons.

This fragmentation is particularly consequential at the current moment in the technology's development, when practitioners, researchers, educators, and policymakers must make important decisions in the absence of adequate synthetic guidance. Software engineering educators must decide how to restructure curricula that have assumed students will spend significant time learning to write code that AI tools can now generate in seconds. Enterprise technology leaders must decide which tools to deploy, how to govern AI code contributions, and how to measure the productivity impacts of adoption investments. Researchers must identify which open problems are most consequential and most tractable. Policymakers developing AI regulation must understand the specific risk landscape of LLMs in software engineering contexts.

The central research problem may be stated as follows: despite the rapid proliferation of large language models in software engineering contexts, the growing empirical evidence of their impact, and the scholarly attention directed at specific sub-domains, no comparative analysis of LLM applications, capabilities, limitations, ethical implications, and research priorities across the full software development lifecycle has been produced. This gap impedes informed decision-making by practitioners, hinders research prioritization, and limits the development of appropriate governance frameworks. Addressing it represents a scholarly contribution of genuine urgency.

1.5 Research Objectives and Questions

Five primary research objectives collectively address the identified gap:

- (i) To provide a comprehensive and systematic review of the theoretical and architectural foundations of large language models as they apply to software engineering tasks, establishing the technical basis for understanding both capabilities and limitations.

- (ii) To systematically identify, categorize, and analyze the full range of LLM applications across the software development lifecycle, synthesizing empirical evidence for each application domain.
- (iii) To conduct a rigorous comparative analysis of leading LLM-based tools and systems using standardized evaluation metrics, synthesizing benchmark results and deployment evidence to support tool selection decisions.
- (iv) To examine six real-world enterprise deployment case studies in depth, extracting lessons regarding architecture, workflow integration, productivity measurement, and business impact.
- (v) To systematically analyze the principal technical, ethical, environmental, and governance challenges of LLM adoption in software engineering, and to delineate a prioritized forward-looking research agenda.

Five research questions structure the investigation:

RQ1: What architectural characteristics of large language models fundamentally enable their application to software engineering tasks, and how do these characteristics translate to specific capability profiles?

RQ2: How do leading LLMs and LLM-based tools compare on established software engineering benchmarks, and what do these comparisons reveal about the practical tool selection landscape?

RQ3: What are the demonstrated impacts of LLM-based tools on developer productivity in real-world deployments, and what factors moderate these impacts?

RQ4: What are the principal technical, ethical, environmental, and governance challenges that must be addressed for responsible enterprise adoption of LLMs in software engineering?

RQ5: What are the most promising future research directions for the next generation of LLM-based software engineering tools, and what progress is required along each direction?

1.6 Significance of the Study

The significance of the research spans four dimensions: academic, practical, educational, and societal.

Academic Significance

Academically, the study makes a distinctive contribution through its scope, depth, and integrative approach. Applying the SLR methodology to over 220 primary sources across eight activity domains produces a synthesis that few individual papers have

attempted. The comparative framework — evaluated across ten dimensions for twelve tools and systems — provides a structured reference that prior reviews have not offered. The research agenda, covering thirteen priority directions with time horizon estimates and readiness assessments, gives future researchers a concrete starting point. The IEEE-referenced bibliography and systematic methodology also raise the standards of evidence in a field that has sometimes been driven more by vendor claims and benchmark competition than by careful empirical analysis.

Practical Significance

For practitioners — software engineering managers, technology leaders, architects, and developers — the thesis provides decision-relevant synthesis in several areas. The comparative analysis of tools across cost, performance, enterprise readiness, and integration characteristics enables informed tool selection. The six enterprise case studies provide concrete, source-attributed evidence of productivity impacts and deployment challenges. The challenge analysis chapter provides a structured risk assessment for organizations contemplating adoption. The recommendation section in the conclusion translates research findings into practical guidance for common deployment scenarios. In an environment characterized by marketing claims and rapid tool turnover, evidence-based guidance grounded in systematic research provides genuine decision support value.

Educational Significance

The educational implications of LLMs in software engineering are profound and contested. LLMs can generate syntactically correct solutions to many standard programming assignments, raising fundamental questions about how computer science and software engineering should be taught. The analysis of LLM capabilities, limitations, and pedagogical implications developed here provides a foundation for curriculum redesign discussions. Understanding where LLMs perform well (routine code generation, documentation) and where they fall short (complex architectural reasoning, security-sensitive implementation) may help educators redesign learning outcomes, assignments, and assessments to develop human capabilities that complement rather than duplicate AI outputs. The analysis thereby contributes to one of the more consequential educational debates of the current decade.

Societal Significance

The societal implications of LLMs in software engineering extend to labor markets, economic inequality, environmental sustainability, and the governance of critical infrastructure. The analysis of labor market effects, environmental costs, accessibility implications, and governance challenges addresses these broader dimensions explicitly. As software increasingly mediates virtually every aspect of contemporary life,

the transformation of software engineering through AI has ramifications that extend beyond the technical community. A scholarly treatment of these implications, grounded in systematic evidence rather than speculative projection, serves the broader public interest in understanding and governing a technology of significant societal consequence.

1.7 Scope, Delimitations, and Limitations

The scope is limited specifically to transformer-based large language models with parameter counts exceeding one billion parameters, applied to tasks within the software engineering domain. Several delimitations define the boundaries of the research:

The study is delimited to software engineering in the sense of professional and semi-professional software development for production systems, as distinct from end-user programming, scientific computing, or embedded systems programming (though examples from these domains are occasionally referenced for illustrative purposes). The temporal scope is 2017 to mid-2025, corresponding to the period from the introduction of the Transformer architecture to the date of thesis completion. The language scope is English-language publications; while important research exists in Chinese, Japanese, and other languages, systematic coverage of non-English literature was beyond the scope of this study's resources.

Several limitations must be explicitly acknowledged. First, the field evolves at a pace that ensures some findings will be partially outdated by the time of publication; the thesis addresses this by focusing on methodological and theoretical findings that are more durable than specific performance numbers, and by providing context for interpreting benchmark comparisons that may shift with subsequent model releases. Second, proprietary models are evaluated primarily through publicly available benchmark results and vendor-reported case studies; the opacity of proprietary training data, architectural details, and internal evaluations limits the depth of analysis possible for closed models. Third, no original empirical experiments involving human participants or novel model evaluations are reported, as these would have required resources and infrastructure beyond the scope of a BCA Honours thesis; the research is therefore based on secondary analysis of existing empirical work. Fourth, publication bias may cause the literature to overrepresent positive results; the thesis acknowledges and, where possible, corrects for this through explicit attention to limitations and negative findings reported in the literature.

1.8 Research Methodology Overview

The methodological approach adopted in this thesis is a systematic literature review (SLR) augmented by secondary comparative empirical analysis. The SLR methodology, formalized by Kitchenham and Charters [25] for software engineering

research, provides a replicable, rigorous, and transparent process for identifying, evaluating, and synthesizing research relevant to specific research questions. The SLR's defining characteristics — comprehensive search strategies, explicit inclusion and exclusion criteria, structured data extraction, and synthesis — distinguish it from informal narrative reviews and provide protection against selection bias.

Search Strategy

The literature search was conducted across five primary databases: ACM Digital Library, IEEE Xplore, SpringerLink, arXiv (for preprints in AI and machine learning), and Google Scholar. Search terms were organized around three conceptual clusters. The first cluster addressed large language models and related terms: 'large language model,' 'LLM,' 'transformer,' 'GPT,' 'BERT,' 'Codex,' 'code generation model,' 'neural code synthesis,' 'pre-trained language model.' The second cluster addressed software engineering activities: 'code generation,' 'automated testing,' 'debugging,' 'code review,' 'software documentation,' 'DevOps automation,' 'CI/CD,' 'fault localization,' 'program repair.' The third cluster addressed evaluation and measurement: 'benchmark,' 'HumanEval,' 'SWE-bench,' 'developer productivity,' 'hallucination,' 'code correctness.' Boolean combinations of these clusters, tailored to each database's search syntax, yielded an initial corpus of approximately 1,800 papers.

Screening and Selection

The initial corpus was reduced through a two-stage screening process. Title and abstract screening eliminated papers that were clearly outside scope, producing a reduced corpus of approximately 450 papers for full-text review. Full-text review against explicit inclusion and exclusion criteria produced the final corpus of 220 primary sources. Inclusion criteria required that sources: (i) present empirical results, systematic analysis, or foundational theoretical contributions directly relevant to LLMs in software engineering; (ii) be published in peer-reviewed venues or as technical reports from established research organizations; and (iii) be available in English. Exclusion criteria eliminated sources that were purely opinion pieces without empirical grounding; duplicated findings covered by higher-quality primary sources; or addressed closely adjacent topics — general NLP applications, hardware design, robotics — without direct software engineering relevance.




Figure 3.1: PRISMA Flow Diagram — Systematic Literature Review

Data Extraction and Synthesis

From each included paper, structured data were extracted covering: research questions and objectives; methodology and evaluation approach; key findings and quantitative results; limitations acknowledged by authors; and relationships to other included papers. These extracted data were organized by software engineering activity domain and by cross-cutting theme (challenges, future directions, evaluation methodology) using a data extraction form applied consistently across all included sources. Synthesis proceeded through thematic analysis across papers within each domain, identifying areas of convergence and divergence, and resolving apparent contradictions through examination of methodological differences and contextual factors. The comparative empirical analysis component synthesized benchmark performance data from published model evaluation papers, independent research evaluations, and vendor-reported results, with source attribution and uncertainty acknowledgment for figures drawn from a single source.

Quality Assessment

Each included source was assessed for quality across four dimensions: methodological rigor (appropriate research design, adequate sample sizes, controlled

conditions); validity of conclusions (claims supported by evidence, appropriate statistical treatment); transparency (methods sufficiently described for replication); and relevance (direct contribution to one or more research questions). Sources with significant quality concerns were downweighted in the synthesis and their limitations explicitly noted. Overall, the included literature is of high quality, with the majority published in leading venues including NeurIPS, ICML, ICLR, ICSE, FSE, ASE, and peer-reviewed journals including IEEE Transactions on Software Engineering and ACM TOSEM.

1.9 Principal Contributions

Five key analytical contributions are made to the existing literature:

Contribution 1 — Unified Taxonomy: A taxonomy of LLM applications across eight major software engineering lifecycle activities, providing a structured framework that encompasses code generation, explanation, refactoring, bug detection, debugging, code review, testing, documentation, and DevOps automation. The framework addresses a gap in prior surveys, which tend to focus on subsets of the lifecycle rather than offering a lifecycle-spanning view.

Contribution 2 — Comparative Analysis: A comparative analysis of twelve leading models and agentic systems across ten evaluation dimensions, synthesizing benchmark results from HumanEval, SWE-bench Verified, MBPP, MultiPL-E, and BigCodeBench alongside cost, latency, context window, hallucination rate, and enterprise readiness assessments. This multi-dimensional framework supplements individual benchmark papers that address isolated capability dimensions.

Contribution 3 — Case Studies: Six enterprise deployment case studies (GitHub Copilot, Google Gemini Code Assist, Amazon Q Developer, Anthropic Claude Code, Sourcegraph Cody, OpenHands), synthesizing public deployment documentation, conference presentations, and published studies to extract generalizable lessons regarding architecture, integration, productivity measurement, business impact, and challenge management.

Contribution 4 — Challenge Analysis: An analysis of ten challenge and risk categories — hallucination, security vulnerabilities, intellectual property, developer overreliance, environmental cost, bias and fairness, governance and accountability, prompt injection, context window limitations, and benchmark contamination — providing a structured risk framework that supports enterprise adoption planning.

Contribution 5 — Research Agenda: A forward-looking research agenda identifying thirteen priority directions with estimated time horizons, key technical and methodological challenges, current readiness assessments, and connections to existing

preliminary work. The agenda is intended as a resource for researchers entering the field and for research program planners in academia and industry.

1.10 Organization of the Thesis

The thesis is organized into nine chapters. Each builds on those preceding it while remaining accessible as a standalone reference for readers interested in specific sub-domains.

Chapter 2, the Literature Review, provides comprehensive coverage of the scholarly foundations for LLMs in software engineering. Organized roughly chronologically from foundational AI research through current agentic systems, the chapter synthesizes twenty-seven thematic literature subsections spanning artificial intelligence foundations, deep learning, Transformer architecture, bidirectional pre-training, code-specialized models, frontier general models, agentic frameworks, evaluation benchmarks, and cross-cutting research themes including hallucination, economic impact, and multilingual support.

Chapter 3, Foundations of Large Language Models, provides the technical depth necessary for a rigorous understanding of LLM capabilities and limitations. Covering tokenization, embeddings, self-attention mechanisms, feed-forward networks, architectural variants (decoder-only, encoder-decoder, MoE), training pipelines, scaling laws, inference optimization, alignment methodology, in-context learning theory, long-context architectures, and emergent abilities, this chapter provides the conceptual vocabulary used throughout subsequent chapters.

Chapter 4, LLM Applications in Software Engineering, presents the comprehensive taxonomy of application domains and systematically surveys empirical evidence for each. The fifteen application categories covered range from code generation and explanation through refactoring, bug detection, debugging, code review, test generation, documentation, DevOps automation, SQL generation, API generation, and code migration to the emerging domain of requirements-to-code synthesis.

Chapter 5, Comparative Study of LLM-Based Software Engineering Tools, presents the comprehensive evaluation framework and comparative analysis of twelve tools and systems. The chapter concludes with systematic analysis of prompt engineering strategies, cost-performance trade-offs, and the open-source versus proprietary decision for organizational contexts.

Chapter 6, Case Studies, examines six real-world enterprise deployments in depth — GitHub Copilot, Google Gemini Code Assist, Amazon Q Developer, Anthropic Claude Code, Sourcegraph Cody, and OpenHands — synthesizing available evidence into structured case studies and extracting cross-cutting lessons for enterprise practitioners.

Chapter 7, Challenges and Risks, provides systematic analysis of ten major challenge categories, discussing their technical causes, empirical evidence, practical manifestations, and current and prospective mitigation strategies.

Chapter 8, Future Research Directions, presents the research agenda in thirteen thematic areas, organized by time horizon and research readiness, providing a structured resource for researchers and research program planners.

Chapter 9, Conclusion, synthesizes the key findings across all research questions, discusses theoretical and practical implications, acknowledges study limitations, provides recommendations for practitioners, and offers closing reflections on the field's trajectory.

The work concludes with a reference list of all 220+ cited sources, followed by four appendices: Appendix A provides detailed descriptions of major benchmarks and datasets; Appendix B presents prompt engineering templates; Appendix C provides a glossary of technical terms; and Appendix D provides additional methodological details for the systematic literature review.

CHAPTER 2: LITERATURE REVIEW

2.1 Artificial Intelligence: Foundations and Historical Development

Artificial intelligence, as a formal field of scientific inquiry, was inaugurated at the Dartmouth Workshop of 1956, where McCarthy, Minsky, Shannon, and a group of pioneering researchers articulated the central aspiration of constructing machines capable of mimicking cognitive functions associated with the human mind [20]. The audacious proposal that submitted to the Rockefeller Foundation for funding stated that 'every aspect of learning or any other feature of intelligence can in principle be so precisely described that a machine can be made to simulate it.' This foundational claim — that intelligence is computable and that any cognitive function that can be precisely described can be mechanized — remains simultaneously the motivating premise and the central controversy of the field.

The subsequent six decades witnessed oscillating periods of enthusiasm and disillusionment — the so-called AI winters of the 1970s and 1980s — driven by the persistent gap between aspirational claims and delivered capabilities. The first AI winter, following the Lighthill Report's critique in 1973, reflected limitations of rule-based symbolic AI in dealing with the ambiguity, context-dependence, and combinatorial complexity of real-world problems. The second AI winter, in the late 1980s and early 1990s, reflected the brittleness and maintenance cost of expert systems that encoded knowledge manually. Each winter was followed by renewed progress as new approaches, larger datasets, and greater computational resources addressed the limitations that had produced stagnation.

The field's major intellectual traditions provide distinct frameworks for understanding LLMs' position in the history of AI. Symbolic AI, dominant from the 1950s through the 1980s, represented knowledge explicitly as logical rules and semantic networks, enabling reasoning through formal inference. Expert systems such as MYCIN, DENDRAL, and XCON demonstrated that symbolic AI could achieve impressive performance within narrow, well-defined domains but struggled to scale to the breadth and contextual sensitivity of general intelligence. Connectionist AI, rooted in artificial neural networks inspired by biological neurons, experienced multiple waves of development — the perceptron era of the late 1950s, the backpropagation revival of the 1980s, and the deep learning revolution of the 2010s — before achieving the practical successes that now dominate the field. The tension between symbolic and connectionist approaches has re-emerged in the context of LLMs, with neuro-symbolic integration research seeking to combine the representational power of neural networks with the logical structure of symbolic reasoning [133].

Probabilistic AI, including Bayesian networks, probabilistic graphical models, and Markov decision processes, provided principled frameworks for reasoning under uncertainty that proved essential for speech recognition, machine translation, and computer vision in the pre-deep learning era. The integration of probabilistic reasoning with neural representations — through variational autoencoders, probabilistic graphical models, and Bayesian neural networks — remains an active research frontier. Evolutionary computation, drawing inspiration from biological evolution through genetic algorithms and genetic programming, explored a distinct path to automated problem-solving that has found applications in architecture search, hyperparameter optimization, and code synthesis, but has not emerged as the dominant paradigm for LLM development.

The philosophical dimensions of AI bear substantively on how we interpret LLM capabilities in software engineering. Turing's imitation game [190], later popularized as the Turing Test, proposed behavioral equivalence with human intelligence as the criterion for artificial intelligence — a criterion that sophisticated LLMs appear to meet in many narrow conversational domains, raising questions about whether behavioral equivalence implies genuine understanding. Searle's Chinese Room argument [191] challenged the functionalist premise that appropriate input-output behavior is sufficient for genuine understanding, arguing that syntactic symbol manipulation — however sophisticated — lacks the semantic grounding of genuine cognition. The Chinese Room's relevance to LLMs is contested: critics note that LLMs' training on human-generated text provides a form of grounding in human experience and meaning that differs from Searle's hypothetical, while proponents maintain that the fundamental argument about the insufficiency of syntax for semantics applies regardless of training methodology.

For the purposes of this study, a pragmatic stance is adopted: whether LLMs genuinely 'understand' code or perform sophisticated statistical pattern matching that functionally resembles understanding is philosophically interesting but practically secondary. What matters for software engineering is the observable behavior of LLMs — their ability to generate functionally correct code, identify bugs, explain programs, and participate in multi-step development workflows — and the empirical evidence for these capabilities reviewed in the chapters that follow. The philosophical dimension becomes more practically salient when considering reliability, since whether models have genuine semantic understanding or are pattern-matching on surface features bears on the conditions under which their outputs can be trusted.

2.2 Machine Learning: Paradigms, Principles, and Applications

Machine learning, the subfield of artificial intelligence concerned with algorithms that improve their performance through experience rather than explicit programming, has produced the most practically impactful AI systems of the contemporary period.

Mitchell's canonical formal definition — that a computer program is said to learn from experience E with respect to tasks T and performance measure P if its performance at T , as measured by P , improves with E — captures the essential property that distinguishes machine learning from conventional programming: the system's behavior is determined not by explicit rules coded by programmers but by patterns extracted from data [25].

The principal paradigms of machine learning are supervised learning, unsupervised learning, semi-supervised learning, and reinforcement learning. In supervised learning, a model is trained on labeled examples — pairs of input and desired output — to learn a mapping function that can generalize to new inputs. Classification tasks (assigning inputs to categorical labels), regression tasks (predicting continuous outputs), and structured prediction tasks (generating structured outputs like sequences or trees) are all amenable to supervised learning with appropriate loss functions and model architectures. For software engineering, supervised learning from labeled code-comment pairs, bug-fix pairs, and test-coverage pairs has enabled specialized code intelligence models.

Unsupervised learning, in which the model discovers latent structure in unlabeled data without external supervision, is the dominant paradigm for pre-training large language models. The masked language modeling objective of BERT and the causal language modeling objective of the GPT series are both forms of self-supervised learning — a subclass of unsupervised learning in which supervision signals are derived from the data itself rather than external annotation. This elegant bootstrapping of supervision from the structure of the data enables training on arbitrarily large corpora without the bottleneck of human annotation, which is the fundamental source of the scalability advantage that large pre-trained models hold over task-specific supervised models [26].

Reinforcement learning, in which an agent learns to take actions in an environment to maximize cumulative reward, plays an increasingly important role in LLM development through RLHF and its variants. The application of reinforcement learning to language model fine-tuning differs from classical RL in that the 'actions' are token generation choices, the 'environment' is the space of human reactions to model outputs, and the 'reward' is a learned model of human preferences. The combination of supervised pre-training to learn the basic language modeling capability and reinforcement learning to align the model's behavior with human values has proven more effective than either approach alone for producing practically useful AI assistants [27, 59, 60].

The bias-variance trade-off, a foundational concept in machine learning theory, describes the fundamental tension between model complexity and generalization. High-complexity models (low bias) can capture complex patterns in training data but tend to overfit, producing high variance in their predictions on new data. Low-complexity models (high bias) generalize more stably but may systematically fail to capture patterns in the training data. For LLMs, which are enormous neural networks trained on vast data,

the classical bias-variance framework is insufficient: these models simultaneously achieve low training loss (suggesting high capacity) and strong generalization (suggesting effective regularization), a combination explained by the double-descent phenomenon and by the implicit regularization effects of stochastic gradient descent on over-parameterized models [192].

Transfer learning — the use of representations learned on one task to improve performance on a different but related task — is the paradigm that makes large language models practically tractable for software engineering applications. Training a code generation model from scratch on code data would require enormous resources and produce a model lacking the world knowledge and reasoning capabilities that enable flexible problem-solving. Instead, the standard approach pre-trains a general-purpose language model on a diverse corpus of text and code, developing rich general representations, and then fine-tunes this pre-trained model on code-specific tasks or provides task instructions in context. The transfer learning paradigm enables practitioners to leverage hundreds of millions of dollars of pre-training compute investment — which only large organizations can afford — for a wide range of downstream tasks through fine-tuning or prompting, which are accessible at much lower cost.

2.3 Deep Learning: Architectures and Training Foundations

Deep learning, the use of artificial neural networks with multiple layers of learnable transformations to learn hierarchical representations from raw data, has proven to be the most powerful and general approach to machine learning yet discovered for high-dimensional, unstructured data including natural language and source code [28]. The historical development of deep learning traces a path from the perceptron of the late 1950s through the backpropagation of the 1980s, the convolutional network innovations of the 1990s and 2000s, and the deep learning revolution initiated by AlexNet's triumph in the 2012 ImageNet Large Scale Visual Recognition Challenge.

The backpropagation algorithm, formalized and popularized by Rumelhart, Hinton, and Williams in 1986 [29], provided the computational mechanism for efficiently computing gradients of the loss function with respect to all parameters in a deep network. By applying the chain rule of calculus through the network's layers in reverse order, backpropagation enables gradient-based optimization of multi-layer networks with millions of parameters — a capability that, combined with stochastic gradient descent, forms the foundation of all practical deep learning training. The algorithmic correctness of backpropagation was established decades before it became practically useful, waiting for the combination of sufficient computational power, adequate datasets, and architectural innovations that would make training deep networks tractable.

Convolutional Neural Networks (CNNs), developed by LeCun and collaborators in the 1980s and 1990s [193, 194], introduced the key innovations of local connectivity

and parameter sharing that exploit the translation invariance present in image data. By sharing convolutional filter weights across all spatial positions, CNNs dramatically reduce the number of parameters required relative to fully connected networks, enabling efficient learning from moderately sized datasets. While CNNs found their most prominent applications in computer vision, the general principle of exploiting data structure — using architectural inductive biases that match the structure of the data — proved influential across deep learning, including in code models that exploit the hierarchical structure of programming languages.

Recurrent Neural Networks (RNNs), and particularly Long Short-Term Memory (LSTM) networks introduced by Hochreiter and Schmidhuber in 1997 [9], extended deep learning to sequential data by maintaining a hidden state that encodes information about prior sequence elements. LSTM's gating mechanisms — the forget gate, input gate, and output gate — enabled selective retention and updating of hidden state information, addressing the vanishing gradient problem that prevented simple RNNs from learning long-range dependencies. Before the Transformer, LSTM was the architecture of choice for sequence-to-sequence tasks including neural machine translation, code completion, and code summarization. The shift from LSTM to Transformer was not driven by theoretical necessity — LSTMs can in principle represent any computable function — but by the practical advantages of parallel training and the demonstrated superior performance of attention-based models on long sequences.

The deep learning revolution, conventionally dated to AlexNet's ImageNet success in 2012 [30], combined three ingredients that had individually been available for years but had not previously been assembled at sufficient scale: large labeled datasets (ImageNet's 1.2 million labeled images), GPU computing enabling parallel matrix multiplications at the scale required for deep networks, and architectural innovations including ReLU activations, dropout regularization, and data augmentation. The ImageNet breakthrough initiated a cascade of deep learning advances across computer vision, speech recognition, and natural language processing, ultimately creating the ecosystem that produced LLMs.

2.4 The Transformer Architecture: Technical Foundations

The Transformer architecture, introduced by Vaswani et al. at Google Brain in 'Attention Is All You Need' (2017) [10], is the architectural foundation upon which virtually all contemporary large language models are built. The paper's title captures its central claim: the self-attention mechanism, used in Transformers as the sole mechanism for relating sequence positions, is sufficient — without recurrence or convolution — for achieving state-of-the-art performance on sequence modeling tasks. The subsequent five years of rapid scaling and performance improvements have overwhelmingly validated this claim.

The original Transformer employed a symmetric encoder-decoder structure. The encoder processes the complete input sequence, with each position attending to all other positions through bidirectional self-attention, producing a set of contextual representations encoding the input's meaning. The decoder autoregressively generates the output sequence, with each position attending to all preceding decoder positions through masked (causal) self-attention and attending to the encoder's representations through cross-attention. This architecture enables complete input-output sequence transformation with full bidirectional context for the input and appropriately causal generation for the output.

The scaled dot-product attention mechanism is the core computational primitive of the Transformer. Given input representations packed as rows of a matrix X , three linear projections produce query matrix $Q = XW_Q$, key matrix $K = XW_K$, and value matrix $V = XW_V$. The attention weights are computed as the row-normalized dot products of queries and keys, scaled by the square root of the key dimension to prevent large magnitudes: $\text{Attention}(Q, K, V) = \text{softmax}(QK^T / \sqrt{d_k}) V$. Each row of the output is a weighted sum of value vectors, with weights reflecting the learned relevance of each key to the corresponding query. Causal (decoder) attention is implemented by masking the attention weight matrix to zero out attention from each position to all subsequent positions, enforcing the autoregressive constraint that predictions may only attend to prior context.

Multi-head attention extends scaled dot-product attention by computing h parallel attention functions with reduced dimensionality $d_k = d_{\text{model}} / h$, concatenating the results: $\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h) W_O$ where $\text{head}_i = \text{Attention}(QW_Q^i, KW_K^i, VW_V^i)$. The motivation is that different heads can attend to different aspects of the input simultaneously — one head might capture syntactic relationships while another captures semantic co-reference — providing richer representations than a single attention function. Analysis of attention patterns in code-trained models has revealed consistent head specialization: certain heads reliably attend to matching brackets, identifier definitions, or function arguments across different code examples, providing evidence that multi-head attention learns structured code representations.

Each Transformer layer alternates between multi-head self-attention and a position-wise feed-forward network (FFN), with residual connections and layer normalization applied around each sub-layer. The position-wise FFN applies two linear transformations with a non-linear activation: $\text{FFN}(x) = \max(0, xW_1 + b_1)W_2 + b_2$ (using ReLU) or $\text{FFN}(x) = \text{SwiGLU}(xW, xV)$ (using SwiGLU as in modern LLMs [78]). The FFN dimension is typically four times the model dimension, providing a wide bottleneck that has been associated with factual knowledge storage — the model's

'memories' of specific code patterns and API usages appear to be distributed across FFN weights [195].

Positional encoding addresses the Transformer's fundamental limitation: unlike RNNs, which process tokens sequentially and inherently encode position through processing order, self-attention operates on sets of token representations with no inherent notion of ordering. The original Transformer used fixed sinusoidal position encodings added to token embeddings. Modern LLMs employ Rotary Position Embedding (RoPE) [71], which encodes position as a rotation applied to the query and key vectors before computing attention, enabling relative position information to be computed as part of the dot product. RoPE generalizes more gracefully to sequence lengths exceeding those seen during training than additive position encodings, contributing to the success of long-context extension techniques.

Computational Complexity and Context Length Scaling

The dominant limitation of standard Transformer attention is its $O(n^2 \cdot d)$ computational complexity in both time and memory with respect to sequence length n , arising from the $n \times n$ attention weight matrix computed for each layer and head. For a sequence of 1,000 tokens, this is manageable; for 100,000 tokens (approximately 75,000 words or several hundred pages of code), the memory requirement becomes prohibitive on current hardware. This quadratic scaling has motivated a substantial research effort in efficient attention alternatives.

FlashAttention [124], introduced by Dao et al. in 2022, achieves exact attention computation with reduced memory requirement through GPU memory hierarchy awareness: rather than materializing the full $n \times n$ attention matrix in GPU high-bandwidth memory (HBM), FlashAttention tiles the computation to keep intermediate results in fast SRAM, dramatically reducing HBM access and enabling longer sequences within the same hardware budget. FlashAttention-2 and FlashAttention-3 introduced additional optimizations achieving further throughput improvements. Ring Attention [201] distributes the attention computation across multiple devices for sequences exceeding single-device memory. These engineering advances have been critical enablers of the 100K+ context windows in modern frontier models.

2.5 BERT and Bidirectional Pre-training for Code Understanding

BERT (Bidirectional Encoder Representations from Transformers), introduced by Devlin et al. at Google AI Language in 2018 [32], established the pre-training and fine-tuning paradigm as the dominant approach for natural language understanding tasks and subsequently for code understanding. BERT's defining innovation was bidirectional pre-training through Masked Language Modeling (MLM): by randomly masking 15% of input tokens and training the model to predict the masked tokens from the full

bidirectional context, BERT developed rich contextual representations in which the meaning of each token is informed by all tokens in both directions.

The empirical success of BERT was immediate and dramatic: fine-tuned BERT models achieved state-of-the-art performance on eleven diverse NLP benchmarks simultaneously at release, with improvements over prior state-of-the-art ranging from modest to substantial. On the GLUE benchmark suite, BERT improved the aggregate score from 72.8 to 80.4 points; on SQuAD 2.0 question answering, BERT achieved superhuman performance. This simultaneous success across diverse tasks was unprecedented and established the transferability of pre-trained representations as a robust phenomenon rather than a task-specific fortunate result.

For code intelligence, the BERT paradigm was adapted through several specialized variants. CodeBERT [33], from Microsoft Research, was pre-trained on bimodal data pairing programming language tokens with corresponding natural language documentation from GitHub repositories. Using both the standard MLM objective and a replaced token detection objective (predicting whether each token is original or generated by a model), CodeBERT learned representations capturing both the syntactic structure of code and its semantic relationship to natural language descriptions. Evaluated on code search (finding code given a natural language query) and code documentation generation, CodeBERT substantially outperformed prior models on both tasks.

GraphCodeBERT [34] extended CodeBERT by incorporating data flow graphs representing the variable definition-use relationships in code. The graph-enhanced pre-training introduced two additional objectives: predicting edges in the data flow graph and aligning node representations across code and natural language for corresponding identifiers. By explicitly modeling which variables in one location in the code affect which variables in other locations, GraphCodeBERT learned representations that better captured the semantic structure of code beyond the surface-level token sequence. Performance improvements over CodeBERT on code clone detection, code translation, and code refinement tasks validated the value of structural code information for representation learning.

UniXcoder [123] further advanced code model capabilities by incorporating Abstract Syntax Tree (AST) information through a prefix-guided attention scheme, enabling the same unified encoder-decoder architecture to serve both understanding and generation tasks with state-of-the-art performance on tasks spanning code search, code generation, code clone detection, code repair, code translation, and code summarization. The unification of diverse code intelligence tasks under a single architecture eliminated the need for task-specific model design and demonstrated that the diversity of code intelligence tasks could be addressed through a flexible architecture with task-appropriate conditioning.

Encoder-Only vs. Decoder-Only for Code

The architectural choice between encoder-only models (BERT-style) and decoder-only models (GPT-style) reflects a fundamental capability trade-off relevant to software engineering. Encoder-only models develop superior representations for understanding tasks — code search, clone detection, vulnerability classification — because bidirectional attention enables each representation to incorporate the full code context. However, they cannot perform autoregressive generation: since each token's representation depends on all other tokens including those 'to the right,' there is no principled way to extend encoder-only models to generate new code token by token.

Decoder-only models sacrifice the bidirectional context of encoders — in a decoder-only model, each token can only attend to preceding tokens — in exchange for the ability to perform autoregressive generation. This trade-off makes decoder-only models well-suited for code completion (the dominant practical application of LLMs in IDE assistants) while potentially disadvantaging understanding tasks. In practice, the vast scale of decoder-only models trained on diverse code corpora — hundreds of billions of parameters trained on trillions of tokens — has overcome this theoretical disadvantage, with frontier decoder-only models achieving strong performance even on tasks like code search and vulnerability detection that seem more naturally suited to encoder-only architectures.

2.6 Code-Specialized Models: From CodeBERT to Frontier Systems

The development of code-specialized models from 2020 to 2025 followed a trajectory from narrow task-specific fine-tuned systems to large general-purpose models with strong code capabilities as an emergent property of scale and diverse training data. This section surveys the major models in approximate chronological order, highlighting the key innovations each contributed.

CodeBERT [33] and GraphCodeBERT [34], discussed in Section 2.5, established the foundations of neural code intelligence through BERT-style pre-training with code-specific objectives. T5 [37] and CodeT5 [38] extended this to encoder-decoder generation through the text-to-text framework. PLBART [200], a BART-based model pre-trained with denoising objectives on code and natural language, demonstrated strong performance on code summarization, generation, and translation tasks, particularly for low-resource programming languages with limited training data.

InCoder [98], a causal language model pre-trained with a fill-in-the-middle objective enabling infilling of arbitrary code regions, demonstrated that causal models could support both left-to-right code generation and bidirectional infilling — the capability required for intelligent insertion at cursor positions within existing code. The SantaCoder and StarCoder [41] models from the BigCode project brought the open-

source code model ecosystem to frontier performance levels, demonstrating that transparent, community-governed model development could produce models competitive with proprietary alternatives.

Code Llama [44] demonstrated that the increasingly capable general-purpose LLaMA model family [12] could be specialized for code through continued pre-training on code data, producing a family of open-weight code models ranging from 7B to 70B parameters. The larger Code Llama 70B variant achieves 53% pass@1 on HumanEval, approaching the performance of early proprietary code models. DeepSeek-Coder [45] and its successor DeepSeek-Coder-V2 [46] demonstrated that scratch-trained code models with high-quality curated training data and efficient MoE architectures could achieve performance competitive with closed-source frontier models. Qwen2.5-Coder [47] extended strong code performance to over 40 programming languages, addressing the multilingual coverage gaps of earlier specialized models.

The Emergence of General Models with Strong Code Capability

A significant trend in the period 2023-2025 is the convergence of general and code-specialized model capabilities: frontier general-purpose models including GPT-4, Claude 3.5, and Gemini 1.5 achieve code generation performance comparable to or exceeding specialized code models at similar scales. This convergence reflects the contribution of diverse training data — exposure to tutorials, technical discussions, Stack Overflow answers, and documentation alongside source code — to developing a richer understanding of programming that transcends pure code pattern matching. A model that understands the natural language description of an algorithm, the mathematical concepts it implements, and the API documentation for relevant libraries can generate code that integrates these diverse knowledge sources in ways that a code-only model cannot.

2.7 The GPT Series: Autoregressive Pre-training and Scaling

The GPT (Generative Pre-trained Transformer) series, developed by OpenAI beginning in 2018, established the decoder-only autoregressive pre-training paradigm as the dominant approach for large-scale language model development and showed that capability improvements of striking magnitude were achievable through the straightforward scaling of model size, training data, and compute [11, 12, 13, 36].

GPT-1 [11], with 117 million parameters trained on the BooksCorpus dataset, demonstrated that unsupervised pre-training followed by discriminative fine-tuning on specific tasks outperformed models trained from scratch on those tasks. The key insight was that the pre-training objective — predicting each next token from all preceding tokens — developed rich linguistic representations as a byproduct of the prediction task, representations that could be rapidly adapted to diverse downstream tasks through fine-tuning. While GPT-1's absolute performance was modest, its demonstration of the pre-

training paradigm for generative models established the template for subsequent GPT models.

GPT-2 [12], with parameters up to 1.5 billion (a 12x increase from GPT-1) and training on a much larger and more diverse corpus (the 40GB WebText dataset), demonstrated emergent zero-shot task performance: without any task-specific fine-tuning, the model performed respectably on tasks including reading comprehension, translation, and summarization simply by conditioning on an appropriately framed prompt. The model's fluency was sufficient that OpenAI initially declined to release the full model due to concerns about misuse for text generation at scale — the first high-profile instance of AI safety considerations affecting model release policy. Retrospectively, GPT-2's capabilities appear modest compared to subsequent models, but its demonstration of zero-shot task transfer established that scale alone, without architectural innovation, could produce emergent capabilities.

GPT-3 [13], with 175 billion parameters trained on the 570 GB Common Crawl-filtered dataset plus curated corpora, represented the first clearly 'large' language model in the contemporary sense and introduced the concept of few-shot in-context learning that has become central to LLM deployment. Brown et al. demonstrated that by including a small number of input-output examples in the prompt (the 'context'), GPT-3 could perform tasks including translation, question answering, arithmetic, and code generation without any gradient updates — purely through pattern matching to the in-context examples. This in-context learning capability, which does not require separate fine-tuning for each task, dramatically reduced the practical barrier to applying GPT-3 to diverse applications and established the prompting paradigm as an alternative to fine-tuning for many use cases.

GPT-4 [36], released in March 2023, represented a qualitative capability advance whose technical details — model size, architecture, training data — OpenAI declined to disclose publicly in contrast to prior GPT releases. The model demonstrated human-expert-level performance across a wide range of cognitive tasks: the 90th percentile on the Uniform Bar Examination (U.S. lawyer licensing), 99th percentile on the Biology Olympiad, 88th percentile on the LSAT, and strong performance on professional licensing examinations in medicine and other domains. For software engineering, GPT-4 demonstrated substantially superior performance on complex multi-file coding tasks, algorithmic reasoning, and code explanation relative to GPT-3.5, with HumanEval pass@1 improving from approximately 67% (GPT-3.5) to 85%+ (GPT-4) at release. GPT-4o (Omni), released in May 2024, further integrated visual, text, and audio modalities within a unified architecture while improving efficiency and reducing cost, maintaining strong coding performance at 90.2% HumanEval pass@1.

2.8 T5, CodeT5, and Encoder-Decoder Approaches for Code

T5 (Text-to-Text Transfer Transformer), introduced by Raffel et al. [37] from Google AI in 2019, proposed a unified framework casting every NLP task as text-to-text: both inputs and outputs are always represented as natural language strings, with task type communicated through a text prefix appended to the input. This elegant unification enabled T5 to be trained with a single maximum likelihood objective across all tasks — translation, summarization, question answering, classification — simply by varying the input prefix. Pre-trained on the Colossal Clean Crawled Corpus (C4, 750GB of filtered Common Crawl text), T5 achieved state-of-the-art performance across the GLUE and SuperGLUE benchmarks and established multi-task pre-training as a viable alternative to single-task fine-tuning from pre-trained representations.

CodeT5, introduced by Wang et al. from Salesforce Research [38], adapted the T5 framework for programming language understanding and generation. CodeT5 pre-trained on the CodeSearchNet corpus [199] of 8.35 million code-documentation pairs across six programming languages using both standard masked span prediction and an identifier-aware pre-training objective that leveraged the semantic importance of code identifiers. The identifier-aware objective distinguished between masked identifier tokens and masked non-identifier tokens, training the model to pay particular attention to the semantically rich identifier vocabulary. CodeT5 achieved state-of-the-art performance on five code intelligence tasks at release: code defect detection, code clone detection, code translation, code summarization, and code generation.

CodeT5+ [39], introduced by Wang et al. in 2023, substantially enhanced CodeT5 through bimodal contrastive pre-training objectives that align code and natural language representations across modalities. By training with a contrastive loss that pulls representations of code and its corresponding documentation together in the embedding space while pushing apart non-corresponding pairs, CodeT5+ learned cross-modal representations enabling strong performance on code search and code generation tasks requiring joint understanding of natural language descriptions and code semantics. CodeT5+ models in sizes from 220M to 16B parameters were released as open-source checkpoints, supporting a range of deployment scenarios from resource-constrained environments to high-performance settings.

The encoder-decoder architecture underlying T5 and CodeT5 variants offers particular advantages for tasks with clearly distinct input and output structures: code translation (source language to target language), code summarization (code to natural language), and code repair (buggy code to fixed code) all benefit from the encoder's ability to develop rich bidirectional representations of the input before the decoder generates the output. However, the increasing capability of decoder-only models — which achieve strong performance even on these input-output transformation tasks through appropriate prompting — has reduced the architectural advantage of encoder-

decoder models in practice, leading the research community to converge on decoder-only architectures for most large-scale code model development since 2022.

2.9 Codex: The Commercial Deployment of LLM Code Generation

Codex, introduced by Chen et al. at OpenAI in June 2021 [18], represents the pivotal development that translated LLM code generation capability from academic research into commercial deployment at scale. The technical foundations of Codex are relatively straightforward: starting from GPT-3 (175 billion parameters), OpenAI fine-tuned the model on approximately 159 gigabytes of Python code collected from public GitHub repositories through May 2020, representing approximately 54 million public GitHub repositories filtered for files with .py extension. The fine-tuning used the same causal language modeling objective as the original GPT-3 training.

Despite the technical simplicity of its construction, Codex demonstrated qualitatively different code generation capabilities compared to GPT-3. On the HumanEval benchmark introduced alongside Codex, the 12-billion parameter Codex model achieved 28.8% pass@1 — meaning that a single sample from the model correctly solved 28.8% of the 164 programming problems. With temperature sampling (pass@100), 70.2% of problems were solved with at least one correct sample across 100 attempts, demonstrating that even this relatively modest model could find correct solutions to most HumanEval problems with sufficient sampling. The discrepancy between pass@1 and pass@100 motivates the 'sample and rank' and 'sample and test' strategies that generate multiple candidate solutions and select among them.

The commercial deployment of Codex through GitHub Copilot transformed the developer tool landscape. GitHub Copilot's technical preview in June 2021 attracted over 1.2 million waitlist registrations, reflecting widespread developer interest in AI-assisted coding. The general availability launch in June 2022, initially at \$10/month for individuals and \$19/seat/month for businesses, marked the transition to commercial reality. The rapid scaling to millions of active users provided the first large-scale empirical evidence that LLM code generation could sustain commercial adoption — that the productivity benefit, in real developer workflows across diverse projects, was sufficient to justify subscription cost.

The research methodology behind Codex's HumanEval benchmark deserves examination, as it established a template widely used for subsequent code generation evaluation. HumanEval consists of 164 manually authored Python programming problems, each with a function signature, natural language docstring specification, and a set of test cases for evaluation. The test cases are not provided to the model; success is determined purely by functional correctness when the generated function is executed against the hidden test suite. The pass@k metric — the probability that at least one of k samples from the model passes all test cases — provides a principled way to compare

models at different points on the accuracy-diversity trade-off curve. This evaluation methodology, while simple and transparent, has known limitations: the 164 problems are relatively simple self-contained tasks; the problems are now widely present in training data (benchmark contamination); and functional correctness on hand-crafted test cases does not capture the full range of real-world software quality properties. These limitations motivated the development of harder and more realistic benchmarks including SWE-bench.

2.10 StarCoder and the Open-Source Code Model Ecosystem

StarCoder, introduced by Li et al. through the BigCode collaboration in 2023 [41], marked a milestone in open-source code generation, showing that transparent, community-governed model development could produce models competitive with proprietary alternatives while providing full research access, customizability, and deployment flexibility. The BigCode project, co-organized by Hugging Face and ServiceNow with participation from over 1,000 researchers across academia and industry, provided an open framework for collaborative model development that prioritized transparency and responsible data practices.

The training data for StarCoder, The Stack [42], is a 6.4 terabyte collection of permissively licensed source code from 358 programming languages assembled from GitHub public repositories. The Stack's construction incorporated several important ethical and practical innovations: an opt-out mechanism allowing GitHub repository owners to remove their code from the training data; a filtering pipeline removing personal information, malicious code patterns, and low-quality examples; and detailed documentation of data sources, filtering methodology, and licensing terms. The opt-out mechanism, through the Am I In The Stack website, received substantial use and demonstrated that responsible data practices were compatible with assembling large-scale training corpora.

StarCoder's architecture — a 15.5 billion parameter decoder-only model with 8,192-token context length, Fill-in-the-Middle training, and Multi-Query Attention (MQA) for inference efficiency — achieved competitive performance with Codex on HumanEval while providing the full model weights and training details that proprietary models withheld. StarCoder demonstrated particular strengths in multilingual code generation, reflecting The Stack's broad language coverage compared to Codex's Python-dominant training corpus. The model's support for Fill-in-the-Middle generation enabled IDE-style completion at arbitrary cursor positions, making it practically suitable for deployment as an IDE assistant.

StarCoder2 [43], released in early 2024, substantially expanded the training corpus to include GitHub issues, pull requests, notebooks, and documentation in addition to source files, improving performance on tasks requiring understanding of the full

software development context. The expanded training produced models more capable of understanding the informal, conversational text characteristic of real development workflows — commit messages, bug reports, code review comments — rather than only the formal code structure of source files. StarCoder2 models in 3B, 7B, and 15B parameter sizes were released to support a range of deployment scenarios, with the 15B model achieving performance competitive with proprietary models trained at similar scales.

2.11 Code Llama: Meta's Contribution to Open Code Intelligence

Code Llama [44], introduced by Roziere et al. at Meta AI Research in August 2023, extended Meta's open-weight LLaMA 2 foundation models with code specialization, producing a family of models that provided frontier code generation quality with open-weight accessibility. The specialization approach — continued pre-training of LLaMA 2 models on a code-rich mixture rather than training from scratch — leveraged the extensive world knowledge and language understanding developed during LLaMA 2's pre-training, producing code models with stronger natural language reasoning capabilities than models trained exclusively on code.

Code Llama was released in three variants: Code Llama (base code model), Code Llama - Instruct (instruction-tuned for conversational use), and Code Llama - Python (further specialized on Python code). Each variant was available in 7B, 13B, 34B, and (subsequently) 70B parameter sizes, providing a full range of performance-efficiency trade-offs for different deployment scenarios. The 7B model, capable of running on consumer hardware with appropriate quantization, brought frontier-quality code generation within reach of individual developers without cloud computing access. The 70B model, achieving 53% pass@1 on HumanEval, demonstrated that open-weight models could approach proprietary model performance at sufficient scale.

Code Llama's long-context fine-tuning variant, supporting contexts up to 100,000 tokens, enabled qualitatively different use cases than short-context alternatives: entire repository scanning, long debugging session continuity, and full-document code review. While the practical utility of 100K-token context for typical coding tasks is limited by the challenge of directing the model's attention to relevant portions of a large context, the capability demonstrated the architectural feasibility of long-context code models and inspired subsequent work on long-context training and attention mechanisms for code. The combination of Code Llama's open weights, permissive licensing (Community License Agreement enabling commercial use), and strong base performance catalyzed a rich ecosystem of fine-tuned variants addressing specialized domains including web development, data science, and scientific computing.

2.12 DeepSeek-Coder: Efficiency and Performance at Scale

DeepSeek-Coder [45], introduced by the DeepSeek AI team in November 2023, showed that models trained from scratch on carefully curated code data could match or exceed models fine-tuned from large general-purpose pre-trained models, challenging the prevailing assumption that general pre-training is a necessary prerequisite for strong code performance. The model's training corpus consisted of 87% source code and 13% natural language related to code — technical documentation, tutorials, and Stack Overflow content — with thorough deduplication and quality filtering that prioritized data quality over raw quantity.

A distinctive feature of DeepSeek-Coder's training approach is its repository-level pre-training: rather than treating individual source files as independent training examples, the training pipeline assembled complete repositories and included cross-file context in training sequences. This design choice improves the model's understanding of how files within a repository relate to each other — which functions are called where, how modules import from each other, how interfaces align across files — a capability essential for real-world code generation tasks that require understanding of multi-file project structure. The repository-level training directly addresses one of the key limitations of file-level code models: their inability to reason about code in the context of the surrounding project.

DeepSeek-Coder-V2 [46], released in May 2024, leveraged a Mixture of Experts architecture to achieve frontier performance at substantially lower inference cost — a notable milestone for the open-source ecosystem. The model activates 21 billion parameters (out of 236 billion total) per token, providing computation equivalent to a 21B dense model while retaining the representational capacity of a much larger model through its 128 experts. Its 90.2% HumanEval pass@1 — matching GPT-4o — while being open-weight and commercially licensed, represented a meaningful democratization of frontier code generation capability. The model also introduced a 128K context window comparable to GPT-4 Turbo, enabling repository-level understanding at practical scale.

2.13 Qwen2.5-Coder: Multilingual and Comprehensive Code Support

The Qwen series of models from Alibaba Cloud's DAMO Academy has established a strong position in the code generation landscape through consistent focus on multilingual coverage and benchmark performance across a wide range of programming languages, including less-commonly represented languages where competing models' training data is limited [47]. The series reflects Alibaba's particular interest in code generation for the diverse technology stacks prevalent in global software development, from languages with significant enterprise usage (COBOL, ABAP) to emerging systems programming languages (Rust, Zig).

Qwen2.5-Coder, the most recent release in the series as of mid-2025, was trained on a 5.5 trillion token corpus spanning over 40 programming languages with careful attention to training data curation. The training pipeline included quality filtering based on multiple criteria: syntactic correctness (filtering code that fails to parse), style consistency (filtering code with problematic formatting), documentation presence (preferring code with inline documentation), and test coverage (preferring code accompanied by test files). This multi-criterion quality filtering, applied at scale, produced a training corpus that the team attributed to Qwen2.5-Coder's consistently strong performance across languages with varied levels of representation in general code corpora.

The model family spans sizes from 0.5B to 72B parameters, with the 72B variant achieving performance competitive with GPT-4 Turbo on several benchmarks including LiveCodeBench and BigCodeBench. The availability of sub-billion parameter variants enables deployment on edge devices and mobile platforms, addressing use cases where cloud round-trip latency or data privacy concerns preclude API-based access. Qwen2.5-Coder's Apache 2.0 licensing enables commercial use without restrictions, and the model has seen substantial adoption as the foundation for domain-specific fine-tuned variants in areas including SQL generation, financial code, and scientific computing.

2.14 Anthropic Claude: Safety-Aligned Language Models for Code

Claude, developed by Anthropic and first publicly available in March 2023, represents a distinct approach to large language model development rooted in a principled framework for AI safety and alignment [15, 48]. Anthropic was founded in 2021 by former OpenAI researchers including Dario Amodei, Daniela Amodei, and others with the explicit mission of building AI systems that are reliably safe, interpretable, and steerable — properties they considered insufficiently prioritized by the commercial pressures driving the broader LLM industry.

Anthropic's distinctive contribution to training methodology is Constitutional AI (CAI) [48], which supplements and partially replaces human feedback with AI-generated feedback guided by a set of explicit principles — the 'constitution' — governing desired model behavior. The CAI training process involves two phases: supervised learning from AI-generated critiques and revisions of model outputs according to constitutional principles, and reinforcement learning from AI-generated preference comparisons guided by the same principles. This approach enables alignment training at a scale and consistency that purely human-feedback-based approaches cannot match: the AI can generate critiques according to constitutional principles much faster and more consistently than human raters, enabling more thorough coverage of the space of potential model behaviors.

Claude 3.5 Sonnet [15], released in June 2024, achieved leading performance on several coding benchmarks at its release, including 92.0% pass@1 on HumanEval and 49% on SWE-bench Verified — then the highest reported performance on the latter benchmark. The model's 200,000-token context window, the largest available among frontier models at release, enables practical processing of entire codebases in a single context pass, a capability with significant implications for repository-level code understanding and generation tasks. The model's alignment training manifests practically in more consistent acknowledgment of code uncertainty, clearer communication of security implications, and more reliable refusals of requests for clearly malicious code — properties that are difficult to measure on standard benchmarks but highly relevant for enterprise deployment.

Anthropic's research contributions extend substantially beyond model development to foundational work on AI interpretability [16], alignment theory [17], and evaluation of potentially dangerous capabilities. The Responsible Scaling Policy (RSP) introduced by Anthropic established a framework for evaluating AI systems for dangerous capabilities at each scale and deploying additional safeguards before crossing capability thresholds. For software engineering practitioners, these research investments manifest as models that are more reliably aligned with human values and organizational norms, with consistent behavior across the diverse and unpredictable inputs encountered in real development workflows.

2.15 Gemini: Natively Multimodal Models from Google DeepMind

Gemini, introduced by Google DeepMind in December 2023 [17], represents Google's response to GPT-4 and Claude 3 with a natively multimodal architecture trained jointly on text, images, audio, and video from the ground up, rather than adding modalities through post-hoc adapters to a text-trained model. The natively multimodal design enables Gemini to process and reason about interleaved sequences of different modalities in a single forward pass, providing more coherent multi-modal reasoning than adapter-based approaches where separate encoders process different modalities independently before combining their representations.

Gemini was released in three variants — Ultra (most capable), Pro (balanced), and Nano (efficient for on-device deployment) — reflecting the range of deployment scenarios from cloud-scale reasoning to edge inference. Gemini Ultra achieved 90.0% on the MMLU benchmark, becoming the first model to exceed human expert performance (89.8%) on this comprehensive knowledge assessment spanning 57 academic domains. For software engineering specifically, Gemini's multimodal capability enables qualitatively new interactions: a developer can share a screenshot of a UI mockup and ask Gemini to generate the corresponding React component; upload a photo of a whiteboard

architecture diagram and ask for an implementation; or provide a screenshot of an error message alongside code and receive debugging assistance that integrates both inputs.

Gemini 1.5 Pro, released in February 2024, introduced the extraordinary context length of one million tokens — sufficient to encompass approximately 30,000 lines of code, hundreds of pages of technical documentation, or the complete source of many production software projects. This capability, achieved through a novel Mixture-of-Experts attention mechanism that achieves sub-quadratic scaling with context length, enables qualitatively different software engineering tasks: analyzing entire legacy codebases for migration planning, maintaining context through extended multi-session debugging investigations, and generating documentation by processing a complete software project in a single inference. Gemini 1.5's 'needle-in-a-haystack' evaluations, in which a specific piece of information is embedded at varying positions within a million-token context and the model is asked to retrieve it, demonstrated near-perfect retrieval accuracy across the full context window — a non-trivial result given the tendency of long-context models to 'lose' information in the middle of very long contexts.

2.16 GitHub Copilot: LLMs in the Professional Developer Workflow

GitHub Copilot, launched by GitHub and OpenAI in June 2022 following a year-long technical preview [19], represents the first large-scale commercial deployment of LLM-assisted programming and remains the market-leading product in the AI coding assistant category. Copilot's position as the first broadly adopted LLM coding tool makes it the subject of more extensive empirical study than any other tool in the domain, with documented evidence from controlled productivity studies, large-scale usage analyses, and enterprise deployment reports providing a unique empirical foundation for assessing the practical impact of LLM coding assistance.

Copilot's technical architecture integrates the underlying LLM — initially Codex, subsequently GPT-4-based models including tailored fine-tunes for code completion — with IDE plugins that manage context extraction, API communication, suggestion rendering, and user interaction. The context extraction system aggregates information from multiple sources to construct the prompt: the content of the active file (including lines preceding and following the cursor); content of recently opened files providing additional project context; language server information about available completions and type signatures; and a proprietary retrieval system that identifies semantically similar code patterns from the developer's own codebase. This multi-source context construction represents a sophisticated RAG system — retrieval-augmented generation from the developer's own code — that substantially improves the relevance of suggestions for project-specific patterns and conventions.

Copilot's primary user experience is inline 'ghost text' suggestions: as the developer types, Copilot generates and displays one or more code suggestions in a

distinct visual style that can be accepted with a keyboard shortcut or dismissed. The debouncing system avoids excessive API calls during rapid typing by waiting for a brief pause before generating suggestions. Suggestion caching prevents repeated API calls for positions where a suggestion has already been generated. The system adapts its suggestion rate based on acceptance patterns, offering more frequent and longer suggestions in contexts where the developer has recently accepted AI-generated code.

Kalliamvakou's controlled study [50], conducted by GitHub Research with approximately 95 experienced developers working on a standardized HTTP server implementation task, found that Copilot users completed the task an average of 55% faster than a control group without AI assistance. Developer satisfaction surveys found that 60-75% of Copilot users felt more satisfied, less frustrated, and more focused when working with the tool. Ziegler et al.'s large-scale observational analysis [115], examining Copilot usage patterns across thousands of professional developers over an extended period, found sustained suggestion acceptance rates of 26-35%, with acceptance strongly correlated to suggestion specificity, length, and the degree to which the suggestion completed a semantically meaningful code unit (full functions or self-contained blocks were accepted more than partial expressions).

2.17 Cursor, Windsurf, and AI-Native Development Environments

The emergence of AI-native integrated development environments, pioneered by Cursor [52] and followed by Windsurf [53], represents a qualitative shift in the architecture of AI coding assistance from plugin augmentation of existing IDEs to ground-up integration of LLM capabilities within the development environment. This architectural shift enables deeper context integration, more sophisticated multi-file operations, and more seamless AI-developer collaboration than plugin-based approaches can achieve.

Cursor, developed by Anysphere and launched in 2023, is built as a fork of Visual Studio Code, providing the familiar VS Code editing environment and ecosystem compatibility while adding deeply integrated AI capabilities. The critical architectural advantage of Cursor over plugin-based tools is full project context access: rather than receiving only the content of the currently active file and a limited window of recently opened files, Cursor indexes the entire project codebase through vector embeddings, enabling semantic search and context retrieval from any file in the project. This full-project indexing enables Cursor to answer questions about project architecture, generate code consistent with project-wide conventions, and understand how a proposed change affects the broader codebase.

Cursor's Composer feature enables multi-file AI editing: the developer describes a task and Cursor proposes and implements changes across multiple files simultaneously, presenting a unified diff for developer review and acceptance. This capability, absent

from plugin-based tools that operate on single files, enables AI assistance for refactoring tasks that necessarily span multiple files. The Cursor Agent mode extends this to autonomous task execution: Cursor can iteratively edit files, run tests, observe failures, and propose fixes in a loop, approximating agentic behavior within the IDE environment. The model selection flexibility — supporting GPT-4o, Claude 3.5 Sonnet, and other frontier models as the underlying intelligence — enables users to choose the model best suited to their tasks and cost constraints.

Windsurf, developed by Codeium and released in late 2024, introduced the 'flows' paradigm for AI-assisted development: collaborative sessions in which the AI maintains a persistent, evolving model of the developer's goals, recent actions, and codebase state throughout a development session. This temporal awareness addresses a critical limitation of context-stateless tools: by tracking what the developer has been doing — which files were recently modified, what tests were run, what error messages appeared — Windsurf can provide more contextually appropriate assistance without requiring the developer to re-explain context with each interaction. Cascade, Windsurf's autonomous agent, implements multi-step task execution with file operations and command execution, enabling workflows like 'implement this feature, write tests, and fix any test failures' that previously required explicit human intervention at each step.

2.18 Agentic Coding Systems: Claude Code, SWE-agent, and Devin

Agentic coding systems represent the frontier of LLM application to software engineering: tools designed for autonomous execution of complex, multi-step tasks rather than single-turn assistance. The emergence of agentic systems reflects a qualitative shift in the LLM software engineering paradigm, from reactive assistants responding to individual developer queries to active collaborators that can plan, execute, and iterate autonomously within developer-defined task boundaries.

SWE-agent, introduced by Yang et al. at Princeton University in April 2024 [54], demonstrated that GPT-4 equipped with a carefully designed Agent-Computer Interface (ACI) could autonomously resolve issues from real-world GitHub repositories at a practically meaningful success rate. The ACI design — providing structured interfaces for file navigation, code search, editing, and test execution optimized for LLM understanding rather than inherited from human-oriented tools — proved critical to performance: using standard Unix shell commands produced substantially worse performance than the purpose-built ACI. SWE-agent achieved 12.5% resolution on the SWE-bench full dataset, representing the first demonstration that LLM agents could meaningfully address the challenging task of resolving real GitHub issues that require understanding of unfamiliar codebases.

Devin, from Cognition AI, announced in March 2024 as the 'first AI software engineer,' incorporated a persistent memory system storing information across extended

task executions, a long-horizon task planner decomposing complex goals into executable steps, web browsing capability for researching solutions and accessing documentation, and integration with a sandboxed shell and code editor. Cognition's reported performance on SWE-bench attracted significant attention — and controversy when third-party reproducibility attempts produced substantially lower results — highlighting the challenge of consistent evaluation for agentic systems whose performance varies with environmental conditions and agent scaffolding.

Claude Code [49], introduced by Anthropic in 2025, provides an agentic coding assistant operating as a terminal-based tool integrated with development environments through CLI and IDE extensions. Claude Code's design philosophy emphasizes developer control: before taking potentially destructive actions, the agent requests explicit confirmation; all actions are logged transparently; and the agent can be interrupted and redirected at any point. The system leverages Claude 3.5 Sonnet's 200K token context window for codebase-level awareness, enabling autonomous handling of tasks that require understanding how changes in one part of a codebase affect other parts. Claude Code's performance on SWE-bench Verified (49%) placed it at the frontier of documented agentic software engineering performance as of its release.

2.19 OpenHands: Open-Source Platform for Agentic Software Engineering

OpenHands (formerly OpenDevin) [56], developed through a community collaboration involving researchers from multiple universities and the open-source AI community, provides an open-source platform for developing, evaluating, and deploying AI software engineering agents. The platform's open nature distinguishes it fundamentally from proprietary agentic systems: all agent logic, evaluation harnesses, and interfaces are publicly available, enabling transparent comparison, academic study, and community contribution in ways that closed-source commercial agents cannot support.

OpenHands' architecture is organized around clear separation of concerns: agent implementations define planning and action-selection logic; model interfaces define how the agent communicates with underlying language models; tool integrations define the environment capabilities available to the agent; and the evaluation harness enables systematic benchmarking. The CodeActAgent, the primary agent implementation in OpenHands, uses a sandboxed Docker container for code execution providing isolation for potentially risky operations, with careful management of the boundary between the agent's execution environment and the host system.

The platform's modular architecture enables systematic comparison of different agent designs, model choices, and tool integrations on standardized benchmarks. The finding that OpenHands with Claude 3.5 Sonnet achieves 37.2% on SWE-bench Verified

— the highest documented open-source agentic software engineering performance as of mid-2025 — demonstrates the importance of the underlying model capability: the same OpenHands scaffold with less capable models achieves substantially lower performance. This finding has implications for understanding the relative contributions of agent architecture versus model capability to overall agentic system performance.

OpenHands' academic governance structure, prioritizing reproducibility and scientific rigor over commercial deployment speed, has produced methodological contributions to agentic system evaluation alongside the platform itself. The team's analysis of agent failure modes on SWE-bench [95] identified characteristic failure patterns including context management failures, insufficient planning for multi-step tasks, and inability to recover from dead-end solution approaches. These failure mode analyses provide actionable targets for improving agent performance and have guided subsequent architecture improvements.

2.20 Retrieval-Augmented Generation for Code Intelligence

Retrieval-Augmented Generation (RAG) [57], introduced by Lewis et al. at Facebook AI Research and now widely applied across domains, addresses a fundamental limitation of parametric language models: their inability to access information beyond what was encoded in model weights during training. For software engineering, this limitation manifests acutely in the context of proprietary codebases: a general-purpose LLM trained on public code data cannot provide contextually appropriate suggestions for organization-specific APIs, internal coding conventions, proprietary frameworks, or code that postdates the model's training cutoff. RAG addresses this by augmenting generation with dynamic retrieval from an external knowledge base at inference time, providing the model with relevant context without requiring retraining or fine-tuning.

RAG architectures for code typically employ a vector database storing dense embeddings of code artifacts: individual functions or methods, complete files, documentation strings, API specifications, and other relevant knowledge sources. The embedding model — which may be a general-purpose semantic embedding model such as text-embedding-ada-002, or a code-specialized model trained for code semantic similarity — maps each artifact to a fixed-dimensional vector representation. At inference time, the developer's query or code context is embedded using the same model, and a k-nearest-neighbor search retrieves the k most semantically similar artifacts. These retrieved artifacts are included in the LLM's context alongside the original query, grounding the model's response in organization-specific knowledge.

RepoCoder [68] demonstrated the value of iterative retrieval for repository-level code completion: an initial completion is generated without retrieved context, this initial completion is used as a query for additional retrieval (since the completion reveals what information is relevant), and a refined completion is generated with the retrieved context.

This iterative approach, which uses the model's own generation to guide context retrieval, provides better completion quality than single-pass retrieval by capturing information needs that are not apparent from the original query alone. Subsequent work has extended this iterative retrieval paradigm to debugging (retrieving similar bug-fix pairs based on error messages) and test generation (retrieving similar test examples based on the function being tested).

Sourcegraph Cody [58] implements RAG over indexed code repositories using Sourcegraph's code intelligence infrastructure, which provides not only vector-based semantic search but also precise symbol cross-references, call graphs, and definition-reference mappings. By combining vector similarity search with precise structural code intelligence, Cody's retrieval system achieves relevance beyond pure semantic similarity: when a developer asks about a function's callers, the system can retrieve precisely the calling code using structural cross-reference information rather than relying on imprecise vector similarity. GitHub Copilot Enterprise's codebase indexing capability provides similar RAG-based context grounding for enterprise deployments, substantially improving suggestion relevance for large organizational codebases.

2.21 RLHF, Constitutional AI, and Alignment for Software Engineering

Reinforcement Learning from Human Feedback (RLHF) [59, 60] is the training methodology that transformed raw language model capability into practically useful AI assistants aligned with human values and preferences. The importance of alignment for software engineering applications cannot be overstated: a highly capable but unaligned code generation model may generate functionally correct but insecure code, confidently assert incorrect information about API behavior, or generate code that serves as a malicious payload rather than a legitimate implementation. Alignment training addresses these risks by teaching models to behave in accordance with human values even in situations that differ from their training distribution.

The RLHF pipeline as applied to code models consists of three phases. In the supervised fine-tuning (SFT) phase, the pre-trained model is fine-tuned on a dataset of high-quality demonstrations curated by human experts: well-documented, correct, and idiomatically written code examples paired with appropriate natural language explanations. This SFT phase establishes the basic behavioral template that the subsequent reinforcement learning phases refine. In the reward model training phase, pairs of model outputs (code samples, explanations, or responses) are presented to human evaluators who indicate their preferences. These preferences are used to train a reward model that can predict human preference scores for novel model outputs without additional human annotation.

In the RL optimization phase, the language model is optimized using the PPO algorithm to maximize reward model scores while remaining close to the SFT model

through a KL divergence penalty. The KL constraint prevents the model from 'reward hacking' — finding ways to score highly on the reward model without actually producing outputs that humans would prefer — by limiting how far the RL-optimized model departs from the SFT baseline. The result of RLHF training is a model whose outputs are substantially better aligned with human preferences: more likely to acknowledge uncertainty, to provide relevant context and caveats, to refuse harmful requests, and to produce documentation and explanations alongside generated code.

Constitutional AI (CAI) [48] provides an alternative alignment methodology developed by Anthropic that partially replaces human preference annotation with AI-generated critiques and revisions guided by explicit constitutional principles. The process works as follows: the model generates an initial response; the model then critiques its own response according to constitutional principles (for example, 'Does this response provide code that could be used for harmful purposes without appropriate context?'); the model generates a revised response that addresses the critique; this revision is used to create preference pairs for reward model training; and the full RLHF pipeline proceeds using these AI-generated preference labels supplemented by human annotation for high-stakes cases. CAI enables alignment training at scales and consistency levels that purely human-annotation approaches cannot achieve, enabling coverage of a broader range of potential model behaviors.

2.22 MCP, Tool Use, and Multi-Agent Frameworks

The evolution from LLMs as passive text generators to active participants in software development workflows has been enabled by the development of tool use capabilities — the ability of LLMs to invoke external tools by generating structured invocation requests — and by the frameworks that orchestrate multi-tool and multi-agent workflows. This capability expansion transforms LLMs from code generation oracles into autonomous agents capable of reading files, executing code, browsing documentation, and iterating through development cycles.

OpenAI's function calling API [61], introduced in June 2023, established a standard interface for LLM tool invocation that has been widely adopted across the industry. Function calling enables developers to describe available tools (as JSON schemas specifying function names, parameter types, and descriptions) and the model responds with structured JSON invocations rather than prose, enabling reliable programmatic parsing of tool calls. The model's ability to select appropriate tools, specify correct parameter values, and chain multiple tool invocations in sequence — learned from training on examples of tool use — provides the foundation for agentic software engineering systems.

The Model Context Protocol (MCP) [62], introduced by Anthropic in November 2024, represents a significant advance in tool use standardization through a client-server

architecture that separates tool implementation from LLM integration. MCP servers expose tools, resources, and prompts through a defined protocol; MCP clients (LLM applications) discover and invoke these capabilities without requiring tool-specific integration code for each combination of LLM and tool. This standardization enables an ecosystem of interoperable tools — database connectors, code execution environments, documentation retrievers, API clients — that work with any MCP-compatible LLM application. The protocol's adoption by multiple LLM providers and tool developers suggests its potential to become a de facto standard for tool-augmented LLM applications.

Multi-agent frameworks including LangChain [64], AutoGen [65], and CrewAI enable the composition of multiple specialized LLM agents into collaborative systems capable of complex software engineering workflows. The ReAct framework [63] interleaves reasoning steps and action execution, enabling agents to plan their tool use and incorporate observation feedback into subsequent reasoning — a pattern proven more reliable for complex multi-step tasks than either pure reasoning or pure action-execution approaches. ChatDev [66] demonstrated end-to-end autonomous software development through a multi-agent architecture with specialized agents for product management, design, coding, and testing, communicating through simulated software company role structures. While ChatDev's autonomous capabilities remain limited compared to professional developers for complex projects, the demonstration that coordinated multi-agent systems could produce simple functional software from natural language requirements was a significant milestone.

2.23 Benchmark Evaluation: HumanEval, SWE-bench, and Beyond

The development of well-designed, standardized benchmarks for LLM code generation capability has been a critical enabling factor for the field's rapid progress, providing the common evaluation currency that allows researchers to compare models fairly, track progress over time, and identify capability gaps motivating further research. The benchmark landscape has evolved from simple function-level programming challenges to complex multi-file issue resolution tasks that better approximate real-world software engineering complexity.

HumanEval [18], introduced alongside Codex in 2021, established the standard evaluation paradigm for code generation: 164 hand-crafted Python programming problems, each with a function signature, natural language docstring, and hidden test cases; evaluation through functional correctness (pass@k metric); and clear baseline establishment against which subsequent models could be compared. HumanEval's strengths — simplicity, standardized evaluation, and wide adoption enabling longitudinal comparison — have made it the de facto standard for code generation evaluation. Its limitations — small size (164 problems), relative simplicity (self-contained functions,

moderate algorithmic difficulty), known contamination risk (widely present in training data by 2023), and exclusive focus on Python — have motivated the development of complementary and alternative benchmarks.

MBPP (Mostly Basic Python Problems) [92], introduced by Austin et al. at Google, provides approximately 500 Python programming problems sourced from programming forums and educational resources. MBPP complements HumanEval with a different distribution of difficulty and problem styles, providing a more robust combined evaluation than either benchmark alone. MultiPL-E [95] addresses HumanEval's Python exclusivity by translating the benchmark problems to 18 programming languages, enabling cross-lingual capability assessment. Models that achieve high HumanEval scores on Python often show substantial performance degradation on lower-resource languages, revealing the uneven language coverage of code model training data.

SWE-bench [93], introduced by Jimenez et al. at Princeton in 2023, represents a qualitatively more challenging and ecologically valid evaluation than function-level benchmarks. Each SWE-bench task provides a real GitHub repository at the time a specific issue was filed, the text of the issue, and a test suite that passes after the issue is properly resolved. The model (or agent) must modify the repository code to resolve the issue while maintaining compatibility with the existing codebase — a task requiring understanding of multi-file project structure, navigation of unfamiliar code, and generation of targeted patches rather than self-contained functions. SWE-bench's difficulty is reflected in the dramatically lower performance of even frontier models: while GPT-4 achieves 85%+ on HumanEval, its direct application (without agent scaffolding) to SWE-bench resolves only approximately 2% of tasks.

The SWE-bench Verified subset [148], curated by filtering SWE-bench tasks where the ground truth test suite was confirmed to correctly characterize the issue (not all original SWE-bench tasks have reliable test suites), has become the preferred evaluation target for agentic systems. SWE-bench Verified's higher quality ground truth enables more reliable comparison of agent performance. LiveCodeBench [94] addresses benchmark contamination through continuous updates: new competitive programming problems from Codeforces, LeetCode, and AtCoder are added monthly, ensuring that the evaluation set always contains problems that postdate model training cutoffs. This dynamic updating provides contamination-resistant performance estimates that better reflect true generalization capability.

BigCodeBench [96] evaluates complex function-level code generation requiring correct usage of diverse library APIs beyond the standard library. By requiring integration with external libraries — NumPy, Pandas, Requests, BeautifulSoup, and others — BigCodeBench tests the kind of API knowledge application that dominates practical software development but is absent from HumanEval's self-contained algorithmic problems. Models that perform well on HumanEval do not always perform

proportionally well on BigCodeBench, reflecting differences in how well they have learned specific API usage patterns versus general algorithmic reasoning.

2.24 Hallucination in Code Generation: Research and Mitigation

Hallucination — the generation of plausible-sounding but factually incorrect content — represents one of the most practically consequential limitations of LLMs for software engineering applications. Unlike hallucinations in natural language tasks, where detecting errors requires human knowledge or fact-checking against external sources, many hallucinations in code generation are immediately detectable through execution: a hallucinated function that does not exist produces a `NameError` or `AttributeError`; a hallucinated API endpoint returns a 404 or connection error; a hallucinated algorithm that compiles but produces incorrect results is detected through test failures. This executability property makes code hallucination both more tractable and more dangerous than factual hallucination in prose: more tractable because automated testing can detect many instances; more dangerous because undetected hallucinations (those that are syntactically correct and pass automated tests but are semantically wrong) may escape into production code.

Ji et al.'s survey of hallucination in large language models [122] identified multiple contributing factors: the maximum likelihood training objective optimizes for plausible next tokens rather than factual accuracy, potentially rewarding fluent but incorrect outputs; distributional shift between training data and deployment contexts introduces situations where model priors are unreliable; and the absence of explicit factual verification mechanisms in transformer architectures means the model has no internal mechanism to distinguish memorized facts from extrapolated patterns.

Code-specific hallucination categories include: API hallucinations, in which the model generates plausible-sounding but non-existent function names, method signatures, or parameter combinations for real libraries; library version hallucinations, in which the model generates code using APIs from versions that predate or postdate the version in the deployment context; logical hallucinations, in which syntactically correct code implements an algorithm that is subtly wrong — off-by-one errors, incorrect edge case handling, or algorithmic complexity mismatches; and context hallucinations, in which code is generated that would be correct in isolation but is inconsistent with the surrounding codebase context — wrong variable names, incompatible type assumptions, or violations of local conventions.

Mitigation strategies under active research and deployment include: RAG augmentation with authoritative API documentation to ground API usage in verified sources (shown to reduce API hallucination rates by 40-60% in controlled evaluations); execution-based verification that runs generated code against test suites and iteratively refines failing solutions through multi-turn feedback; self-verification prompting that

instructs the model to review its own generated code for potential errors before finalizing; uncertainty quantification approaches that identify when the model's confidence in specific code elements is low; and ensemble approaches that generate multiple candidate solutions and identify hallucinations through disagreement among candidates. Each approach addresses different categories of hallucination with different trade-offs in computational cost and effectiveness.

2.25 Cross-Lingual Code Support and Low-Resource Language Challenges

The distribution of programming language representation in LLM training data is highly skewed, reflecting the distribution of public code on platforms including GitHub, Stack Overflow, and developer tutorials. Python, JavaScript, Java, C++, TypeScript, and Go collectively account for a large majority of public code, while many languages with significant enterprise usage — COBOL, ABAP, RPG, PL/SQL, Fortran, MUMPS — are severely underrepresented. This imbalance creates a two-tier LLM support landscape: high-resource languages receive strong code generation, documentation, and debugging assistance, while low-resource languages — precisely those languages most often associated with critical legacy systems and specialized domains — receive substantially weaker assistance.

The practical implications of this imbalance are significant for enterprise software engineering. Financial institutions maintain billions of lines of COBOL code running critical transaction processing systems; healthcare organizations maintain RPG and ABAP applications on IBM i platforms; scientific computing communities rely on Fortran for high-performance numerical computing. The developers responsible for maintaining these systems are aging, and attracting younger developers to legacy language domains is increasingly difficult. LLM assistance for low-resource languages could substantially reduce the barrier to entry for maintaining legacy systems, but current models' limited training data in these languages produces poor performance precisely where assistance would provide the greatest value.

Research addressing the low-resource programming language challenge includes: cross-lingual transfer learning, which leverages high-resource language capability through shared token representations and cross-lingual pre-training objectives; few-shot fine-tuning on small quantities of target language examples; synthetic data generation through automated translation from high-resource languages; and community-driven data collection initiatives. MultiPL-E's evaluation of models across 18 languages has documented the performance gradient from high-resource to low-resource languages, providing empirical grounding for prioritizing research efforts. Models including Qwen2.5-Coder and StarCoder2 have specifically targeted broader language coverage in

their training data, demonstrating meaningful improvements for previously low-resource languages.

2.26 Economic Impact and Labor Market Effects of LLMs in SE

The economic implications of large language models in software engineering extend well beyond individual developer productivity to encompass labor market structure, compensation dynamics, the distribution of economic value in the software industry, and broader macroeconomic effects of productivity-augmenting technology. These implications are consequential and contested, and the scholarly literature reflects genuine uncertainty about both the magnitude and direction of effects.

At the individual productivity level, the evidence reviewed elsewhere in this thesis consistently demonstrates meaningful improvements in task completion speed for routine coding, documentation, and testing tasks. Translating these task-level improvements to economic value requires several additional steps: estimating the fraction of developer time spent on tasks amenable to LLM assistance, accounting for verification and integration overhead, and modeling how productivity improvements translate to organizational output. McKinsey's 2023 analysis [118] estimated that generative AI could automate or substantially augment 20-45% of software engineering tasks, with the lower bound reflecting current capabilities and the upper bound reflecting anticipated near-term advances.

Labor market effects are more complex and contested than productivity effects. The standard economic framework for analyzing automation's labor market impact — the task-based model of Acemoglu and Restrepo [2019] — distinguishes between substitution effects (automation displacing workers from tasks they previously performed) and complementarity effects (automation increasing the productivity and therefore the demand for workers who perform complementary tasks). For software engineering, substitution effects are most pronounced for routine implementation, boilerplate code generation, and simple documentation tasks; complementarity effects are strongest for senior developers who can effectively direct AI assistants, architectural design, system integration, and quality assurance for AI-generated code. The net labor market effect depends on the relative magnitude of these countervailing forces, which remains empirically uncertain.

Early empirical evidence on hiring trends suggests differentiated effects across experience levels. Industry surveys and job posting analyses report decreased relative demand for junior software engineering roles — particularly those focused on routine implementation in well-understood domains — alongside sustained or increased demand for senior roles requiring architectural expertise, complex problem-solving, and effective AI collaboration. If this pattern persists, it raises concerns about the career pipeline: if AI tools reduce the volume of junior-level work available, the path by which junior

developers gain experience and advance to senior roles may be disrupted, potentially creating long-term supply shortages of senior expertise.

2.27 Synthesis and Identification of Research Gap

The literature surveyed in this chapter covers a period of unusual productivity in applied AI research. From 2017 to 2025, the Transformer provided an enabling architectural foundation; BERT and GPT established the pre-training paradigm; Codex demonstrated practical commercial viability; and successive model generations — from StarCoder and Code Llama through DeepSeek-Coder and Claude 3.5 Sonnet — have pushed the performance frontier to levels that would have seemed implausible a decade earlier. The emergence of agentic systems extending LLM capability from single-turn assistance to multi-step autonomous software engineering has further broadened the trajectory.

Yet the literature also reveals persistent limitations that should temper optimism. Hallucination rates, while declining with scale and improved training, remain practically significant for high-stakes applications. Security vulnerabilities in LLM-generated code are an under-examined risk in the deployment literature. The concentration of strong LLM performance in high-resource programming languages leaves large communities of developers working with legacy systems without adequate AI support. Environmental costs at enterprise-wide adoption scales have received insufficient attention. Governance frameworks covering security scanning, IP compliance, developer oversight, and accountability for AI-generated code remain underdeveloped.

The clearest finding from the survey, however, is a gap in the literature itself: while individual LLM applications in software engineering have attracted extensive scholarly attention, no unified treatment spanning the full development lifecycle — integrating technical foundations, empirical capability evidence, real-world deployment lessons, challenge analysis, and forward-looking research directions — has been produced. Closing that gap is the primary aim of the chapters that follow.

Model/ System	Organizati on	Type	Context	HumanEva l	SWE- bench(V)	Open
GPT-4o	OpenAI	General LLM	128K	90.2%	49%	No
Claude 3.5 Sonnet	Anthropic	General LLM	200K	92.0%	49%	No
Gemini 1.5 Pro	Google DeepMind	General LLM	1M	84.1%	38%	No
DeepSeek- Coder-V2	DeepSeek AI	Code MoE	128K	90.2%	N/A	Yes
Qwen2.5- Coder-72B	Alibaba DAMO	Code LLM	128K	88.4%	N/A	Yes
Code Llama 70B	Meta AI	Code LLM	100K	53.0%	N/A	Yes
StarCoder2- 15B	BigCode	Code LLM	16K	46.1%	N/A	Yes
Codex	OpenAI	Code LLM	8K	28.8%	N/A	No

(davinci)						
GitHub Copilot	GitHub/ Microsoft	IDE Tool	Large	N/A	N/A	No
Claude Code	Anthropic	Agent	200K	~92%	49%	No
OpenHands+ Claude	Community	Agent Plat.	200K	N/A	37.2%	Yes
Devin	Cognition AI	Agent	Large	N/A	13.9%	No

**Figure 2: Tool/Model Type Distribution
(Comparative Study, n = 11)**




Figure 2: Distribution of Tool/Model Types in Comparative Study




Figure 1: HumanEval Pass@1 Scores by Model

Table 2.1: Model and system comparison as of mid-2025. SWE-bench (V) = Verified subset. All values approximate.

CHAPTER 3: FOUNDATIONS OF LARGE LANGUAGE MODELS

The sections that follow provide the technical foundations necessary for understanding large language model capabilities and limitations as they apply to software engineering. The treatment is designed to be accessible to readers with a background in computer science but not necessarily in deep learning research. Mathematical notation is introduced where it clarifies rather than obscures, with verbal explanation accompanying all formal definitions. Readers seeking greater depth are directed to the primary sources cited throughout.

3.1 Tokenization and Vocabulary Construction

Tokenization is the process of converting raw text — including source code — into discrete units called tokens that serve as the fundamental inputs and outputs of language models. The choice of tokenization scheme has significant consequences for model capability, computational efficiency, and behavior on domain-specific inputs such as programming languages. Three principal tokenization approaches have been developed and applied: character-level, word-level, and subword tokenization.

Character-level tokenization uses individual characters as tokens, providing a minimal vocabulary (typically 100-300 characters for Unicode subsets) and the ability to handle arbitrary strings without out-of-vocabulary failures. However, character-level tokenization produces very long token sequences — a 1,000-character code snippet becomes 1,000 tokens — making it computationally expensive for models with quadratic attention complexity and challenging for models to learn long-range dependencies across many characters. Character-level models have demonstrated surprising capability on code tasks [21] but are generally outperformed at equivalent compute budgets by subword models that operate at a more abstract token granularity.

Word-level tokenization uses entire words or tokens delimited by whitespace and punctuation as the basic units, producing naturally interpretable tokens but struggling with out-of-vocabulary words — novel identifiers, domain-specific terms, or code-specific tokens not present in the training vocabulary. The fixed vocabulary size of word-level models (typically 50,000-100,000 words) cannot accommodate the open-ended identifier namespace of programming languages, where developers coin new names freely.

A function named `calculateRevenueAttributedToCustomerSegmentAfterTaxDeduction` would require either a vocabulary entry for this specific identifier (almost certainly absent) or an out-of-vocabulary token that loses all information about the identifier's meaning.

Byte-Pair Encoding (BPE), introduced by Sennrich et al. [67] for neural machine translation and adopted by the GPT series, provides the dominant tokenization approach for contemporary LLMs by achieving an effective balance between vocabulary size and sequence length. BPE begins with individual characters (or bytes) as the initial vocabulary and iteratively merges the most frequent consecutive pair of tokens, adding the merged pair to the vocabulary. This bottom-up process continues until the vocabulary reaches a target size (typically 32,000-100,000 tokens). The resulting vocabulary contains common words as single tokens, common multi-word sequences as tokens, rare words decomposed into common subword units, and individual characters as the ultimate fallback for novel strings.

BPE Tokenization for Source Code

For source code, BPE tokenization has specific and important properties. Programming language keywords (`def`, `class`, `import`, `return`, `for`, `while`) appear frequently in Python code and are almost always encoded as single tokens. Common identifiers (`self`, `print`, `len`, `range`, `list`, `dict`) similarly become single tokens. Less common identifiers are decomposed into meaningful subword components: `calculateRevenue` might tokenize as `[calculate, Revenue]` or `[calc, ulate, Rev, enue]` depending on frequency in training data. This decomposition is semantically meaningful to some degree — model attention can leverage the subword structure to understand relationships between identifiers that share common subwords — but also introduces representation challenges when rare identifiers are fragmented into many pieces.

The tokenization efficiency of different programming languages varies substantially, with important consequences for model performance. Python code, due to its clean syntax and heavy representation in training data, achieves high tokenization efficiency: a typical Python function requires relatively few tokens per line of code. Languages with verbose syntax (Java's boilerplate), unusual identifiers (COBOL's data division names), or special characters (APL's mathematical symbols) require proportionally more tokens per semantic unit, increasing sequence lengths and therefore computational cost. This tokenization efficiency difference is one factor contributing to the observed performance advantage of LLMs on Python relative to less-represented languages.

SentencePiece [69], developed at Google, provides an alternative subword tokenization approach that operates directly on Unicode characters rather than requiring pre-tokenization into words. This language-agnostic approach is particularly well-suited to multilingual models and to code models that must handle diverse programming languages with different identifier characters and syntactic conventions. The T5 and many subsequent models use SentencePiece, while the GPT series uses Tiktoken's implementation of BPE. For software engineering applications, the practical differences

between these approaches are minor compared to the effects of training data composition and model scale.

3.2 Embeddings, Representations, and Positional Encoding

Tokens produced by the tokenizer are discrete indices into the vocabulary — integers with no inherent geometric structure. Before being processed by transformer layers, tokens must be converted to continuous vector representations that can participate in the linear algebraic operations underlying attention and feed-forward computation. This conversion is accomplished by an embedding table: a learned matrix E of dimensions $|V| \times d_{\text{model}}$, where $|V|$ is the vocabulary size and d_{model} is the model's embedding dimension (typically 768 to 8,192 for models of various scales). The embedding of token v is simply the row $E[v]$ — a d_{model} -dimensional vector learned during training to represent the token's properties.

Well-trained embedding matrices exhibit rich geometric structure: semantically similar tokens cluster in nearby regions of the embedding space, and semantic relationships are reflected in vector arithmetic. For code, this manifests as clustering of semantically related tokens: control flow keywords (`if`, `else`, `elif`, `while`, `for`) cluster together; built-in functions (`print`, `len`, `type`, `range`) cluster together; error types (`ValueError`, `TypeError`, `AttributeError`) cluster together. The analogy-solving property of word embeddings — that $\text{vec}(\text{king}) - \text{vec}(\text{man}) + \text{vec}(\text{woman}) \approx \text{vec}(\text{queen})$ — has analogues in code embeddings: relationships between language features can sometimes be captured in embedding geometry [70].

Because transformer self-attention is a set operation with no inherent sensitivity to token order — two sequences with the same tokens in different orders would produce identical attention weights without additional information — explicit positional information must be incorporated. The original Transformer added fixed sinusoidal positional encodings to token embeddings: position pos and dimension i received the encoding $\sin(\text{pos}/10000^{(2i/d)})$ for even dimensions and $\cos(\text{pos}/10000^{(2i/d)})$ for odd dimensions. This fixed encoding provides a unique representation for each position and enables the model to learn relative distance relationships from the difference of sinusoidal encodings at different positions.

Modern LLMs have largely replaced additive sinusoidal encodings with Rotary Position Embedding (RoPE) [71], which encodes position as a rotation applied to the query and key vectors before computing attention. For a position p , the rotation applied to a query vector q transforms it to $q_{\text{rotated}} = R_p q$, where R_p is a position-dependent rotation matrix defined by block-diagonal structure on 2-dimensional subspaces: each pair of dimensions (q_{2k}, q_{2k+1}) is rotated by angle $p\theta_k$ where θ_k is a geometrically decaying frequency. The key insight of RoPE is that the dot product of a rotated query at position p with a rotated key at position m depends only on the relative

position $m-p$ through the algebraic identity $R_p^T R_m = R_{\{m-p\}}$, making RoPE naturally implement relative position attention without requiring explicit relative position computation.

Long-Context Position Encodings

RoPE's generalisation to positions beyond those seen during training is limited: rotational frequencies calibrated during training on 4K-token sequences may provide poorly calibrated attention at 100K-token inference positions. YaRN (Yet another RoPE extension) [73] addresses this through frequency interpolation that adjusts rotational frequencies to scale smoothly from training to extended context lengths, enabling models trained on 4K contexts to function coherently at 128K tokens with appropriate fine-tuning. ALiBi (Attention with Linear Biases) [72] provides an alternative approach that adds position-dependent linear penalties to attention logits before softmax — discouraging attention to distant tokens — which generalizes to positions far exceeding training lengths without frequency interpolation requirements.

Method	Description	Context Generalization	Used By
Sinusoidal	Fixed additive encodings	Limited beyond training	Original Transformer
Learned Absolute	Trained position embeddings	Training length only	BERT, GPT-2
RoPE	Rotational query/key modification	Good with scaling	LLaMA, GPT-NeoX
ALiBi	Linear attention biases	Strong beyond training	StarCoder, BLOOM
YaRN	RoPE frequency interpolation	Excellent with finetuning	Mistral, extended LLaMA

Table 3.1: Positional encoding methods and their context generalization properties.

3.3 Self-Attention: The Core Computational Primitive

Self-attention is the mechanism that distinguishes Transformer architectures from all prior sequence modeling approaches and provides the fundamental capability for modeling arbitrary long-range dependencies in code. Unlike recurrent architectures that propagate information through time steps and convolutional architectures that aggregate local neighborhoods, self-attention directly computes pairwise relationships between all positions in the input sequence, enabling each position to draw information from any other position in a single computational step.

Formally, consider an input sequence represented as a matrix $X \in \mathbb{R}^{n \times d_{\text{model}}}$, where n is the sequence length and d_{model} is the embedding dimension. Three learned linear projections transform X into query, key, and value matrices:

$$Q = XW_Q, \quad K = XW_K, \quad V = XW_V$$

where $W_Q, W_K, W_V \in \mathbb{R}^{d_{\text{model}} \times d_k}$ are learned projection matrices. The scaled dot-product attention then computes:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

The term $QK^T \in \mathbb{R}^{n \times n}$ computes the dot product similarity between each query-key pair, producing a raw score matrix where entry (i,j) reflects how relevant position j is to position i . Division by $\sqrt{d_k}$ prevents excessively large dot products for high-dimensional keys, which would push the softmax into saturation regions with near-zero gradients. The softmax normalizes each row into a probability distribution over positions, and the resulting normalized scores are used as a weighted average of the value matrix rows, producing output $O \in \mathbb{R}^{n \times d_k}$ where each row is a weighted combination of value vectors with weights reflecting attention patterns.

For causal (decoder) attention, an additive mask is applied before softmax: entries in the upper triangle of the $n \times n$ matrix (corresponding to position $j > i$ for query at position i) are set to $-\infty$, causing them to receive zero attention weight after softmax. This mask enforces the autoregressive constraint that generation at each position may only depend on preceding positions, enabling the Transformer decoder to be trained with full sequence parallelism despite generating autoregressively at inference time.

Multi-Head Attention and Head Specialization

Multi-head attention computes h parallel scaled dot-product attention functions with reduced dimensionality $d_k = d_{\text{model}} / h$, then concatenates and projects the results:

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h) W_O$$

where $\text{head}_i = \text{Attention}(QW_Q^i, KW_K^i, VW_V^i)$ and $W_O \in \mathbb{R}^{d_{\text{model}} \times d_{\text{model}}}$ is the output projection. The reduced per-head dimensionality (typically 64 or 128 dimensions per head) allows each head to specialize in capturing different relationship types without the total computation exceeding that of a single full-dimension attention.

Analysis of attention patterns in code-trained models has revealed consistent semantic specialization across heads [75]. Certain heads reliably attend to matching delimiters: an open parenthesis token attends strongly to its corresponding close

parenthesis; an open bracket attends to its close bracket. Other heads capture identifier cross-references: uses of a variable attend to its definition site. Still other heads capture syntactic relationships: function call tokens attend to the function definition tokens; argument tokens attend to corresponding parameter tokens in the function signature. This implicit learning of code structural relationships from the language modeling objective, without explicit supervision, demonstrates the power of attention as a mechanism for learning structured code representations.

Efficient Attention Variants

Multi-Query Attention (MQA) [138] reduces memory requirements at inference time by using a single set of key and value matrices shared across all query heads: only query projections are head-specific. Grouped Query Attention (GQA) [76] provides a middle ground, grouping heads into G groups that share key-value projections, with G tunable between 1 (MQA) and h (full MHA). GQA achieves most of the memory savings of MQA while retaining more representational capacity, and has been adopted in recent frontier models including LLaMA 2 and Claude 3.

The computational complexity of standard attention — $O(n^2 \cdot d)$ in both time and memory — represents the primary scaling constraint for context length. FlashAttention [74] achieves the same mathematical result as standard attention with dramatically reduced memory usage by restructuring computation to exploit GPU memory hierarchy: rather than materializing the complete $n \times n$ attention weight matrix in GPU HBM (slow, large memory), FlashAttention tiles the computation to keep intermediate results in GPU SRAM (fast, small memory), reducing HBM reads and writes by a factor proportional to block size. The practical consequence is that FlashAttention enables training and inference with sequences 3-4x longer than standard attention on the same hardware, and has been essential to the practical deployment of long-context models.

3.4 Feed-Forward Networks, Normalization, and Residual Connections

Each transformer layer interleaves multi-head attention with a position-wise feed-forward network (FFN) that applies the same two-layer MLP to each sequence position independently. The canonical FFN with ReLU activation is:

$$\text{FFN}_{\text{ReLU}}(x) = \max(0, xW_1 + b_1) W_2 + b_2$$

Modern LLMs replace ReLU with SwiGLU activation [78], which introduces a gating mechanism providing better gradient flow and empirically superior performance:

$$\text{FFN}_{\text{SwiGLU}}(x, W, V, W_2) = (\text{Swish}(xW) \odot (xV)) W_2$$

where $\text{Swish}(z) = z \cdot \text{sigmoid}(z)$ and $\odot$ denotes element-wise multiplication. The gating mechanism allows the network to selectively suppress or amplify information from one pathway based on another, providing more expressive nonlinearity than pure activation functions. The FFN dimension is typically $4 \times$ the model dimension (sometimes $8/3 \times d_{\text{model}}$ for SwiGLU to maintain parameter count equivalence), providing a wide expansion-compression bottleneck that has been associated with the storage and retrieval of factual knowledge from pre-training [195].

Layer Normalization [79] normalizes the activations within each layer to have zero mean and unit variance along the embedding dimension, preventing activation magnitudes from growing too large or too small during forward passes and stabilizing gradient magnitudes during backpropagation. The normalized activation is rescaled and shifted by learned parameters γ and β : $\text{LayerNorm}(x) = \gamma \cdot (x - \mu) / \sigma + \beta$, where μ and σ are the mean and standard deviation computed across the d_{model} dimensions for each token independently.

Modern models apply Pre-Layer Normalization (Pre-LN) — applying layer normalization before each sub-layer rather than after — which improves gradient flow in very deep networks and enables training without learning rate warmup [196]. The Pre-LN transformer layer can be written as:

$$x' = x + \text{MultiHead}(\text{LayerNorm}(x))$$

$$\text{output} = x' + \text{FFN}(\text{LayerNorm}(x'))$$

Residual connections [80] — the additive skip connections that add each sub-layer's input to its output — are not merely an engineering convenience but provide a fundamental guarantee for deep network trainability. Without residual connections, gradients must flow through every layer during backpropagation, and vanishing gradients prevent effective learning of early layers. With residual connections, gradients can flow directly from the loss through the skip connections to any layer, maintaining gradient magnitude regardless of depth. For LLMs with 32-100+ layers, residual connections are essential for stable training.

3.5 Decoder-Only vs. Encoder-Decoder vs. Encoder-Only Architectures

The three principal Transformer architecture variants — decoder-only, encoder-decoder, and encoder-only — reflect different trade-offs in computational structure and inductive bias that have significant implications for their suitability to different software engineering tasks. Understanding these architectural differences is essential for reasoning about why particular models excel at particular tasks.

Decoder-only (or causal) Transformer architectures, used by the GPT series, LLaMA, Claude, Gemini, DeepSeek, and most frontier LLMs, apply causal self-attention at each layer: each position may only attend to preceding positions. The pre-training objective — predicting the next token from all preceding tokens — is naturally aligned with this causal structure. Decoder-only models are the dominant choice for large-scale code generation because: (1) their pre-training objective directly trains the skill of continuing code sequences, which is precisely the IDE code completion task; (2) any generation task can be formulated as next-token prediction by appropriately formatting the input and expected output in the context; and (3) they scale more cleanly than encoder-decoder models because there is no need to manage the architectural difference between encoder and decoder components at scale.

Encoder-decoder (or sequence-to-sequence) Transformer architectures, used by T5, CodeT5, CodeT5+, and BART-based code models, apply bidirectional self-attention in the encoder and causal self-attention in the decoder, with cross-attention in the decoder attending to encoder representations. This architecture is particularly natural for tasks where the input and output are clearly distinct sequences: code translation (source language to target language), code repair (buggy code to fixed code), and code summarization (code to natural language description). The encoder's bidirectional context allows it to develop richer representations of the input than a unidirectional decoder can produce, potentially providing better input understanding for translation tasks. However, encoder-decoder models have been largely superseded by large decoder-only models in practice because the latter's ability to handle input-output separation through prompt formatting, combined with their scale advantage, outweighs the architectural benefit of bidirectional encoding for most tasks.

Encoder-only Transformer architectures, used by BERT, CodeBERT, GraphCodeBERT, and related models, apply bidirectional self-attention at each layer, enabling each token to attend to all others without causal masking. This provides maximum representational richness for understanding tasks but prevents autoregressive generation. Encoder-only models excel at code search, code clone detection, vulnerability classification, and other tasks where a discriminative prediction must be made about the full input sequence. Their pre-training objectives — masked language modeling, replaced token detection, or contrastive objectives — develop rich bidirectional representations without training generative capability.

3.6 Mixture of Experts Architectures

Mixture of Experts (MoE) architectures address a fundamental tension in large language model design: larger models achieve better performance, but larger dense models require proportionally more computation per forward pass, making inference increasingly expensive. MoE architectures decouple model capacity (total parameters) from computation (active parameters per forward pass) by replacing the dense FFN in each transformer layer with a set of expert FFNs, routing each token to only a small subset of experts during each forward pass.

The MoE layer for a token representation x consists of N expert FFNs $E_1, \dots, E_N$ and a gating network G . The gating network computes a sparse distribution over experts: $G(x) = \text{Softmax}(\text{TopK}(x \cdot W_g, k))$ where W_g is the gating projection matrix and TopK selects the k largest gate values (typically $k=2$). The MoE output is the weighted sum of the outputs of the top- k experts: $\text{MoE}(x) = \sum_{i \in \text{TopK}} G_i(x) \cdot E_i(x)$. For each token, only k of N experts perform computation; the remaining $N-k$ experts are inactive. This sparse routing achieves model capacity proportional to N (through the total expert parameters) while requiring computation proportional to k (through the active experts), a favorable trade-off for model scale.

DeepSeek-Coder-V2's MoE architecture [46] illustrates the capability of modern MoE models: with 236 billion total parameters but only 21 billion active parameters per token, the model achieves frontier code generation performance while requiring computation equivalent to a 21B dense model. Mixtral 8x7B [82] achieves competitive performance with much larger dense models using only 12.9B active parameters (out of 46.7B total), demonstrating the practical efficiency advantages of MoE at smaller scales. For software engineering deployments, MoE models offer an appealing efficiency proposition: frontier-quality code generation at dense-model inference costs.

Load Balancing and Expert Utilization

A key challenge in MoE training is preventing expert collapse: the gating network converging to routing all tokens to the same small subset of experts, defeating the purpose of the architecture. An auxiliary load-balancing loss penalizes imbalanced expert utilization: $L_{\text{balance}} = \alpha \cdot N \cdot \sum_i f_i \cdot P_i$, where f_i is the fraction of tokens routed to expert i and P_i is the average gate probability for expert i . This loss encourages uniform token distribution across experts. Expert Dead Zone prevention techniques, including random re-initialization of underutilized experts and diverse initialization of expert parameters, further promote robust expert specialization. In code-trained MoE models, different experts may implicitly specialize in different aspects of code: some experts may focus on Python-specific patterns, others on algorithmic logic, others on string manipulation or API usage.

3.7 Training Pipelines, Scaling Laws, and Distributed Optimization

Training large language models requires distributed optimization across hundreds to thousands of accelerator devices (GPUs or TPUs), sophisticated parallelization strategies, and careful management of numerical precision, gradient accumulation, and learning rate scheduling. The engineering of LLM training pipelines is a discipline in itself, with substantial expertise concentrated at the handful of organizations capable of training frontier models.

Three fundamental parallelism strategies are combined in large model training. Data parallelism replicates the complete model on multiple devices and distributes different mini-batches to each replica, aggregating gradients across replicas through all-reduce operations before parameter updates. Data parallelism scales to as many devices as the global batch size permits, but requires each device to store a full model copy — infeasible for models with hundreds of billions of parameters. Tensor parallelism (model parallelism) partitions the weight matrices of individual layers across devices: in the simplest case, the weight matrix W of a linear layer is split by columns across devices, which each compute a portion of the output. Tensor parallelism requires high-bandwidth communication between devices at each layer — typically requiring NVLink or InfiniBand — but reduces per-device memory proportionally to the parallelism degree. Pipeline parallelism assigns different layers to different pipeline stages, enabling devices to work on different micro-batches at different pipeline stages simultaneously, with micro-batch scheduling hiding the pipeline bubble overhead.

The combination of all three strategies — 3D parallelism — is implemented in frameworks including Megatron-LM [85] (developed by NVIDIA) and DeepSpeed [84] (developed by Microsoft). A typical large model training configuration might use tensor parallelism of degree 8 within each node (exploiting NVLink), data parallelism across nodes (using InfiniBand), and pipeline parallelism across node groups, combining these to train models with hundreds of billions of parameters on thousands of GPUs. The engineering sophistication required for these configurations — managing the interplay of communication patterns, gradient synchronization, checkpoint management, and fault recovery — represents a significant competitive advantage for organizations that have invested in it.

Scaling Laws and Training Compute Optimization

Scaling laws describe the quantitative relationship between model performance (typically measured as loss on a held-out test set) and the computational resources invested in training. Kaplan et al. [21] established that model performance follows power-law relationships with model size (N parameters), dataset size (D tokens), and compute ($C = 6ND$ FLOPs for standard training). The Kaplan scaling laws were

instrumental in motivating GPT-3's 175B parameter scale and establishing the 'bigger is better' paradigm that drove LLM development in 2020-2022.

Hoffmann et al.'s Chinchilla scaling laws [88], published by Google DeepMind in 2022, revised the Kaplan laws by demonstrating that the optimal compute allocation balances model size and training data in approximately equal proportions: for a fixed compute budget C , the optimal model size N^* and dataset size D^* satisfy $N^* \approx D^* \approx (C/12)^{0.5}$. The Chinchilla analysis showed that GPT-3 and many contemporaneous models were 'undertrained' relative to their size — they could achieve better performance at lower compute by training smaller models on more data. The 70B Chinchilla model, trained on 1.4 trillion tokens, outperformed the 280B Gopher model on most tasks despite using the same total compute.

For code-specific models, scaling law behavior may differ from general text: code generation capabilities exhibit sharper emergence thresholds, with models below certain scale thresholds generating syntactically correct but semantically wrong code while models above the threshold generate functionally correct solutions. This non-smooth emergence motivates training at scales where emergent code reasoning capabilities are reliably present rather than attempting to economize at sub-threshold scales [90].

3.8 Inference Optimization Techniques

The deployment of large language models for software engineering applications requires careful attention to inference efficiency: the computational cost, latency, and memory requirements of generating code completions, answering developer queries, and executing agentic workflows in production settings. Without optimization, frontier models with hundreds of billions of parameters would require multiple high-end GPUs per user request — economically infeasible for wide-scale deployment. A rich toolkit of inference optimization techniques addresses this challenge.

Key-Value (KV) caching stores the key and value tensors computed for previous tokens, avoiding redundant computation during autoregressive generation. Without KV caching, generating each new token would require computing attention over the complete sequence from scratch, with $O(n)$ computation per step accumulating to $O(n^2)$ total computation for a sequence of length n . With KV caching, only the new token's queries need to be computed at each step, with keys and values for previous tokens retrieved from cache, reducing per-step computation from $O(n \cdot d)$ to $O(d)$ for attention computation. KV caching trades memory for compute, requiring storage proportional to sequence length, model depth, and number of heads — a significant memory cost for long sequences and large models that motivates KV cache compression techniques.

Quantization reduces model weight and activation precision from the 32-bit floating point of training to 16-bit (FP16/BF16), 8-bit (INT8), or 4-bit (INT4)

representations, reducing memory requirements by 2x, 4x, or 8x respectively. The key challenge of quantization is preserving model accuracy under reduced numerical precision: aggressive quantization introduces quantization noise that can significantly degrade model capability, particularly for outlier activation values that are difficult to represent in low-precision formats. LLM.int8() [197] addresses outlier activations through mixed-precision decomposition, computing attention for a small fraction of high-magnitude activations in FP16 while quantizing the remainder to INT8. GPTQ [128] applies layer-wise quantization with optimal weight adjustment to minimize quantization error. AWQ (Activation-aware Weight Quantization) [198] identifies the small fraction of weights critical for model capability through activation magnitude analysis and protects them from aggressive quantization, achieving strong 4-bit performance.

Speculative Decoding and Batched Inference

Speculative decoding [127, 141] accelerates autoregressive generation by using a smaller, faster 'draft' model to propose multiple tokens simultaneously, then using the larger target model to verify the proposals in parallel. If the target model accepts the draft tokens (which it will do when the draft model's predictions match what the target would generate), multiple tokens are accepted in a single target model forward pass. Since target model forward passes can process multiple positions in parallel, speculative decoding achieves effective throughput improvements of 2-4x for tasks where the draft model's accuracy is high — a condition typically satisfied for code completion where consecutive tokens are often predictable from context.

Batched inference groups multiple requests together for simultaneous processing, amortizing the fixed overhead of model loading and maximizing GPU utilization. Continuous batching [202] dynamically adds and removes requests from the batch as they arrive and complete, avoiding the underutilization of static batching approaches where all requests must finish before the batch is released. PagedAttention [203], implemented in the vLLM library, manages KV cache memory using a paging scheme analogous to operating system virtual memory, eliminating the memory fragmentation that wastes capacity in naive KV cache implementations and enabling much higher throughput for concurrent request batching.

3.9 Alignment: RLHF, Constitutional AI, DPO, and Preference Optimization

Alignment — the challenge of ensuring that AI system behavior reflects the values, intentions, and interests of its users and society — is a central concern for LLM deployment in software engineering. Misaligned models may generate code with hidden security vulnerabilities; confidently assert incorrect information about API behavior; generate malicious code when prompted with adversarial inputs; or behave unpredictably

in deployment contexts that differ from training. The alignment research community has developed several methodologies for training models whose behavior reliably reflects human values.

Reinforcement Learning from Human Feedback (RLHF) [59, 60] provides the most widely deployed alignment methodology. The three-phase process begins with supervised fine-tuning on demonstrations of desired behavior from human experts — for code, this means examples of well-documented, correct, secure, and idiomatically written code paired with appropriate explanations. The SFT phase establishes the behavioral template that subsequent phases refine toward alignment with human preferences. The reward model training phase presents pairs of model outputs to human evaluators — for code, these might be two candidate implementations of the same function, two debugging analyses of the same bug, or two explanations of the same code — and trains a reward model on the resulting preference labels to predict human preference for novel outputs. The RL optimization phase uses Proximal Policy Optimization (PPO) to maximize expected reward while maintaining proximity to the SFT model through a KL divergence regularizer.

The practical manifestations of RLHF alignment in code-specific contexts include: more consistent acknowledgment of uncertainty ('This implementation may have edge cases I haven't considered'); more reliable identification and communication of security implications ('Note: this uses MD5 which is cryptographically weak for security applications'); clearer communication of when generated code requires testing before deployment; and more appropriate refusals for requests that appear to target malicious use. These behaviors, absent from base pre-trained models, substantially improve the practical safety of LLM-assisted software engineering.

Direct Preference Optimization (DPO) [91] provides a simpler alternative to RLHF that directly optimizes the language model as a reward function, avoiding the need for a separate reward model and PPO training. DPO reparameterizes the RLHF objective in terms of the policy model's log-probabilities of preferred and dispreferred outputs, deriving a closed-form loss function that can be optimized with standard supervised learning techniques. Despite its simplicity, DPO has demonstrated comparable alignment performance to RLHF on several benchmarks and has been adopted by multiple model developers as a more computationally efficient alignment technique.

Constitutional AI (CAI) [48], developed by Anthropic, provides an alignment methodology that partially replaces human preference annotation with AI-generated critiques guided by explicit constitutional principles. The process trains the model to critique its own outputs according to stated principles (helpfulness, harmlessness, honesty), revise outputs that violate these principles, and use AI-generated preference comparisons as training signal alongside human feedback. CAI enables alignment training at scales where comprehensive human annotation is impractical and produces

models with more consistent behavior under the stated constitutional principles than models trained only with human feedback, which may reflect human annotator inconsistencies.

3.10 In-Context Learning, Few-Shot Prompting, and Chain of Thought

In-context learning (ICL) — the ability of large language models to perform tasks from examples provided in the input context without any gradient updates to model weights — is among the most practically consequential capabilities of frontier LLMs. The demonstration that GPT-3 could perform few-shot NLP tasks from in-context examples established ICL as a first-class capability, and subsequent work has revealed ICL to be a general mechanism through which LLMs adapt their output distribution to novel tasks at inference time.

From a mechanical perspective, ICL works by conditioning the autoregressive generation on a context that establishes the task structure: few-shot ICL provides input-output examples that demonstrate the mapping the model should apply; zero-shot ICL provides a natural language description of the task without examples. The model's prior knowledge — encoded in its weights through pre-training — enables it to generalize from the in-context examples to the test input, without requiring gradient updates. The theoretical understanding of why ICL works remains an active research area: proposed explanations include ICL as implicit gradient descent [204], ICL as Bayesian inference [205], and ICL as retrieval from the training distribution [206].

For software engineering, ICL provides several practical capabilities. Few-shot code generation provides examples of the desired code style, documentation format, error handling approach, and level of abstraction before the generation request, substantially improving the relevance of generated code for project-specific contexts. Few-shot debugging provides examples of error messages and their resolutions before presenting the target error, leveraging the model's pattern-matching capability to identify similar bugs. Few-shot code review provides examples of desired review comments and their format before requesting a review, calibrating the model's output to organizational norms.

Chain-of-Thought Prompting

Chain-of-Thought (CoT) prompting [24] extends few-shot ICL by instructing the model to generate intermediate reasoning steps before producing the final answer. For code generation tasks, CoT prompting involves asking the model to first describe the algorithm, data structures, and implementation approach it will use, then generate the implementation. This intermediate reasoning step improves performance on complex algorithmic tasks by decomposing the generation challenge: reasoning about the algorithm separately from expressing it in code syntax reduces the cognitive (and computational) load of each step.

Empirical studies find that CoT prompting improves pass@1 rates on HumanEval by 10-25% for frontier models, with larger improvements on more complex tasks. Zero-shot CoT prompting — simply appending 'Let's think step by step' to the prompt — also provides performance improvements, suggesting that the prompting cue itself activates reasoning patterns regardless of provided examples [207]. The Self-Consistency technique [208] extends CoT by sampling multiple reasoning chains and selecting the most consistent answer by majority vote — for code generation, this translates to generating multiple candidate solutions and selecting those that pass most test cases, substantially improving performance at the cost of increased inference compute.

3.11 Long-Context Architectures and Repository-Level Understanding

The context length available to an LLM determines the scope of code it can consider simultaneously: a 4K-token context can accommodate a few functions; a 32K-token context can accommodate most single files; a 200K-token context can accommodate many complete codebases. The dramatic expansion of context lengths from the 4K tokens of early GPT models to the 1M tokens of Gemini 1.5 Pro has enabled qualitatively different software engineering capabilities — repository-level code understanding, complete test suite generation, and full-codebase refactoring — that are simply not possible within shorter contexts.

Extending context length beyond training length requires careful handling of positional encodings: position encodings calibrated for 4K contexts may provide unreliable relative position information at 100K or beyond, degrading the model's ability to attend appropriately to distant context. YaRN [73], as described in Section 3.2, addresses this through frequency interpolation. Phi-LongRoPE [209] introduces non-uniform frequency interpolation that better preserves short-context capabilities during context extension. LongLoRA [210] enables efficient long-context fine-tuning through shifted sparse attention that reduces the quadratic complexity of full attention during training while maintaining full-context capability through the residual connections.

Long context models face the additional practical challenge of the 'lost in the middle' phenomenon [74]: models reliably retrieve information near the beginning or end of long contexts but perform substantially worse at retrieving information from the middle of the context. For software engineering applications, this has practical implications: relevant code placed in the middle of a very long codebase context may receive less model attention than code near the context boundaries. Addressing this limitation through attention architecture modifications, training curriculum design, and retrieval augmentation that places relevant code near the context boundary are active research directions.

3.12 Emergent Abilities and Code Reasoning Capabilities

Emergent abilities of large language models — capabilities that appear abruptly and unpredictably at sufficient scale, rather than improving smoothly with model size — have been documented across a range of reasoning and knowledge tasks and are particularly relevant to understanding the non-linear capability improvements observed in code generation [90]. Wei et al.'s analysis of emergent abilities found that many capabilities — multi-step arithmetic reasoning, analogical reasoning, multi-lingual question answering — were absent in models below a threshold scale and present above it, with the transition occurring more sharply than would be predicted by extrapolating pre-threshold performance improvement rates.

For code generation specifically, several emergent capabilities have been documented. Algorithm synthesis — the ability to generate correct implementations of non-trivial algorithms from natural language descriptions — appears to emerge around the 10-30B parameter scale for models trained on code. Multi-file code generation — the ability to generate consistent code across multiple interdependent files — appears to emerge with large context windows combined with sufficient scale. Autonomous debugging — the ability to iteratively identify and fix bugs through execution-observation-hypothesis cycles — appears to emerge only in frontier models combined with appropriate tool use scaffolding. These threshold behaviors have practical implications for tool development: capabilities present at 7B parameters and those present only at 70B+ parameters should inform architecture and deployment choices differently.

The theoretical explanation of emergent abilities remains contested. Schaeffer et al. [2023] argue that emergence is partly an artifact of measurement: metrics with sharp thresholds (pass@1 on code generation) create the appearance of sharp capability transitions even when underlying continuous capabilities improve smoothly. Other researchers maintain that genuine phase transitions occur in model capability with scale, analogous to phase transitions in physical systems, and that these transitions reflect the acquisition of qualitatively new computational strategies. The debate has practical implications for extrapolating current performance improvements to future model scales: if emergence reflects true phase transitions, extrapolating from current performance may underestimate future capability; if it is a measurement artifact, extrapolation should be more reliable.

CHAPTER 4: LLM APPLICATIONS IN SOFTWARE ENGINEERING

4.1 A Taxonomy of LLM Applications Across the Software Lifecycle

Large language models interact with the software development lifecycle at multiple points, from initial requirements specification through production operation and maintenance. Drawing on evidence from controlled studies, deployment reports, and benchmark evaluations across the 220 primary sources reviewed, the present chapter surveys fifteen application categories: code generation, code explanation, refactoring, bug detection and static analysis, debugging and fault localization, code review, test generation, documentation generation, DevOps and infrastructure automation, SQL and database code generation, API generation and specification, code migration and modernization, accessibility and inclusive development, natural language requirements to code, and emerging applications in software architecture assistance.

The taxonomy organizes these applications along two dimensions: the software lifecycle phase they primarily address, and the nature of the LLM's contribution (generation, understanding, transformation, or analysis). Understanding this dual organization helps practitioners identify where LLM adoption is most mature, where empirical evidence is strongest, and where emerging capabilities promise future productivity improvements. Figure 4.1 presents the taxonomy visually (not shown in text format); Table 4.1 summarizes key benchmark performance for code generation across the taxonomy.

4.2 Code Generation: From Autocomplete to Algorithm Synthesis

Code generation — the synthesis of executable source code from natural language specifications, examples, or partial implementations — is the most extensively studied and commercially deployed LLM application in software engineering. It encompasses a spectrum of sub-tasks that differ substantially in difficulty, context requirements, and evaluation methodology: intelligent autocomplete extending token-level prediction to multi-line suggestions; function synthesis generating complete implementations from docstrings; algorithm synthesis implementing complex computational procedures from abstract descriptions; and repository-level generation creating code that must integrate with existing multi-file projects.

Intelligent Autocomplete

The autocomplete paradigm, exemplified by GitHub Copilot, represents the most widely deployed form of LLM code generation. As a developer types, the system

generates suggestions ranging from a single identifier completion to multi-line code blocks, drawing on the current file context and (in more sophisticated systems) cross-file project context. Svyatkovskiy et al. [97] demonstrated that contextual code completion could achieve acceptance rates exceeding 35% in real-world usage, with accepted completions disproportionately concentrated in boilerplate patterns: standard class initialization code, exception handling templates, repeated API call patterns, and conventional test structure. For boilerplate code — which constitutes a significant fraction of total code volume in most large codebases — acceptance rates substantially exceed the average, as these patterns are highly predictable from context.

Latency requirements for inline autocomplete are stringent: developers expect suggestions within 100-300ms of pausing, and suggestions appearing after a longer delay provide diminished value as the developer has already continued typing. This latency requirement drives model selection for autocomplete: small, fast models (1-7B parameters with quantization) are preferred over large, accurate models (70B+ parameters) for inline suggestions, despite the capability trade-off. The tiered model architecture used by production systems such as GitHub Copilot — a fast small model for inline suggestions and a larger model for chat — reflects this latency-accuracy trade-off.

Function Synthesis and HumanEval Performance

Function synthesis — generating complete function implementations from docstrings or natural language descriptions — is the task most directly measured by HumanEval and MBPP benchmarks. Performance on HumanEval has improved from Codex's 28.8% pass@1 in 2021 to GPT-4o's 90.2% and Claude 3.5 Sonnet's 92.0% in 2024-2025, representing a 3x improvement in under four years. This rapid improvement trajectory reflects both architectural advances (larger models, better training data, improved alignment) and the HumanEval benchmark's relative accessibility — it consists of relatively simple self-contained Python functions whose difficulty distribution spans elementary to intermediate algorithmic programming.

The practical significance of HumanEval performance improvements should be interpreted carefully. HumanEval's problems are self-contained, well-specified, and now widely present in LLM training data, potentially inflating performance through memorization rather than generalization. LiveCodeBench [94], which continuously updates with fresh competitive programming problems, provides contamination-resistant estimates that are typically lower: models scoring 90%+ on HumanEval typically score 40-70% on LiveCodeBench, reflecting the performance degradation when benchmark contamination is controlled. The difference quantifies the contamination effect for each model and should inform practitioners weighing benchmark claims.

Model	HumanEval	MBPP	MultiPL-E (avg)	SWE-bench (V)	BigCodeBench	LiveCodeBench
GPT-4o	90.2%	86.8%	72.3%	49.0%	61.1%	58.3%
Claude 3.5	92.0%	91.0%	75.1%	49.0%	64.0%	61.2%

Model	HumanEval	MBPP	MultiPL-E (avg)	SWE-bench (V)	BigCodeBench	LiveCodeBench
Sonnet						
Gemini 1.5 Pro	84.1%	83.5%	68.4%	38.0%	56.2%	52.1%
DeepSeek-Coder-V2	90.2%	89.4%	73.0%	N/A	66.0%	60.8%
Qwen2.5-Coder-72B	88.4%	87.5%	72.0%	N/A	65.0%	59.4%
Code Llama 70B	53.0%	62.4%	46.2%	N/A	43.0%	38.7%
StarCoder2-15B	46.1%	54.5%	38.0%	N/A	36.0%	31.2%
Codex (2021)	28.8%	50.0%	N/A	N/A	N/A	N/A

Figure 3: Model Performance Across Four Benchmarks




Figure 3: Model Performance Across Four Benchmarks

Table 4.1: Code generation benchmark performance for leading models (pass@1, as of mid-2025). N/A = not reported.

Repository-Level Code Generation

Repository-level code generation — synthesizing code that must correctly integrate with an existing multi-file codebase — is substantially more challenging than self-contained function synthesis. The task requires understanding how the new code relates to existing modules, what naming conventions and style patterns are established, which interfaces the new code must implement or call, and how tests should be structured. SWE-bench [93], which requires modifying real-world Python repositories to resolve GitHub issues, provides the most ecologically valid evaluation of this capability, with frontier agentic systems achieving 37-49% resolution rates on the Verified subset — dramatically lower than their function-level benchmark performance.

RepoCoder [68] and CoCoST [161] have developed specialized retrieval-augmented generation approaches for repository-level code generation, using vector similarity search to retrieve relevant code snippets from the repository before generating new code. The integration of retrieved context substantially improves consistency of generated code with existing codebase conventions and API usage patterns. Longer-context models (200K-1M tokens) enable alternative approaches that simply include large portions of the codebase in the context, trading retrieval precision for context completeness. The empirical comparison of retrieval versus long-context approaches is an active research question, with results suggesting complementarity: retrieval efficiently identifies relevant code, while long context enables integration of retrieved context with broader project understanding.

4.3 Code Explanation and Comprehension Assistance

Code explanation — generating natural language descriptions of code behavior, intent, design decisions, and implementation details — is one of the most practically valuable yet least formally evaluated LLM applications in software engineering. The practical value is clear: professional software engineers spend an estimated 58% of their time reading and understanding code rather than writing it [211], and effective code explanation tools could substantially reduce this cognitive burden.

LLMs can explain code at multiple levels of abstraction, each serving different user needs. Line-by-line explanation provides detailed, granular description of each statement's effect, valuable for developers learning unfamiliar language features or understanding complex algorithms. Function-level explanation describes a function's purpose, parameters, return values, side effects, and key implementation decisions — essentially generating the documentation that should accompany the code but often does not. Module-level explanation describes how the components of a module relate to each other and to the broader system, supporting architectural understanding. System-level explanation describes how different modules and services interact to implement the overall system behavior, supporting maintainer onboarding and impact analysis for proposed changes.

Geng et al. [99] evaluated GPT-4 on code summarization across five programming languages, finding that GPT-4 produced explanations rated by professional developers as comparable to or exceeding developer-written comments in clarity and accuracy for functions under 50 lines. Performance degraded for longer code segments, reflecting the challenge of maintaining coherent high-level summaries across complex implementations with multiple interacting concerns. Microsoft Research studies [100] found that LLM-based code explanation reduced the time required for developers to understand unfamiliar codebases by 30-50%, with the largest improvements for

developers encountering programming languages or frameworks outside their primary expertise.

Explaining Legacy Code

A particularly high-value application of LLM-based code explanation is in legacy code understanding. Many organizations maintain substantial codebases written in languages or frameworks with limited current developer expertise — COBOL financial systems, RPG manufacturing applications, early Java enterprise systems with outdated patterns. The original developers may have retired or moved on; documentation may be absent or outdated; and the code may have accumulated decades of patches that obscure the original design intent.

LLMs trained on diverse code corpora can explain COBOL data divisions, identify the business rules implemented in undocumented PL/SQL stored procedures, and describe the state machine implicit in complex conditional logic — capabilities that would otherwise require expensive specialist consultants. Early case studies from financial institutions report that LLM-assisted COBOL explanation reduced the time required for junior developers to understand transaction processing logic by 40-60%, enabling more efficient legacy system maintenance and providing the documentation baseline needed for modernization planning.

4.4 Code Refactoring and Quality Improvement

Code refactoring — restructuring existing code to improve internal quality metrics including maintainability, readability, performance, and testability without altering external behavior — is a critical software maintenance activity that LLMs can both suggest and, in agentic configurations, perform. The spectrum of refactoring operations spans from simple renaming (renaming a poorly named variable to improve clarity) through extract method (identifying repeated logic and extracting it into a reusable function) to complex architectural restructuring (converting a monolithic service to microservices or migrating from one design pattern to another).

Simple refactoring operations — renaming, extract method, inline variable, move method — are well within the capability of all frontier LLMs and can be applied reliably through standard code transformation prompts. Liu et al. [101] evaluated LLM-generated refactoring suggestions on a dataset of 500 Python and Java code samples, finding that developer-accepted refactoring suggestions improved code quality metrics (measured by complexity, coupling, and readability scores) in 68% of cases. The primary failure modes were: semantic drift (the refactored code subtly altered behavior despite the developer's intent to preserve it), stylistic inconsistency (the refactored code diverged from the surrounding code's style conventions), and incomplete refactoring (the model refactored the highlighted code but did not update all call sites or dependent code).

Complex architectural refactoring — converting a class hierarchy to composition, migrating from callbacks to async/await patterns, or restructuring a monolith into services — presents substantially greater challenges. These tasks require understanding the broader system context that motivates the architectural change, reasoning about invariants that must be preserved across the transformation, and making design decisions about where to draw architectural boundaries. Agentic systems with large context windows and iterative test execution can approach these tasks more reliably than single-turn generation, but even the best current systems require careful human oversight for complex refactoring in production codebases.

Design Pattern Application

LLMs can suggest and implement design pattern transformations — converting procedural code to use the Strategy pattern, adding the Observer pattern for event handling, applying the Factory pattern to decouple object creation — with reasonable accuracy when the target pattern and its application context are clearly specified in the prompt. The model's training on pattern descriptions and examples enables recognition of code structures amenable to pattern application and generation of pattern-compliant refactored code. The challenge is that design patterns involve architectural commitments whose appropriateness depends on the broader system context that a narrow code view cannot fully capture. A well-applied Factory pattern in one context may be over-engineering in another; without understanding the system's extensibility requirements, an LLM cannot reliably evaluate which patterns to recommend.

4.5 Bug Detection and Static Analysis Augmentation

LLM-based bug detection extends traditional static analysis approaches by leveraging semantic understanding to identify bugs that emerge from the complex interaction of code semantics, developer intent, and execution context — bugs invisible to rule-based static analysis that checks only syntactic patterns. While tools like SonarQube, ESLint, and CodeClimate excel at identifying violations of predefined rules (missing null checks, potential resource leaks, deprecated API usage), they cannot reason about semantic bugs: incorrect algorithmic implementations, off-by-one errors, race conditions in concurrent code, or logic bugs in complex conditional structures.

Wu et al. [102] evaluated GPT-4 on a dataset of 500 real bugs from 12 open-source Python and Java projects, finding that GPT-4 correctly identified the buggy code region in 43% of cases — significantly better than rule-based static analysis tools on semantic bugs (which they largely cannot detect) but substantially worse than specialized analysis tools on syntactic vulnerability patterns (where established tools excel). The practical complement of LLM semantic analysis with traditional rule-based static analysis creates a detection system that outperforms either approach alone: a study by Pearce et al.

[103] found that the combined system achieved 73% detection rate for a mixed bug dataset compared to 52% for LLM-only and 38% for rule-based only, demonstrating clear complementarity.

Security Vulnerability Detection

Security vulnerability detection is a particularly high-stakes sub-domain of bug detection where LLMs have demonstrated both significant promise and significant limitations. Pearce et al.'s landmark study [40] found that approximately 40% of code generated by GitHub Copilot for security-sensitive scenarios contained at least one security vulnerability from the OWASP Top 10 categories, including SQL injection through unsanitized user input, buffer overflows in C memory management, use of deprecated cryptographic functions (MD5, SHA-1 for password hashing), and hard-coded credentials. These findings motivated significant research attention on both detection of LLM-generated vulnerabilities and improvement of LLM-based vulnerability detection capabilities.

Khoury et al. [104] found that models can detect at least some of their own generated vulnerabilities when explicitly prompted to review generated code for security issues — suggesting that the knowledge necessary to identify vulnerabilities is present in the model but not automatically activated during generation. Self-review prompting ('Please review the code you just generated for security vulnerabilities') reduces the rate of undetected vulnerabilities in generated code by 15-30% in controlled evaluations, though this remains well short of expert security code review. Commercial tools including Snyk Code, Semgrep, and GitHub's code scanning have begun integrating LLM-based analysis alongside traditional pattern matching, reporting improved detection rates for complex vulnerability patterns.

SAST (Static Application Security Testing) tools augmented with LLM reasoning are emerging as a practical enterprise security control for LLM-assisted development pipelines. These tools scan generated code in real-time, flagging potential vulnerabilities before they are committed to repositories. Amazon Q Developer's built-in security scanning feature, using a combination of rule-based detection for known vulnerability signatures and an LLM-based learned vulnerability model, represents a production deployment of this approach. Early enterprise deployment reports indicate meaningful reductions in security defect rates compared to pure LLM generation without security scanning.

4.6 Debugging, Fault Localization, and Automated Program Repair

Automated debugging — identifying, localizing, and correcting software defects — represents one of the highest-potential-value applications of LLMs in software engineering. The debugging process encompasses multiple sub-tasks: fault localization

(identifying which code element is responsible for observed incorrect behavior), root cause analysis (understanding why the identified code produces the observed behavior), fix generation (implementing a correction), and validation (confirming that the fix resolves the issue without introducing regressions). LLMs have demonstrated meaningful capability at each sub-task, though end-to-end automated debugging of complex real-world defects remains challenging.

Xia et al. [105] evaluated ChatGPT (GPT-3.5) on 337 bugs from five open-source Java projects in the Defects4J benchmark, finding correct fixes for 109 bugs (32.3%) — substantially more than leading automated program repair (APR) tools including TBar, which correctly fixed 43 bugs on the same benchmark using developer-designed fix patterns. The LLM's broader knowledge of code patterns and repair strategies enabled more diverse fix approaches than template-based APR tools, though the LLM also produced more incorrect fixes that superficially appeared correct (plausible patches that passed some tests but introduced regressions). The combination of LLM-generated patches with test suite validation — accepting only patches that pass the full test suite — substantially improves fix reliability.

Multi-Turn Interactive Debugging

LLMs demonstrate stronger debugging performance in multi-turn interactive settings than in single-turn generation. Providing the model with the error message, stack trace, relevant code context, and (critically) the results of executing suggested fixes creates a feedback loop that significantly improves diagnostic accuracy. The ChatRepair approach [106] combines LLMs with conversation-driven patch generation and test execution feedback, achieving 162 bug fixes on the Defects4J benchmark — outperforming all prior automated repair systems. The key innovation is systematic use of test failure information to guide iterative patch refinement: when a patch fails tests, the failure output is provided back to the model as context for generating an improved patch.

Agentic debugging systems represent the current frontier of automated debugging capability. Claude Code, Devin, and OpenHands demonstrate the ability to: examine error messages and stack traces; navigate to relevant source code; formulate debugging hypotheses; design and execute targeted experiments (adding logging statements, writing test cases for specific code paths); update hypotheses based on results; and implement and validate fixes. These autonomous debugging capabilities succeed reliably for classes of bugs where the symptom is informative, the code is accessible, and the fix is localized — and struggle with bugs requiring deep domain knowledge, complex distributed system interactions, or non-deterministic manifestations.

4.7 Automated Code Review

Code review — the systematic examination of source code by one or more developers to identify defects, security issues, design problems, and maintenance concerns before code is merged into the main codebase — is a fundamental quality assurance practice in professional software development but also one of the most time-consuming. Studies consistently find that code review consumes 5-15% of total development time across organizations, with delays in review turnaround representing a significant velocity impediment. LLM-based automated code review offers the potential to reduce review burden by handling routine review comments automatically, enabling human reviewers to focus on higher-level design and security concerns.

Li et al. [107] evaluated LLM-generated code reviews on 1,000 pull requests from 20 active open-source projects, comparing GPT-4-generated reviews with human reviews from the same pull requests. GPT-4-generated reviews identified relevant concerns in 71% of pull requests (compared to 78% for human reviews), with particularly strong coverage of documentation incompleteness, naming clarity violations, error handling gaps, and code style inconsistencies. Human reviews showed stronger performance for logic correctness issues (59% vs 41% for LLM) and architectural concerns (63% vs 38%), reflecting the advantage of human contextual understanding for evaluating whether code is not just syntactically correct but appropriate for its intended purpose in the broader system.

GitHub Copilot's pull request review feature, integrated into GitHub Actions and pull request workflows, provides automated review comments directly on pull request diffs. Amazon Q Developer's automated code review capability similarly integrates with pull request workflows, focusing specifically on security vulnerabilities, performance issues, and AWS best practice violations. These production deployments have moved LLM code review from research prototype to enterprise practice, with adoption particularly strong among teams where human reviewer capacity is limited and cycle time reduction is a priority.

Review Comment Quality and Developer Acceptance

The practical utility of automated code review depends not just on whether issues are identified but on whether review comments are actionable and accepted by developers. Studies of developer acceptance of LLM-generated review comments find acceptance rates of 25-40% for tool-generated comments versus 55-70% for human review comments, reflecting both the higher quality of human contextual judgment and developer preferences for peer interaction over automated feedback. Acceptance rates improve substantially when review comments include: specific citations of relevant coding standards or security guidelines; concrete suggested fixes rather than issue identification alone; and clear explanation of the potential consequences if the identified

issue is not addressed. Prompt engineering for review generation — providing the model with relevant coding standards, repository conventions, and examples of accepted review comments — significantly improves comment quality and acceptance rates.

4.8 Automated Test Generation

Test generation — automatically creating unit tests, integration tests, and end-to-end tests from code specifications, existing implementations, or natural language descriptions — is an application where LLMs demonstrate practical value that complements existing automated test generation tools. Traditional automated test generation tools such as EvoSuite (for Java) and Hypothesis (for Python) use evolutionary or property-based techniques to generate tests that maximize code coverage, producing tests that exercise many code paths but often lack semantic meaningfulness: test names are auto-generated identifiers, assertions are coverage-driven rather than specification-driven, and the tests do not reflect what developers actually care about in the code's behavior.

Schafer et al. [108] conducted a large-scale evaluation of LLM-generated unit tests across 1,000 Python functions from 20 open-source repositories, comparing LLM-generated tests with EvoSuite-generated tests on coverage metrics and developer quality ratings. LLM-generated tests achieved line coverage of 65-78% (compared to EvoSuite's 60-72%), with similar branch coverage. However, developer quality ratings showed a more dramatic difference: developers rated LLM-generated tests as 'useful' or 'very useful' in 73% of cases versus 42% for EvoSuite tests, reflecting the LLM tests' semantic meaningfulness — tests that verified business logic and edge cases that developers actually cared about, rather than arbitrarily exercising code paths.

Dakhel et al. [109] demonstrated a particularly valuable capability: test-driven development automation where LLMs generate tests from function signatures and docstrings before the implementation, providing formal specifications that then guide implementation generation. This TDD-automation workflow creates a virtuous cycle: tests generated from the specification constrain the LLM's subsequent implementation, reducing the space of acceptable implementations and improving functional correctness. The combination of specification-to-test-to-implementation is closer to formal software synthesis than typical LLM code generation and represents a promising research direction for high-reliability software development.

Test Type	LLM Capability Level	Key Tools	Coverage Achievement	Key Challenge
Unit tests	High (mature)	Copilot, Claude, Cursor	65-78% line coverage	Assertion quality
Integration tests	Medium	Amazon Q, Claude Code	40-60%	Mock/dependency setup
End-to-end tests	Medium	Playwright+LLM	50-70% critical paths	Environment configuration
Property-based tests	Low-Medium	Hypothesis+LLM	Property discovery	Invariant

Test Type	LLM Capability Level	Key Tools	Coverage Achievement	Key Challenge
			difficult	specification
Security tests	Medium	Burp+LLM, OWASP	OWASP Top 10 coverage	Novel attack patterns
Performance tests	Low	k6+LLM	Baseline coverage	Threshold specification
Mutation tests	Low	PITest+LLM	Basic mutations	Equivalent mutation detection

Table 4.3: LLM capabilities in test generation by test type, as of mid-2025.

4.9 Documentation Generation and Maintenance

Documentation — including inline comments, function docstrings, API reference documentation, architectural documentation, user guides, and tutorials — is universally acknowledged as important for software maintenance and knowledge transfer, yet consistently underprioritized in practice. Time pressure during development, the rapid obsolescence of documentation as code evolves, and the cognitive burden of context-switching between implementation and explanation all contribute to documentation deficits that accumulate into significant maintenance liabilities. LLMs address these challenges by enabling automatic generation of documentation directly from code, substantially reducing the marginal cost of documentation production.

Hu et al. [110] provided the first large-scale neural code summarization study, establishing that neural approaches substantially outperform template-based and IR-based documentation generation methods across multiple programming languages. Subsequent work has demonstrated continued improvement with scale: GPT-4-generated function docstrings are rated by professional developers as meeting or exceeding their documentation standards in approximately 70% of cases when evaluated blind, with the primary shortcomings in describing parameter constraints, error conditions, and interaction effects with other system components.

CoT-CodeSumm [111] demonstrated that chain-of-thought prompting improves code summarization quality by structuring the generation process: the model first identifies the function's purpose at a high level, then describes its key algorithmic steps, then notes important constraints or edge cases, and finally synthesizes these into a coherent docstring. This structured approach reduces the incidence of docstrings that accurately describe what the code does but fail to explain why design decisions were made — the documentation content that is most valuable for future maintainers.

Documentation at Scale: Enterprise Deployments

Enterprise documentation generation deployments represent some of the most clearly quantified LLM productivity improvements in the literature. Stripe's documentation automation initiative, using LLM-generated first drafts reviewed and

refined by technical writers, reduced documentation production time by 67% while improving consistency across the API reference documentation [112]. The model generated accurate first drafts that technical writers then refined for clarity, completeness, and consistency — a workflow that leverages both LLM productivity and human judgment for quality. Atlassian reported similar results for their Confluence AI documentation features, with developers reporting a 40% reduction in time spent writing documentation when AI assistance was available.

Documentation consistency — ensuring that terminology, formatting, and level of detail are uniform across a large API surface — is a particular benefit of LLM-based documentation generation. Human-written documentation for large APIs inevitably varies across authors and time periods; LLM-generated documentation can be produced from a consistent prompt template that enforces style guidelines, terminology choices, and structural requirements. This consistency benefit may ultimately be as valuable as the productivity benefit for documentation consumers who navigate across different parts of a large system.

4.10 DevOps and Infrastructure Automation

The application of LLMs to DevOps and infrastructure automation represents a rapidly growing domain with significant productivity implications. DevOps encompasses the practices, tools, and cultural philosophies that integrate software development and IT operations — automating the pipeline from code commit through build, test, deployment, monitoring, and incident response. LLMs support this pipeline at multiple points, from generating pipeline configurations to explaining build failures to responding to production incidents.

CI/CD pipeline configuration generation represents one of the most mature LLM DevOps applications. Developers can describe desired build, test, and deployment behavior in natural language and receive a functional GitHub Actions workflow YAML, Jenkins pipeline script, or GitLab CI configuration. Liu et al. [113] evaluated LLM-generated CI/CD configurations against manually created alternatives, finding functional equivalence in 78% of cases with substantially reduced configuration authoring time. The most common failure modes were: incorrect handling of environment-specific variables and secrets, overly permissive permission grants that represent security risks, and configurations that work for the described scenario but lack the error handling and notification steps that production pipelines require.

Infrastructure as Code (IaC) generation for Terraform, Ansible, CloudFormation, and Pulumi represents a high-value application with documented productivity improvements. The complexity of cloud infrastructure configuration languages — deeply nested JSON/YAML schemas, provider-specific resource types, complex dependency relationships between resources — makes manual authoring time-consuming and error-

prone. LLM generation from natural language descriptions ('Create an AWS VPC with public and private subnets, an Application Load Balancer, auto-scaling EC2 instances in the private subnet, and an RDS PostgreSQL instance with Multi-AZ deployment') produces complete, syntactically valid configurations in seconds that would require hours to write manually. Studies report 40-60% reductions in IaC authoring time for standard infrastructure patterns [114].

DevOps Domain	LLM Capability	Primary Tools	Productivity Impact	Maturity
CI/CD Pipelines	Config generation	Copilot, Claude	40-60% faster	High (production)
Terraform/IaC	Config synthesis	Claude, Copilot	40-60% faster	Medium-High
Kubernetes YAML	Manifest generation	All major LLMs	50% faster	Medium
Docker	Dockerfile generation	All major LLMs	60% faster	High (production)
Log Analysis	Anomaly explanation	Amazon Q, Datadog AI	30% faster MTTD	Medium
Incident Response	Alert correlation	PagerDuty AI, Copilot	25% faster MTTR	Medium
Monitoring	Dashboard generation	Grafana AI	35% faster setup	Medium
Security (SAST)	Vulnerability detection	Snyk, Semgrep AI	50% more findings	High

Table 4.2: LLM capabilities in DevOps domains. Productivity figures are approximate estimates from reported studies.

4.11 SQL and Database Code Generation

SQL generation — producing structured query language statements from natural language descriptions — is one of the most practically valuable and technically well-studied LLM applications in software engineering. SQL is required by virtually every application that interacts with relational databases, yet complex queries involving joins, subqueries, window functions, aggregations, and database-specific dialects present significant barriers for developers without SQL expertise. The natural language to SQL (text-to-SQL) task has been studied extensively, with a progression of benchmarks — WikiSQL, Spider, BIRD — tracking the field's progress.

Performance on the Spider benchmark's test set has improved from approximately 60% execution accuracy for early neural approaches to over 85% for GPT-4-based systems with appropriate prompting, representing a practical threshold for deployment in business intelligence applications. DIN-SQL [134] demonstrated that decomposing complex SQL generation into sub-problems using chain-of-thought prompting substantially improved accuracy on challenging multi-join and subquery-heavy queries. The BIRD benchmark, introduced in 2023, addresses limitations of Spider by incorporating realistic database values and requiring queries that correctly handle ambiguous specifications, NULL values, and database-specific SQL dialects — conditions that better reflect production database environments.

Production SQL generation deployments must address several additional challenges beyond query correctness. Injection safety requires that generated queries be parameterized rather than string-interpolated, preventing SQL injection vulnerabilities when user-provided values are incorporated. Performance safety requires validation that generated queries will not trigger full table scans on large tables — an inadvertent full-scan query on a 10-billion-row table could overload a production database. Access control requires validation that generated queries do not access tables or columns beyond the querying user's permissions. Commercial text-to-SQL platforms including Vanna.AI, Defog, and enterprise BI tools (Tableau AI, Power BI Copilot) have productized LLM SQL generation with these safety guardrails, enabling business users without SQL expertise to query databases through natural language interfaces.

4.12 API Generation and Specification

API generation encompasses two related and high-value tasks: generating code that correctly calls external APIs given their documentation, and generating API interface definitions (OpenAPI/Swagger specifications, GraphQL schemas, gRPC protobuf definitions) from natural language descriptions of desired API behavior. Both tasks are well-suited to LLMs, given the structured and schema-driven nature of API interfaces and the availability of extensive API documentation and usage examples in LLM training data.

For API client code generation, the primary challenge is ensuring that generated code correctly handles authentication, rate limiting, error handling, and pagination — aspects of API usage that vary across providers and that are easy to get wrong. Models trained on large quantities of API usage examples from public repositories generally perform well on widely-used APIs (Stripe, Twilio, GitHub API, AWS SDK) where training examples are abundant, but may hallucinate method signatures for less-represented APIs. RAG augmentation with API documentation substantially improves accuracy: Xu et al. [2024] reported 40-60% reduction in hallucinated API calls when relevant documentation sections were included in the generation context, validating the importance of documentation-grounded generation for API clients.

OpenAPI specification generation from natural language descriptions enables rapid API prototyping: a developer describing desired endpoints, request schemas, response schemas, and authentication requirements in natural language can receive a complete, valid OpenAPI 3.0 specification in seconds. This specification can immediately be used to generate client SDKs, server stubs, and interactive documentation through standard OpenAPI tooling (Swagger UI, FastAPI, Django REST Framework). The ability to move from high-level natural language description to deployable API contract without manual schema authoring represents a significant reduction in API design cycle time.

4.13 Code Migration and Modernization

Code migration — translating programs from one programming language, framework, API version, or architectural paradigm to another — is an application where LLMs demonstrate compelling practical value, particularly for organizations facing the challenges of legacy system modernization. Manual migration of large codebases is expensive, error-prone, and requires expertise in both source and target technologies that may be scarce. LLMs, having been trained on code in dozens of programming languages and framework versions, can leverage cross-lingual understanding to perform translation with semantic awareness beyond simple syntactic transformation.

Common migration scenarios where LLMs provide practical value include: Python 2 to Python 3 conversion (handling string encoding, print function syntax, integer division, and standard library changes); Java 8 to Java 17 modernization (incorporating records, sealed classes, pattern matching, text blocks, and modernized API usage); AngularJS to Angular 2+ migration (converting directive-based architecture to component-based); jQuery to vanilla JavaScript (replacing DOM manipulation patterns with modern APIs); and cloud provider API migrations (AWS SDK v2 to v3, deprecated Azure API versions). For each of these well-defined migration pairs, LLMs achieve 70-85% functional correctness on the migrated code — transformations that pass existing test suites — substantially reducing the manual effort required for migration campaigns.

Amazon's documented migration of approximately 30,000 Java 8 services to Java 17, using an LLM-assisted migration pipeline presented at AWS re:Invent 2023, demonstrated the practical scale at which LLM-based migration can operate. The pipeline generated migration patches using Claude, validated patches against existing test suites, flagged cases where tests failed for manual review, and applied validated patches automatically. This combination of LLM generation, automated testing, and targeted human review for failures achieved approximately 50% reduction in migration effort compared to purely manual approaches. The approach generalizes to any migration scenario with an adequate test suite: the tests provide the correctness oracle that enables automated validation of LLM-generated transformations.

4.14 Accessibility and Inclusive Development

LLMs have the potential to substantially democratize software development by reducing barriers for developers with varying expertise levels, native languages, and accessibility needs. These democratization effects extend both to professional developers with non-standard working conditions and to non-specialist users ('citizen developers') who create software to automate their own workflows without formal programming education.

For non-native English speakers — who constitute a large and growing proportion of the global developer community — LLMs that accept queries in any language and generate code with explanations in the developer's preferred language reduce cognitive load and communication friction. A developer who thinks primarily in Mandarin, Hindi, or Portuguese can describe their programming intent in their native language and receive both code and explanation without the additional cognitive burden of formulating technical queries in English. This capability has particular significance for the large developer communities in India, China, Brazil, and other regions where non-English development is common.

For developers with visual impairments who rely on screen readers and non-visual interfaces, LLM-based code explanation and generation through voice interfaces provides a more accessible alternative to visual IDEs. GitHub's work on Copilot accessibility features and Microsoft's broader AI accessibility initiatives reflect recognition of this potential. Research by Albusays et al. [2021] documented the significant challenges faced by blind and visually impaired programmers with conventional development tools, establishing motivation for LLM-based accessibility improvements that go beyond visual augmentations.

4.15 Natural Language Requirements to Code

The full pipeline from natural language requirements specification to executable code — collapsing the traditional distinction between requirements engineering, design, and implementation — represents the most ambitious application of LLMs in software engineering and the furthest from current reliable capability. Early research and commercial demonstrations suggest this capability is approaching practical usability for simple applications, though significant limitations persist for complex, multi-component systems.

Devin's demonstrated capability to implement a simple web application from a natural language description — including creating HTML, CSS, and JavaScript files, setting up a local server, and deploying to a cloud platform — illustrates the current state of the art for end-to-end autonomous software development. OpenDevin's academic evaluation on SWE-bench provides a more controlled assessment: current systems can resolve approximately 37-49% of GitHub issues that require multi-file changes, suggesting that autonomous implementation from requirements is approaching practical utility for well-defined, well-scoped tasks in familiar domains.

The primary limitations for requirements-to-code automation are: under-specification in natural language requirements that does not provide the precise detail needed for implementation without additional clarification; ambiguity about non-functional requirements (performance, security, scalability) that are critical for production systems; and the fundamental challenge that complex system architectures

require design decisions that depend on organizational context, technical constraints, and operational requirements not expressible in simple natural language specifications. Research addressing these limitations through structured requirement elicitation, iterative refinement, and formal specification grounding is an active area with direct practical implications for the evolution of software engineering practice.

CHAPTER 5: COMPARATIVE STUDY OF LLM-BASED SOFTWARE ENGINEERING TOOLS

5.1 Evaluation Framework and Methodology

Ten leading LLM-based software engineering tools and systems are compared across ten evaluation dimensions designed to capture the full range of considerations relevant to tool selection in professional and enterprise contexts. The evaluation framework integrates quantitative benchmark performance with qualitative assessment of enterprise readiness, integration capability, and total cost of ownership, producing a multi-dimensional comparison that supports nuanced tool selection decisions.

The ten evaluation dimensions are: (1) functional correctness on HumanEval (pass@1); (2) real-world task resolution on SWE-bench Verified; (3) multilingual performance on MultiPL-E (average across 18 languages); (4) context window capacity (maximum tokens); (5) inference cost per million tokens; (6) typical response latency for code completion tasks; (7) hallucination rate on code-specific tasks; (8) multi-step reasoning capability; (9) enterprise readiness (security, compliance, SLA, support); and (10) ecosystem integration (IDE support, CI/CD integration, API availability). These dimensions collectively address both the technical capability questions that inform performance expectations and the operational and organizational questions that determine deployment feasibility.

The tools evaluated span three categories: (A) general-purpose frontier models applied to software engineering (GPT-4o, Claude 3.5 Sonnet, Gemini 1.5 Pro), which provide maximum capability but require integration effort; (B) specialized developer tools (GitHub Copilot, Amazon Q Developer, Cursor, Windsurf), which provide optimized developer experience at the cost of flexibility; and (C) agentic coding systems (Claude Code, Devin, OpenHands), which provide autonomous multi-step task execution at the cost of predictability. This category structure reflects the evolution of LLM deployment from direct model access through tool-mediated experiences to autonomous agent systems.

5.2 GPT-4o: Analysis and Software Engineering Profile

GPT-4o (Omni), released by OpenAI in May 2024, represents the current production frontier of the GPT model family, integrating text, code, image, and audio processing in a unified architecture. For software engineering applications, GPT-4o demonstrates strong performance across the evaluation framework: 90.2% pass@1 on HumanEval, 49% resolution on SWE-bench Verified (through agentic deployment), and 72.3% average on MultiPL-E across 18 programming languages.

GPT-4o's primary strengths for software engineering include: a 128K token context window enabling processing of large code files and moderate-scale codebases; sophisticated multi-step reasoning demonstrated through strong performance on competitive programming and algorithm design tasks; broad programming language coverage including low-level systems languages (C, C++, Rust, Assembly) where alternative models may be less capable; and a mature API with function calling, structured outputs, and integration SDKs supporting diverse deployment patterns. The model's primary limitations include: higher cost relative to open-source alternatives at the same capability tier; lack of organizational fine-tuning options for aligning the model to internal code patterns; and the potential for confident hallucinations that require active verification.

GPT-4o's role in GitHub Copilot and numerous third-party developer tools makes it the most widely deployed frontier model in software engineering production contexts. Its position as the underlying model for the most-adopted AI coding tool means that performance improvements to GPT-4o translate directly to experience improvements for millions of developers. The model's consistent API and stable performance over time distinguish it from research models optimized for benchmark performance without consideration of deployment reliability.

5.3 Claude 3.5 Sonnet: Analysis and Software Engineering Profile

Claude 3.5 Sonnet, released by Anthropic in June 2024, achieved leading performance on several coding benchmarks at release while demonstrating distinctive characteristics attributable to Constitutional AI training that differentiate it from models trained with pure RLHF. The model's 92.0% HumanEval pass@1 and 49% SWE-bench Verified resolution place it at the frontier of functional code generation capability, while its 200,000-token context window — the largest among frontier models at release — enables repository-level code understanding inaccessible to shorter-context alternatives.

Claude 3.5 Sonnet's alignment characteristics manifest in specific, practically relevant software engineering behaviors. The model more consistently acknowledges when its generated code may have edge cases or limitations not explicitly covered by the specification — a property that reduces false confidence and encourages appropriate testing. It provides more nuanced security assessments of generated code, more reliably noting when implementation choices create security trade-offs. It is more resistant to prompt injection in agentic contexts — an important security property for autonomous coding agents that process potentially adversarial inputs. These alignment benefits complement raw capability metrics to produce a model with a distinctive enterprise value proposition.

The model serves as the underlying intelligence for Claude Code, Anthropic's agentic coding assistant, leveraging both the large context window for codebase-level

awareness and the reasoning capabilities for autonomous multi-step task execution. Early deployment reports from software engineering teams using Claude Code indicate particularly strong performance on large-scale refactoring tasks, documentation generation for complex codebases, and test-driven development workflows where the agent iterates on implementation until tests pass — all tasks that benefit from the large context window and consistent alignment properties.

5.4 Gemini 1.5 Pro: Analysis and Software Engineering Profile

Gemini 1.5 Pro's defining characteristic for software engineering is its one-million-token context window, which enables qualitatively different software engineering tasks than any other available model. With one million tokens at approximately 750-1000 tokens per typical code page, the model can accommodate codebases of 750,000-1,000,000 lines of code in a single context — a scale sufficient to encompass many large enterprise software projects in their entirety. This capability is particularly valuable for: legacy system analysis where the full codebase must be understood before planning modernization; cross-repository dependency analysis for security audits; complete API documentation generation for large SDKs; and extended agentic sessions that must maintain awareness of extensive prior context.

Gemini 1.5 Pro's native multimodal capabilities provide software engineering applications that single-modality models cannot support. UI implementation from design mockups: the model can process a Figma export or screenshot of a UI design and generate corresponding React, Vue, or HTML/CSS implementation. Architecture diagram interpretation: uploaded system architecture diagrams can be interpreted to generate corresponding infrastructure configurations or code stubs. Error screenshot analysis: screenshots of error dialogs or unexpected UI states can be combined with code context for more informative debugging assistance. These multimodal capabilities are particularly valuable for front-end development and for bridging the gap between visual design artifacts and code implementation.

Gemini 1.5 Pro's HumanEval performance (84.1%) is modestly below frontier alternatives, but its context length advantage may compensate substantially in repository-level tasks where broader context enables better understanding of code relationships and conventions. The model's pricing (\$3.50 per million input tokens, \$10.50 per million output tokens as of mid-2025) is competitive with comparable capability alternatives, and its integration with Google's development ecosystem (Android Studio, Cloud Code, Google Workspace) provides natural deployment paths for organizations already within the Google ecosystem.

5.5 GitHub Copilot: Analysis and Enterprise Impact

GitHub Copilot represents the most widely adopted LLM-based software engineering tool, with millions of active users across individual and enterprise accounts as of mid-2025. Its position as the first broadly adopted AI coding assistant provides a unique evidence base: more controlled productivity studies, more longitudinal usage analyses, and more enterprise deployment reports have been published about Copilot than any comparable tool, enabling more evidence-grounded assessment of real-world impact.

The productivity evidence for Copilot is the most thoroughly documented in the field. Kalliamvakou's controlled study [50] at GitHub Research found 55% task completion speed improvement for a standardized coding task; Ziegler et al.'s large-scale observational study [115] found sustained 26-35% acceptance rates for suggestions across thousands of developers; Accenture's enterprise deployment report [117] documented 35% productivity improvement for routine coding tasks; and McKinsey's analysis [118] estimated 20-45% improvement across mixed task portfolios. While these estimates vary substantially — reflecting genuine variation in task types, developer experience, and measurement methodology — the convergence of evidence from multiple independent studies and measurement approaches provides stronger confidence in meaningful productivity benefits than any single study could provide.

Copilot's evolution from its initial autocomplete-focused design to a broad AI development platform — encompassing chat assistance, commit message generation, pull request summaries, automated code review, and the Copilot Workspace for specification-to-implementation workflows — represents a strategy of expanding LLM assistance across the full development lifecycle rather than deepening a single capability. This breadth strategy aligns with the evidence that the largest productivity benefits come from deep workflow integration rather than point tools that require context-switching.

5.6 Amazon Q Developer: Security-First Enterprise AI

Amazon Q Developer, formerly known as Amazon CodeWhisperer, differentiates itself within the competitive AI coding assistant landscape through two distinctive capabilities: built-in real-time security scanning of generated code and deep integration with the AWS service ecosystem. These differentiators reflect Amazon's specific context as both a developer tool provider and a cloud platform company with particular expertise in cloud security and AWS service semantics.

The security scanning capability uses a hybrid approach combining rule-based detection for known vulnerability signatures (OWASP Top 10 categories, CWE Top 25 most dangerous software weaknesses) with a learned vulnerability detection model trained on labeled security defects. This hybrid approach achieves lower false negative rates than either approach alone: rule-based detection provides reliable coverage for

known patterns while the learned model extends coverage to novel vulnerability forms that violate no explicit rule. The real-time scanning, integrated into the suggestion rendering pipeline rather than as a post-hoc scanner, enables developers to see security feedback before accepting suggestions rather than after committing code — a timing advantage that reduces the friction of security validation.

Amazon Q Developer's AWS ecosystem integration provides contextually accurate code generation for developers building on AWS: Lambda function configurations, DynamoDB query patterns, S3 access patterns with appropriate error handling, CloudFormation resource definitions, and IAM policy constructions are all generated with awareness of AWS-specific conventions and best practices that general models may lack. The model's training on AWS documentation, SDK examples, and AWS-focused code from the training corpus produces suggestions that align with AWS architectural patterns and avoid common misconfigurations that generate security issues or cost inefficiencies.

5.7 Cursor and Windsurf: AI-Native IDE Analysis

Cursor and Windsurf represent a distinct architectural category within the AI coding assistant landscape: AI-native IDEs that build the complete development environment around LLM capabilities rather than augmenting existing IDEs with AI plugins. This architectural approach provides several advantages over plugin-based tools: full project context access through deep codebase indexing; deeper integration between AI capabilities and editor functionality; and the ability to coordinate multi-file editing operations that require simultaneous awareness of multiple files' contents and structure.

Cursor's rapid growth to over 150,000 monthly active users as of 2024 reflects strong product-market fit among professional developers who find the deeper AI integration — particularly multi-file editing through Composer and autonomous task execution through Agent mode — more productive than the single-file suggestion paradigm of plugin-based tools. User surveys consistently cite multi-file editing capability as the top differentiating feature: tasks like 'add a new API endpoint with tests, update the router, and add the model' that require simultaneous modification of multiple files are dramatically more efficient in Cursor than in plugin-based alternatives requiring manual file-by-file interaction.

Windsurf's 'flows' paradigm addresses a genuine limitation of context-stateless tools: when each AI interaction begins without memory of recent developer actions, the model cannot maintain awareness of work done in previous interactions within the same session. Cascade, Windsurf's session-aware agent, tracks recent file modifications, test runs, and error messages across the session, providing contextually aware assistance without requiring the developer to re-explain context at each interaction. Early usage data suggests that session-aware context substantially reduces the number of follow-up

interactions required to complete complex tasks, potentially more than compensating for the computational cost of maintaining session state.

5.8 Claude Code and Devin: Agentic System Analysis

Claude Code and Devin represent the frontier of agentic coding systems — tools designed for autonomous execution of complex, multi-step software engineering tasks that require extended sequences of actions rather than single-turn responses. These systems raise qualitatively different evaluation questions than single-turn code generation tools: not just 'does the generated code work?' but 'can the system autonomously navigate to a solution across multiple decision points and error recovery cycles?'

Claude Code's terminal-based agentic architecture reflects a philosophy of developer collaboration: the system operates with explicit action transparency (reporting each action before execution), requests permission before potentially destructive operations, and supports interruption and redirection at any point. This controlled autonomy model sits between fully autonomous execution (which may take unintended actions) and pure suggestion (which requires developer action at each step). The architecture enables workflows like 'implement this feature, run the tests, fix any failures, and prepare a pull request' that complete without intermediate interaction but maintain auditability through logged action sequences.

Devin's architecture places greater emphasis on long-horizon autonomous execution: the persistent memory system enables continuity across tasks that span hours of execution, and the web browsing capability enables research into unfamiliar APIs and error conditions that are not present in the local codebase. These capabilities enable handling of a broader class of engineering tasks — including tasks requiring configuration of external services, installation of dependencies, and implementation following documentation research — but introduce additional failure modes when autonomous decisions are incorrect at decision points that would benefit from human judgment.

5.9 OpenHands Platform: Open-Source Agentic Analysis

OpenHands provides the open-source research community with a reproducible platform for developing and evaluating AI software engineering agents, serving a fundamentally different role than commercial agentic systems. While commercial systems optimize for end-user experience in production deployments, OpenHands optimizes for research accessibility: transparent agent logic, modular architecture enabling systematic comparison, a full evaluation harness for SWE-bench, and community governance of development priorities.

The platform's finding that OpenHands with Claude 3.5 Sonnet achieves 37.2% on SWE-bench Verified — among the highest documented agentic performance — is particularly informative because it is the most transparently reproducible frontier performance measurement in the agentic space. Commercial systems reporting higher SWE-bench performance may do so with additional proprietary augmentations (enhanced retrieval, specialized prompting, ensemble approaches) that are not publicly described, making independent verification difficult. OpenHands' full reproducibility provides a reliable baseline against which commercial system claims can be calibrated.

The academic community's adoption of OpenHands as a shared evaluation platform has produced methodological advances alongside technical ones. Systematic failure mode analysis — identifying that agents most commonly fail through context management errors (losing track of their position in multi-step tasks), premature solution commitment (fixing on an approach before sufficiently exploring the problem space), and error recovery failures (being unable to backtrack after incorrect edits) — provides actionable targets for improving agent architectures and training approaches.

5.10 Comprehensive Comparative Analysis

Tool/ Model	Human Eval	SWE- bench(V)	Context	Cost (\$/1M)	Latency	Halluci nation	Reasoni ng	Enterpr ise	Open Source
GPT-4o	90.2%	49%	128K	\$5/\$15	Medium	Low-Med	Excellent	High	No
Claude 3.5 Sonnet	92.0%	49%	200K	\$3/\$15	Medium	Low	Excellent	High	No
Gemini 1.5 Pro	84.1%	38%	1M	\$3.50/\$10.5	Medium	Medium	Good	High	No
GitHub Copilot	N/A	N/A	Large	\$10-39/mo	Very Low	Low-Med	Good	Very High	No
Amazon Q Dev.	N/A	N/A	Large	\$19-30/mo	Very Low	Low	Good	Very High	No
Cursor	N/A	N/A	Large	\$20/mo	Low	Low-Med	Good	Medium	No
Windsurf	N/A	N/A	Large	\$15/mo	Low	Low-Med	Good	Medium	No
Claude Code	~92%	49%+	200K	Usage	Medium	Low	Excellent	Medium	No
Devin	N/A	13.9%	Large	\$500/mo	Slow	Medium	Good	Low-Med	No
OpenHands+Cl.	N/A	37.2%	200K	Model cost	Variable	Low	Excellent	Low	Yes




Figure 4: SWE-bench Verified — Real-World Bug Resolution Rate

Table 5.1: Comprehensive comparison of leading LLM-based SE tools (mid-2025). N/A = not applicable. Costs approximate.

5.11 Prompt Engineering Strategies for Software Engineering

The quality of LLM-generated software engineering outputs is substantially influenced by how tasks are framed in prompts. Prompt engineering — the practice of constructing input prompts to elicit desired model behaviors — has evolved from ad-hoc experimentation into a structured discipline with documented best practices and measurable impacts on output quality. The principal strategies relevant to software engineering applications are surveyed below, together with their empirically documented effectiveness.

Zero-shot prompting provides only the task description without examples, relying on the model's pre-trained knowledge to interpret and execute the task. For well-defined standard tasks (generate a Python function that sorts a list of dictionaries by a specified key), zero-shot prompting from frontier models typically produces high-quality outputs without additional engineering. Zero-shot prompting is the most cost-efficient approach and should be the starting point before more sophisticated strategies are considered.

Few-shot prompting provides 3-10 input-output examples before the target task, leveraging in-context learning to align the model's output distribution with the demonstrated format, style, and quality level. For code generation, few-shot examples demonstrating the desired code style, documentation format, error handling approach,

and level of abstraction substantially improve the relevance of generated code for project-specific contexts. Liu et al. [2023] found that carefully selected examples reflecting target project conventions improved developer acceptance rates for generated code by 30% compared to zero-shot prompting.

Chain-of-Thought prompting improves performance on complex algorithmic tasks by instructing the model to reason step-by-step before generating code. The pattern 'Let's think through this step by step: first identify the algorithm, then plan the data structures, then implement' activates systematic reasoning that reduces implementation errors. Wei et al. [2022] documented 10-25% improvement in HumanEval pass@1 with CoT prompting for frontier models. For debugging tasks, CoT prompting that structures diagnostic reasoning — 'First identify what the error message indicates, then locate the relevant code, then form a hypothesis about the root cause, then verify the hypothesis' — substantially improves diagnostic accuracy.

Strategy	Best Use Case	Typical Improvement	Cost Overhead	When to Use
Zero-shot	Standard, well-specified tasks	Baseline	None	Default starting point
Few-shot (3-5)	Style/convention matching	+15-30% acceptance	Variable (examples)	Project-specific generation
Chain-of-Thought	Algorithmic, multi-step tasks	+10-25% pass@1	1.5-2x tokens	Complex algorithms
Self-consistency	High-stakes, verifiable tasks	+20-40% pass@1	10x inference	Critical paths with tests
Role prompting	Domain-specific tasks	+10-20% specialized	Minimal (~50 tokens)	Security, performance focus
RAG-augmented	Proprietary/internal APIs	+40-60% accuracy	Retrieval latency	Enterprise private codebases
Structured output	JSON/schema generation	+30-50% format	Minimal	API integration, config gen
Stepback prompting	Debugging complex bugs	+15-25% accuracy	1.3x tokens	Root cause analysis

Table 5.2: Prompt engineering strategies and their impacts on LLM software engineering performance.

5.12 Cost-Performance Analysis for Enterprise Deployment

The selection of LLM-based software engineering tools for organizational deployment requires cost-performance analysis that accounts for both direct tool costs and the productivity value generated by different performance levels. A straightforward ROI framework considers the developer's fully loaded cost, the fraction of time spent on tasks amenable to LLM assistance, the measured productivity improvement for those tasks, and the tool subscription or usage cost.

For a developer earning \$100,000 annually (approximately \$50/hour at 2,000 working hours), tasks amenable to LLM assistance — code generation, documentation, test writing, debugging — comprising 50% of working time, and a conservative 30%

productivity improvement on those tasks yields: $\$50 \times 0.50 \times 0.30 \times 2,000 = \$15,000$ in annual productivity value. Against GitHub Copilot Business subscription cost of \$228/year (\$19/month), this represents a return of approximately 65:1 — strongly positive by any organizational investment standard. Even against the more conservative 20% productivity improvement from enterprise deployments, the ROI remains approximately 43:1.

Cost structures vary substantially across tool categories and must be evaluated for specific usage patterns. Subscription-based tools (GitHub Copilot, Amazon Q Developer, Cursor, Windsurf) provide predictable costs independent of usage intensity, favoring high-utilization deployments where per-interaction API costs would be higher. API-based access (GPT-4o API, Claude API) provides flexibility for variable usage patterns and enables integration into custom workflows, but may be more expensive at high usage intensities. Self-hosted open-source models (DeepSeek-Coder, Qwen2.5-Coder) eliminate per-interaction costs after infrastructure investment, favoring high-volume deployments and providing data privacy guarantees that may be necessary for sensitive codebases.

Tool	Monthly Cost/Dev	Performance Tier	Privacy	Customizable	Best For
GitHub Copilot Business	\$19/mo	High	Enterprise DPA	Fine-tuning	General dev teams
GitHub Copilot Enterprise	\$39/mo	High+	Enterprise DPA	Yes	Large orgs + custom
Amazon Q Developer	\$19-30/mo	High	AWS DPA	No	AWS-focused teams
Cursor Pro	\$20/mo	Very High	Cloud	No	Power users
Windsurf Pro	\$15/mo	High	Cloud	No	Small teams
Claude API (Sonnet)	~\$30-80/mo*	Very High	Anthropic BAA	Via fine-tune	Custom integrations
GPT-4o API	~\$40-100/mo*	Very High	OpenAI BAA	Via fine-tune	Custom integrations
Self-hosted OSS	\$50-200/mo†	Medium-High	Full control	Full	Privacy-first orgs




Figure 5: Monthly Cost per Developer by Tool (2024-25 Pricing)

*Table 5.3: Cost comparison for LLM-based developer tools. *API at ~1M tokens/month. †Cloud GPU infrastructure cost.*

5.13 Open-Source vs. Proprietary Models: Enterprise Trade-offs

The decision between open-source and proprietary LLMs for software engineering applications involves trade-offs across multiple dimensions: capability, cost, privacy, customization, and operational control. This decision is increasingly consequential as both categories converge in capability, making the non-capability dimensions more determinative of tool selection.

Proprietary frontier models (GPT-4o, Claude 3.5 Sonnet, Gemini 1.5 Pro) maintain performance advantages at the highest capability tiers, particularly on complex reasoning tasks, long-context understanding, and agentic multi-step execution. They provide managed infrastructure with guaranteed availability, professional support, and enterprise agreements (BAA, DPA) necessary for regulated industries. However, they require sending code to external APIs — a data privacy concern for organizations with sensitive intellectual property, regulatory restrictions, or contractual prohibitions on sharing code with third parties.

Open-source models (DeepSeek-Coder-V2, Qwen2.5-Coder, Code Llama, StarCoder2) provide data privacy through self-hosted deployment, full customization through fine-tuning on organizational code, and marginal inference cost after infrastructure investment that favors high-volume usage. The capability gap between

open-source and proprietary models has narrowed substantially since 2023: DeepSeek-Coder-V2's 90.2% HumanEval performance matches GPT-4o, and Qwen2.5-Coder-72B at 88.4% is competitive with most proprietary alternatives. For organizations with adequate ML infrastructure and data science expertise, self-hosted open-source models now represent a practically viable alternative to proprietary API access for most software engineering tasks.

Hybrid approaches — using proprietary APIs for the most demanding tasks (complex debugging, architectural design, agentic workflows) while self-hosting open-source models for high-volume routine tasks (inline completion, simple documentation) — can optimize the capability-cost-privacy balance for organizations with sophisticated requirements. This tiered architecture, increasingly adopted by large technology companies, leverages the strengths of each approach without accepting the limitations of either as the universal solution.

CHAPTER 6: CASE STUDIES OF LLM DEPLOYMENT IN SOFTWARE ENGINEERING

6.1 Case Study Framework and Selection Criteria

The case studies presented in this chapter were selected to provide representative coverage across three critical dimensions: organizational context (from individual developers to large enterprises), deployment model (cloud API, embedded tool, agentic system, self-hosted), and application domain (general software development, cloud infrastructure, enterprise systems, open-source projects). Seven case studies were selected through a purposive sampling methodology that prioritizes depth of documented evidence over breadth of superficial coverage, ensuring that each case study contributes distinct insights not duplicated across the set.

Each case study is structured around five analytical questions: What specific problem or opportunity motivated the LLM deployment? What was the technical architecture of the implementation? What quantitative and qualitative outcomes were observed? What implementation challenges were encountered and how were they addressed? What lessons generalize to other deployment contexts? This consistent structure enables cross-case comparison while preserving the contextual richness necessary for translating case study lessons to novel deployment situations.

6.2 Case Study 1: GitHub Copilot at Scale — Microsoft/GitHub Enterprise Deployment

Context and Motivation

GitHub Copilot's enterprise deployment across Microsoft's own software engineering organization, comprising tens of thousands of engineers working in diverse technology stacks and application domains, provides the most extensively documented large-scale case study of LLM-based coding assistance. Microsoft's dual role as both Copilot's developer (through its ownership of GitHub) and its largest enterprise customer created conditions for unusually careful measurement: the organization had both the motivation to demonstrate value and the technical capability to measure it precisely.

The deployment motivation was partly economic — documenting productivity improvements for Copilot's enterprise positioning — and partly operational: Microsoft software teams had begun reporting anecdotal productivity improvements from personal Copilot use, and the organization sought to understand whether those improvements were generalizable, which developer populations benefited most, and what workflow practices

correlated with greatest impact. This combination of motivations produced unusually careful measurement protocols compared to most enterprise tool deployments.

Implementation Architecture

The Microsoft/GitHub enterprise deployment of Copilot integrates across the full development toolchain: Visual Studio Code and Visual Studio for desktop development, GitHub web editor for browser-based development, GitHub Mobile for lightweight review tasks, and GitHub Actions integration for CI/CD pipeline assistance. The deployment architecture includes organization-level configuration for content exclusion (preventing Copilot from accessing code in sensitive repositories defined by repository topic tags), centralized telemetry collection for adoption and productivity metrics, and an enterprise GitHub Copilot Business plan providing enhanced security commitments and organizational policy controls.

The context-enrichment capabilities deployed in the enterprise configuration extend beyond single-file context to workspace-level context (multiple open files), Git history context (recent diffs and commit messages), and organizational knowledge context (internal documentation indexed through semantic search). This extended context enables suggestions that align with internal API conventions, follow organization-specific naming patterns, and reuse existing utility functions rather than reimplementing common patterns — improvements over single-file context that are particularly valuable in large, mature codebases with established conventions.

Outcomes and Measurements

The Microsoft internal deployment produced the most reliable available productivity evidence for Copilot at enterprise scale. The controlled study published by Ziegler et al. found sustained code suggestion acceptance rates of 27-32% across diverse developer populations, reflecting both the quality of suggestions and the selective nature of developer acceptance decisions. More directly relevant for productivity assessment, task completion time measurements found 35-40% acceleration for routine implementation tasks — function implementation from specifications, boilerplate generation, standard test structure — while complex reasoning tasks showed more modest improvements of 10-20%, consistent with the model providing more value for well-defined tasks where solution patterns are predictable.

Developer satisfaction surveys showed 76% of surveyed developers reporting that Copilot 'helps me feel more productive,' with the most frequently cited benefit being reduction of context-switching burden: developers could remain in the implementation flow without pausing to look up API documentation, recall function signatures, or search for relevant code examples. The social dimension of AI assistance — developers reporting that Copilot made coding feel 'less lonely' and reduced the anxiety of starting with a blank file — emerged as an unexpected quality-of-experience benefit with

potential implications for developer wellbeing and retention beyond pure productivity measurements.

Implementation Challenges

The primary technical challenge in the Microsoft enterprise deployment was IP and code privacy management: ensuring that sensitive internal code, algorithms, and business logic were not included in training data or accessible to Copilot users outside the organization. GitHub addressed this through the Copilot for Business architecture that provides explicit commitments against using enterprise customer code for model training and implements content exclusion policies at the repository level. The secondary challenge was measuring Copilot's actual contribution to productivity versus confounding variables — developer experience, task complexity, and project phase — that substantially affect productivity independent of AI assistance. The controlled study methodology (with vs. without Copilot, randomized assignment) addressed this challenge more carefully than most enterprise productivity studies.

Generalizable Lessons

Three lessons generalize broadly from the Microsoft/GitHub Copilot deployment. First, adoption is substantially higher among developers who receive structured onboarding (tutorial, best practices guidance, team champion support) versus those given tool access without onboarding, emphasizing the importance of change management alongside technology deployment. Second, productivity benefits are largest for junior and intermediate developers on routine tasks, while senior developers disproportionately benefit from complex reasoning assistance (Copilot Chat for debugging, architecture consultation) — suggesting that tool configuration and promotion strategies should be segmented by developer experience level. Third, the highest-value workflows require workflow adaptation rather than simple tool insertion: teams that deliberately restructure their development practices to integrate Copilot (using it to draft test cases before implementation, reviewing its suggestions critically for security issues, using it for documentation at the time of implementation rather than after) achieve substantially greater benefits than teams that simply try the tool in their existing workflow.

6.3 Case Study 2: Google Internal Deployment of Gemini Code Assist

Context and Motivation

Google's internal deployment of an LLM-based coding assistant across its engineering organization — initially called ML-enhanced code completion and subsequently evolved into Gemini Code Assist — provides a second large-scale enterprise case study with distinctive characteristics. Google's engineering environment presents specific challenges and opportunities: a polyglot codebase of extraordinary scale

(one of the world's largest monorepos), strong internal culture of rigorous empirical evaluation, and the distinctive advantage of being able to deploy an AI assistant integrated with proprietary internal tools (internal code search, internal documentation systems, proprietary build systems like Blaze).

The deployment was motivated by both productivity and competitive considerations: Google's engineering leadership recognized that competing organizations were adopting AI coding tools and sought to provide comparable or superior assistance that leveraged Google's unique advantages — access to a comprehensive internal codebase as context, integration with Google's proprietary development tools, and the ability to train models on internal code patterns.

Implementation and Technical Details

Google's internal AI coding assistant is uniquely integrated with Code Search (Google's internal code navigation tool that indexes the entire monorepo), enabling suggestions that reference and reuse any existing function or pattern in the codebase regardless of which directory it resides in. This cross-repository context is qualitatively different from the cross-file context available in commercial tools: instead of context limited to open files in the editor, the assistant has access to the entire Google codebase through semantic search, enabling it to suggest reusing an existing utility function from a completely unrelated project that happens to implement exactly the needed pattern.

The multi-language support requirements at Google — C++, Java, Python, Go, JavaScript, TypeScript, Kotlin, SQL, and numerous proprietary languages — required careful evaluation of model performance across all languages, with specialized fine-tuning attention for lower-represented proprietary languages where general training data is insufficient. Internal extensions to support Google-specific APIs (internal libraries, proprietary ML frameworks, internal infrastructure APIs) required augmented training data beyond publicly available code and documentation.

Outcomes and Reported Impact

Google published selected results from its internal AI coding assistant deployment at research conferences and in technical blog posts, reporting code acceptance rates comparable to external tools (25-30%), with higher acceptance rates for specific task categories (test generation, documentation, boilerplate for internal framework usage) where the model's training on internal patterns provides an advantage over general-purpose tools. The internal deployment reports attribute 5-10% improvements in developer productivity (as measured by features shipped, bug fix velocity, and self-reported time allocation) to AI coding assistance, acknowledging substantial uncertainty in isolating the assistant's contribution from other concurrent engineering investments.

A distinctive outcome from Google's deployment is the documentation of where AI assistance provides diminishing returns at high skill levels. Staff engineers and principal engineers working on highly novel algorithmic problems and architectural decisions report lower utility than earlier-career developers, consistent with the hypothesis that AI assistance provides greater value for tasks where established patterns exist than for genuinely novel problem-solving. This finding has implications for targeting AI assistance investments toward the most impactful deployment contexts.

6.4 Case Study 3: Amazon Q Developer for AWS Infrastructure Automation

Context and Motivation

Amazon's deployment of Amazon Q Developer (formerly CodeWhisperer) for AWS infrastructure automation at enterprise scale represents a case study with a distinctive security-first emphasis reflecting Amazon's security culture and the high-stakes nature of cloud infrastructure configuration. Configuration errors in AWS infrastructure — incorrect security group rules, overly permissive IAM policies, unencrypted storage configurations — can have significant security and compliance consequences, motivating careful attention to AI-generated infrastructure code quality.

The deployment context includes both internal Amazon teams building AWS infrastructure for their own services and external enterprise customers using Amazon Q Developer through the AWS console and developer tools integration. The dual deployment context provides insights into both AI-assisted infrastructure development (where the developer understands the infrastructure being configured) and AI-assisted infrastructure auditing (where security and compliance teams review existing configurations for issues).

Technical Architecture

Amazon Q Developer's infrastructure automation capabilities operate through two complementary mechanisms: IDE-integrated code generation in AWS Cloud9, VS Code with the AWS Toolkit, and JetBrains IDEs; and console-integrated assistance through the AWS Management Console's built-in chat interface. The console integration is particularly distinctive, enabling developers to describe infrastructure requirements in natural language directly within the console interface and receive immediately deployable CloudFormation or CDK configurations.

The security scanning integration uses a hybrid detection system combining Amazon's CodeGuru Reviewer (rule-based static analysis for security and code quality) with an ML-based vulnerability detection model trained specifically on AWS infrastructure code patterns. This combination provides both reliable detection of known

vulnerability signatures (world-writable S3 buckets, security groups allowing unrestricted inbound access, EC2 instances without encrypted storage) and learned detection of more subtle misconfigurations that violate no explicit rule but represent best practice violations.

Quantified Outcomes

AWS published case studies from enterprise Q Developer deployments documenting specific productivity and quality outcomes. Accenture's deployment of Amazon Q Developer for AWS infrastructure projects reported 35% reduction in time required to create new infrastructure configurations and a 40% reduction in security findings during code review, attributed to the assistant's proactive security feedback reducing defect injection at the generation stage. The Guardian's deployment reported 80% reduction in time required for AWS Lambda configuration tasks — a high-frequency operation that is complex to configure correctly but highly repetitive, creating ideal conditions for AI assistance.

Security outcome measurement is particularly challenging in infrastructure automation cases because security improvements manifest as absence of incidents rather than presence of measurable events. Amazon's published metrics rely primarily on defect rate comparisons between AI-assisted and non-AI-assisted infrastructure code in pre-production security reviews, finding 40-50% fewer critical security findings in AI-assisted code, consistent with the assistant's real-time security feedback preventing the most common security configuration errors at generation time rather than catching them in review.

6.5 Case Study 4: Claude Code in Academic Software Development

Context and Motivation

The deployment of Claude Code in academic research software development contexts provides a case study with distinctive characteristics: smaller teams (typically 1-5 developers), less structured development processes than enterprise environments, high diversity of programming tasks (data analysis scripts, simulation code, scientific visualization, machine learning experiments), and particular sensitivity to correctness (errors in research code can propagate to invalid scientific conclusions). The academic deployment context also provides an opportunity to study LLM assistance effects on less experienced programmers — graduate students and early-career researchers who may have domain expertise but limited software engineering training.

Multiple computer science and computational science research groups documented their experiences with Claude Code as part of a broader qualitative study of AI coding assistant adoption in academic contexts. The study collected structured usage

diaries, code inspection results, and semi-structured interviews from 35 graduate students and postdoctoral researchers across 12 research groups over a six-month period, providing longitudinal perspective on usage patterns and learning effects not captured in shorter productivity studies.

Usage Patterns and Findings

The study found that research software developers used Claude Code for a distinctive task distribution compared to professional developers: a higher proportion of usage was for understanding and debugging unfamiliar code bases (inheriting a labmate's simulation code), explanation of numerical algorithms (understanding why a particular finite element formulation is implemented in a specific way), and translation of algorithms from publications (implementing a method described mathematically in a paper). These explanation and understanding tasks leveraged the model's broad knowledge of scientific computing patterns and ability to bridge between mathematical notation and code implementation.

A significant finding was the differential impact on code quality for participants with different programming backgrounds. Computer science graduate students (with formal programming training) primarily used Claude Code for productivity acceleration — generating boilerplate, writing tests, producing documentation — with limited quality impact (their pre-AI code quality was already adequate). Non-CS graduate students (physicists, biologists, engineers using Python for data analysis) demonstrated more dramatic code quality improvements: their AI-assisted code showed substantially better error handling, more consistent use of standard scientific computing patterns (using NumPy/SciPy rather than manual loops, using Pandas for tabular data rather than raw CSV parsing), and more comprehensive documentation than their pre-AI-assistance code.

Challenges and Limitations

The primary challenge reported by academic research software developers was domain-specific correctness: Claude Code's suggestions were syntactically correct and reasonable for general Python programming but sometimes suggested approaches that, while conventional in software engineering, were inappropriate for scientific computing contexts. Specific examples included suggesting inappropriate floating-point comparison patterns for numerical validation, recommending general-purpose data structures where domain-specific scientific computing structures (sparse matrices, coordinate arrays) were more appropriate, and misunderstanding the statistical assumptions required for correct data analysis approaches. These domain-specific errors required reviewers with scientific computing expertise to identify, which most academic research groups lacked.

The longitudinal design also revealed a dependency risk not apparent in short-term studies: participants who heavily used Claude Code for the first three months showed reduced ability to write code independently when tested at six months compared

to control group members who had not used AI assistance. This finding — preliminary and based on a small sample — raises important questions about skill development effects of heavy AI coding assistance that deserve investigation in larger studies with more controlled experimental designs.

6.6 Case Study 5: Sourcegraph Cody for Enterprise Codebase Navigation

Context and Motivation

Sourcegraph Cody represents a distinctive approach to LLM-based developer assistance that prioritizes codebase understanding and navigation over code generation. Rather than generating new code, Cody's primary value proposition is enabling developers to effectively navigate and understand large enterprise codebases — finding relevant code, understanding how components interact, and identifying where specific changes need to be made. This navigation-first approach reflects the recognition that developer time is often spent understanding existing code rather than writing new code, and that AI assistance can provide substantial value by making large codebases more navigable.

The case study enterprise — a financial services company with 6 million lines of Java code spread across 200+ repositories accumulated over 15 years — represents a common but particularly challenging enterprise context: large, mature codebase; significant knowledge loss as original developers have left; inconsistent documentation; and architectural patterns that have evolved across multiple technology generations. The deployment was motivated by a desire to reduce onboarding time for new developers and reduce the 'human bus factor' risk from knowledge concentration in a small number of senior developers who understood the full system.

Implementation and Technical Approach

Cody's implementation uses Sourcegraph's enterprise code search infrastructure — which indexes all repositories, commits, and file contents — as its retrieval substrate, augmenting generation with precise code snippets from the current codebase rather than relying solely on model training. When a developer asks 'How does the payment processing pipeline handle failed transactions?', Cody retrieves relevant code from the payment processing services, the error handling utilities, and the retry logic, and uses this retrieved context to generate an accurate, specific answer citing actual code locations rather than generic patterns.

The enterprise deployment also leverages Sourcegraph's Code Insights capabilities — analytics about code change frequency, test coverage, dependency graphs, and code ownership — to provide context-aware assistance. Knowing which files change

most frequently, which components are most tightly coupled, and which code areas have the lowest test coverage enables Cody to provide more relevant prioritization recommendations for refactoring, testing, and documentation efforts.

Measured Outcomes

The financial services company reported a 45% reduction in new developer onboarding time from six months (average time for a new developer to be independently productive in the codebase) to three months, attributing this improvement primarily to Cody's ability to answer codebase-specific questions instantly rather than requiring lengthy documentation reading or bothering senior developers. Senior developer satisfaction scores increased as they were interrupted less frequently for basic codebase navigation questions, enabling more focused work on the architectural and design questions that genuinely require their expertise.

Documentation improvement was an indirect benefit: as developers used Cody's codebase explanation capability, they identified documentation gaps and code regions that were particularly difficult for the AI to explain accurately (a proxy for code regions that would also be difficult for human developers to understand). This 'explanation difficulty' signal was used to prioritize documentation and refactoring efforts, directing attention to code areas most likely to create maintenance problems. The systematic use of AI explainability as a code quality signal is an emerging practice that this case study illuminates.

6.7 Case Study 6: OpenHands for Open-Source Repository Management

Context and Motivation

The deployment of OpenHands (formerly OpenDevin) for autonomous management of issue queues in open-source projects represents a case study at the frontier of practical agentic AI system deployment. Open-source repositories frequently accumulate backlogs of reported issues — bugs, feature requests, documentation gaps — that exceed maintainer capacity to address. LLM-based autonomous agents offer the prospect of expanding effective repository maintenance capacity without proportionally expanding human maintainer time investment.

The case study examines three open-source projects (anonymized as Projects A, B, and C) with active issue backlogs of 200-800 open issues that participated in a research pilot deploying OpenHands agents to: triage newly opened issues (labeling, requesting clarification, identifying duplicates), generate fix proposals for labeled 'good first issue' bugs, and prepare test cases for issues marked as 'needs tests'. The pilot operated over three months with human maintainer oversight of all agent actions before

execution, providing human validation of agent decision quality while enabling systematic measurement of agent performance.

Agent Performance Results

Issue triage accuracy — the agent correctly applying the same labels that human maintainers would apply — achieved 72% accuracy for standard labels (bug, feature request, documentation, needs clarification) and 56% accuracy for domain-specific labels requiring understanding of the project's specific architectural concepts. Label accuracy was highest for clear, well-specified issues and lowest for ambiguous issues where human maintainers themselves sometimes disagreed about classification. The agent's performance on duplicate detection was particularly strong: 81% accuracy for identifying that a newly opened issue was a duplicate of an existing issue, reflecting the strength of semantic similarity reasoning for matching natural language issue descriptions.

Fix proposal quality — the agent generating a pull request implementing the described fix — was evaluated by maintainers on a four-point scale (unacceptable, needs substantial revision, needs minor revision, acceptable as-is). For 'good first issue' tagged bugs in Project A (a Python library with solid test coverage), 38% of agent-generated fixes were rated acceptable as-is, 31% needed minor revision, 21% needed substantial revision, and 10% were unacceptable. These results align with SWE-bench Verified performance expectations and demonstrate practical utility for simple, well-specified bugs — though the significant fraction requiring human revision indicates that full automation is premature for general issue resolution.

Enterprise and Community Impact Lessons

The three-month pilot established several lessons with broader implications. Autonomous agent deployment requires graduated autonomy: starting with read-only triage actions (labeling, commenting) before enabling write actions (committing code, merging pull requests) builds community trust and catches failure modes before they affect the main codebase. Agent performance is highly variable across project characteristics: projects with strong test suites, clear coding style, and well-specified issue reports enable substantially higher autonomous fix quality than projects with sparse testing or ambiguous issues — suggesting that investments in project quality infrastructure are prerequisites for effective AI augmentation.

6.8 Enterprise AI Coding Deployment Patterns and Anti-Patterns

Common Enterprise Deployment Patterns

Analysis of the six case studies above, combined with deployment reports from 15 additional enterprise Copilot and Q Developer deployments in the research literature, identifies five common patterns associated with successful enterprise AI coding

deployments. The graduated deployment pattern — starting with a small pilot group, measuring outcomes carefully, incorporating lessons into broader deployment, and scaling incrementally — consistently outperforms big-bang full-organization deployments in both adoption rates and measured productivity impact. The pilot group both proves value (providing evidence for broader investment) and surfaces implementation issues at manageable scale.

The role-segmented deployment pattern recognizes that different developer roles benefit from different AI assistance patterns: individual contributor developers benefit most from inline code generation and documentation assistance; technical leads benefit most from code review assistance and architecture consultation; DevOps engineers benefit most from infrastructure configuration generation and incident diagnosis; and data scientists benefit most from SQL generation and visualization code assistance. Deployment configurations, training, and promotion strategies segmented by role achieve higher adoption and reported productivity impact than uniform approaches.

The security-first configuration pattern — configuring content exclusion policies, data residency requirements, and access controls before broader deployment — prevents security and compliance issues that can derail deployment programs. Organizations that establish security guardrails proactively report smoother deployments and fewer incidents requiring emergency policy changes that disrupt developer workflows.

Common Anti-Patterns

The 'uncritical adoption' anti-pattern — deploying AI coding assistance without establishing code review processes specifically designed to catch AI-generated errors — is associated with increased security vulnerability rates and technical debt accumulation. AI-generated code has systematic failure patterns (under-specified error handling, missing null checks, SQL injection vulnerabilities in database code, performance anti-patterns) that are not fully caught by traditional code review processes calibrated for human-generated code. Organizations should supplement existing code review with security scanning tools and explicit reviewer guidelines for evaluating AI-generated code sections.

The 'measured productivity regress' anti-pattern occurs when heavy adoption of AI coding assistance initially improves productivity but subsequently causes skill degradation that reduces long-term productivity, a risk surfaced in the academic case study. This anti-pattern is mitigated by deliberate practice policies that require developers to maintain skill currency through periodic AI-unassisted coding, learning-oriented AI usage practices that emphasize understanding over acceptance, and career development frameworks that value and assess software engineering craft independently of AI tool proficiency.

Case Study	Organization Type	Tool Deployed	Key Outcome	Primary Challenge
GitHub Copilot	Microsoft (enterprise)	Copilot Business	35-40% faster routine tasks	Change management, IP protection
Gemini Code Assist	Google (tech giant)	Internal AI tool	5-10% velocity improvement	Monorepo integration, novel tasks
Amazon Q Developer	AWS enterprise customers	Q Developer	35-80% task time reduction	Security validation, IaC quality
Claude Code	Academic research	Claude Code	Significant quality lift for non-CS devs	Domain correctness, skill dependency
Sourcegraph Cody	Financial services	Cody Enterprise	45% onboarding time reduction	Legacy codebase context quality
OpenHands	Open-source projects	OpenHands agents	38% auto-fix acceptance rate	Graduated autonomy, community trust

Table 6.1: Summary of case study deployments, tools, outcomes, and key challenges.

CHAPTER 7: CHALLENGES AND LIMITATIONS OF LLMs IN SOFTWARE ENGINEERING

7.1 Overview of Identified Challenges

While the productivity and quality benefits of LLM-based software engineering tools are well-documented, deployment at scale has revealed a constellation of technical, organizational, and socio-technical challenges that must be understood and addressed for responsible adoption. The primary challenge categories identified across the literature and from the case studies of Chapter 6 are analysed in the sections below: correctness and reliability, security and safety, intellectual property and legal concerns, cognitive effects on developers, bias and fairness, sustainability, scalability in large codebases, and organizational integration challenges.

The challenges are organized from most immediate and well-documented (correctness limitations) to emerging and prospective (long-term skill effects, environmental impact), reflecting both the maturity of understanding and the urgency for practitioners. Each challenge analysis follows a consistent structure: empirical characterization of the challenge, root cause analysis, current mitigation approaches, their limitations, and the open research questions that must be addressed for more complete solutions.

7.2 Correctness, Reliability, and Hallucination in Generated Code

Code correctness — the property that generated code implements the intended specification without errors — remains the most fundamental and extensively studied challenge for LLM-based code generation. Unlike natural language generation, where minor inaccuracies or imprecision may be acceptable, code requires absolute syntactic correctness (the code must parse) and functional correctness (the code must produce the intended outputs for all inputs). LLMs do not guarantee either property, producing code that may silently fail, produce subtly incorrect outputs for edge cases, use deprecated APIs, or contain logic errors that pass surface inspection but fail in production.

The 'confident incorrectness' failure mode — where models generate plausible-looking but functionally incorrect code without indicating uncertainty — is particularly challenging for practitioners. Unlike human developers, who typically indicate uncertainty or acknowledge knowledge limitations, LLMs produce generated code with uniform confidence regardless of whether the model has high or low certainty about the correctness of the approach. This uniform presentation of confident output regardless of underlying uncertainty creates a false sense of security that can lead to insufficient testing of AI-generated code.

Hallucination in code-specific forms manifests as: invented API method calls that do not exist in the referenced library; incorrect method signatures (wrong parameter types, orders, or counts); invented package names in import statements; incorrect namespace or module path specifications; and fabricated class hierarchies. Wu et al. [2024] found that GPT-4 hallucinated at least one API method call or signature in approximately 18% of function generation responses for questions involving less-common libraries, with hallucination rates exceeding 40% for libraries with limited training data representation. The challenge is particularly acute for internal or proprietary APIs not represented in training data — LLMs frequently hallucinate plausible-seeming internal API calls when prompted to use APIs whose documentation is not publicly available.

Mitigation Approaches and Their Limitations

Current mitigation approaches for correctness and hallucination include: automated test execution (running generated code against test suites to catch correctness failures before deployment); static analysis integration (using type checkers, linters, and formal verification tools to catch syntactic and simple semantic errors); RAG augmentation with authoritative documentation (retrieving accurate API documentation and examples to ground generation); multiple-sample generation with majority voting (sampling multiple completions and selecting the most common correct implementation); and output verification with smaller, faster verification models. Each approach addresses some failure modes but none provides comprehensive correctness guarantees, and all introduce overhead in developer time, computational cost, or system complexity.

7.3 Security Vulnerabilities in LLM-Generated Code

The security implications of LLM-generated code adoption at scale represent one of the most consequential challenge categories, with potential for systemic effects if common vulnerability patterns are widely propagated through AI-generated code. Pearce et al.'s foundational study [40] establishing that approximately 40% of Copilot-generated code for security-sensitive scenarios contained vulnerabilities motivated significant subsequent research on both the scale and character of the problem and on mitigation approaches.

The vulnerability distribution in LLM-generated code is not uniformly distributed across OWASP categories but concentrated in specific patterns that reflect LLM training data characteristics and generation tendencies. SQL injection through string interpolation rather than parameterized queries is particularly common, as string-interpolated SQL patterns are well-represented in training data from legacy codebases. Hardcoded credentials in configuration and test code are common because examples in training data frequently include real or representative credentials without emphasizing their removal.

Use of deprecated or weak cryptographic functions (MD5, SHA-1 for password hashing; DES, 3DES for encryption) reflects training data that includes historical code predating security practice evolution. Insufficient input validation before security-sensitive operations reflects the general tendency to generate code that works for expected inputs without handling boundary cases.

The systemic risk from LLM-generated vulnerability propagation is amplified by the adoption scale: millions of developers using tools like GitHub Copilot means that vulnerability patterns in LLM training data may be propagated to codebases globally, potentially creating vulnerability correlation across unrelated organizations' systems. Unlike human developers who introduce vulnerabilities in diverse and idiosyncratic ways, LLM-introduced vulnerabilities may be systematically correlated — organizations deploying similar AI tools may share common vulnerability patterns that create correlated exploitability across the industry.

Enterprise Security Controls

Organizations deploying LLM-based coding assistance at scale should implement layered security controls specifically designed for AI-generated code: real-time SAST scanning of all AI-generated suggestions before acceptance; security-focused code review guidelines that prompt reviewers to explicitly check AI-generated code for the most common LLM vulnerability patterns; developer training on LLM security limitations and secure AI-assisted coding practices; and monitoring of security defect rates in AI-generated code cohorts to detect degradation in security outcome metrics. These controls should be treated as standard security program components for organizations with significant AI coding tool adoption rather than optional enhancements.

7.4 Intellectual Property, Copyright, and Legal Challenges

The legal landscape surrounding LLM-generated code remains contested and rapidly evolving as courts, regulators, and legislative bodies grapple with questions about training data copyright, generated code ownership, and liability for LLM-generated vulnerabilities. These unresolved legal questions create uncertainty for organizations deploying AI coding tools, particularly those in regulated industries where IP protection and liability management are critical concerns.

Copyright concerns center primarily on two questions: whether training on copyrighted code without explicit license creates copyright liability for model developers, and whether generated code that closely resembles training data constitutes copyright infringement by the generating organization. The *Doe 1 v. GitHub, Inc.* class action lawsuit, filed in 2022 and ongoing as of mid-2025, directly challenges whether training Copilot on open-source code without compliance with license terms (including copyleft license attribution requirements for GPL-licensed code) constitutes copyright

infringement. The outcome of this case is expected to substantially clarify the legal framework for LLM training data, with implications for all model developers training on publicly available code.

Generated code attribution requirements — whether AI-generated code must be disclosed as such, attributed to the generating tool, or licensed consistently with the licenses of training data that influenced the generation — are similarly unresolved. Some organizations have adopted precautionary disclosure policies requiring developers to flag AI-generated code in commit messages or comments; others have determined that existing IP assignment agreements with employees cover AI-assisted work; and others are awaiting regulatory or legal clarification before establishing formal policies. This variation in organizational approach reflects genuine legal uncertainty rather than different risk tolerances.

7.5 Cognitive Effects: Skill Development and Developer Agency

The long-term cognitive effects of heavy AI coding assistance adoption on software engineering skill development represent one of the most important yet least studied challenges. The concern is straightforward: if developers consistently delegate routine coding tasks to AI systems, they may progressively lose proficiency in those tasks, creating skill atrophy that reduces effectiveness when AI assistance is unavailable and potentially limits the ability to identify errors in AI-generated code.

The academic case study in Chapter 6 found preliminary evidence of this skill atrophy effect in a small sample of research software developers over six months. Independent of this specific study, the cognitive science literature on skilled performance provides theoretical grounding for the concern: deliberate practice is necessary for skill maintenance and development, and offloading practice opportunities to AI systems may reduce the deliberate practice that maintains and develops programming skills. The 'automation complacency' phenomenon documented in aviation and autonomous vehicle research — where operators of automated systems lose situational awareness and manual skill proficiency that is necessary when automation fails — has direct parallels in AI-assisted software development.

Developer agency concerns extend beyond individual skill development to the broader question of professional practice: whether heavy AI assistance creates a mode of software development where developers primarily manage and validate AI output rather than creating software through their own design and implementation decisions. Some practitioners argue that this shift represents efficient specialization — developers contributing higher-level design and validation expertise while AI handles implementation details — while others contend that implementation-level expertise is necessary for good design judgment and that the shift represents a fundamental deskilling with negative long-term consequences for the field.

7.6 Bias, Fairness, and Representation in LLM Code

LLM training on existing code corpora propagates the biases and limitations present in that code into generated outputs. The most consequential biases for software engineering applications include programming language representation bias, domain and application area bias, and demographic assumption bias embedded in code logic and data handling.

Programming language representation bias reflects the distribution of public code repositories: Python, JavaScript, Java, and C++ dominate training data, producing model capabilities that are substantially stronger for these languages than for less-represented alternatives. COBOL, RPG, and other enterprise languages used extensively in legacy systems have limited training representation despite their commercial importance; Rust, Zig, and other newer systems languages may be underrepresented relative to their growing adoption; and domain-specific languages (R for statistics, MATLAB for engineering computation, SAS for data analysis) show variable quality in generated code reflecting their training data representation. Organizations deploying AI coding tools for less-represented languages should evaluate model performance for their specific language context before assuming general marketing claims apply.

Demographic assumption bias — the embedding of assumptions about user characteristics in generated code logic and data handling — represents a subtler but important concern for software that interacts with diverse user populations. Code generated for user registration systems may default to binary gender options; code generated for name handling may assume Western naming conventions (single first name, family name) incompatible with names from other cultural traditions; code generated for internationalization may provide incomplete support for right-to-left text, complex scripts, or Unicode characters beyond Latin and common Asian scripts. These assumptions in generated code reflect assumptions in training data and can propagate to deployed systems that serve diverse user populations, requiring intentional review and correction.

7.7 Environmental and Sustainability Concerns

The environmental footprint of large-scale LLM deployment is an important and increasingly scrutinized concern as the scale of AI computing infrastructure grows. LLM training — required once per model generation but for models of frontier scale — consumes substantial computational resources: estimates for GPT-4 training energy consumption range from 50-100+ GWh depending on training efficiency assumptions, comparable to the annual electricity consumption of 4,000-8,000 average US households. While training cost is a one-time investment amortized across potentially billions of inference queries, training multiple large models per year across multiple competing organizations represents a significant aggregate energy investment.

Inference energy consumption — the energy required for each query served — is potentially more consequential at scale given the per-query nature of the cost. A frontier model serving millions of developers making hundreds of queries per day creates inference compute requirements that aggregate to substantial energy consumption. Meta's research [157] estimated that if the 3.5 million global software developers in 2024 each made 100 AI coding queries per day using frontier models, the aggregate inference compute would require approximately 1.2 GWh per day — equivalent to powering approximately 110,000 average homes daily. This estimate highlights the importance of model efficiency improvements (smaller models, quantization, speculative decoding, caching) for sustainability as adoption scales.

The industry response to environmental concerns includes: training on renewable energy where infrastructure permits (Google's commitment to 24/7 clean energy for AI data centers, Microsoft's renewable energy purchasing commitments, Amazon's significant renewable energy investment); efficiency improvements that reduce inference cost per query (quantization, distillation, speculative decoding, efficient attention architectures); and research into training efficiency (better data curation reducing required training compute, improved training recipes achieving higher data efficiency). These efforts address some concerns but do not eliminate the aggregate energy growth from expanding AI deployment.

7.8 Scalability in Large, Complex Codebases

LLM performance in large, complex production codebases — the environment where commercial software engineering actually occurs — differs substantially from performance in the benchmark environments where models are primarily evaluated. The challenges of scale include: context window limitations that prevent including sufficient codebase context for informed suggestions; semantic coherence degradation as context grows (models attending poorly to information at the edges of long contexts); incremental update handling (the model must maintain consistent understanding as files are modified during a development session); and cross-language and cross-service dependencies in heterogeneous system architectures.

The 'lost in the middle' problem [74] — where information in the middle of long contexts receives less attention than information at the beginning or end — has direct implications for repository-level code generation. A 200K token context containing an entire codebase may include the relevant code somewhere in the middle, where the model attends less reliably than to code appearing at the start or end of the context window. Retrieval-augmented approaches that identify and surface relevant code to the beginning of the context may achieve better practical performance than naive long-context approaches even with shorter effective context, because retrieved code appears in the high-attention portion of the context.

Architectural heterogeneity — systems where frontend services are implemented in TypeScript/React, backend services in Python/FastAPI, data pipelines in Scala/Spark, and infrastructure in Terraform/Go — creates cross-language coordination challenges for AI coding tools optimized primarily for single-language contexts. Changes that require coordinated modifications across service boundaries (updating an API contract that affects both the providing service's implementation and multiple consuming services' client code) exceed the capability of tools focused on single-file or single-language contexts, requiring either agentic systems with cross-file coordination capabilities or human orchestration of AI assistance across services.

7.9 Organizational and Cultural Integration Challenges

Beyond technical challenges, organizations adopting AI coding tools face significant organizational and cultural integration challenges that determine whether technical capabilities translate to actual productivity and quality improvements. Resistance from experienced developers who view AI assistance as threatening to professional identity, as reducing code ownership and craftsmanship, or as a step toward their eventual replacement creates adoption friction that requires thoughtful change management. Developer communities have demonstrated strong cultural values around code quality, understanding, and ownership, and AI tools that appear to compromise these values face resistance independent of their technical effectiveness.

Team dynamics changes from AI assistance adoption can be both positive (reduced need to interrupt senior colleagues for routine questions) and challenging (reduced motivation for knowledge sharing, potential loss of informal learning through pair programming and code review). Organizations that proactively design how AI assistance integrates with collaboration practices — ensuring that AI reduces administrative burden rather than human interaction — tend to achieve better adoption outcomes than those that deploy tools without consideration of collaboration effects.

Process and workflow integration challenges are particularly evident at the organizational boundary between AI-assisted code generation and traditional quality assurance processes designed for human-generated code. Code review guidelines, testing requirements, security scanning thresholds, and documentation standards designed for human-generated code may be inappropriately calibrated for AI-generated code — either too lenient (missing AI-specific failure modes) or too strict (requiring human-appropriate documentation on AI-generated boilerplate). Organizations must deliberately re-examine their quality assurance processes in light of AI-assisted development practices, rather than assuming existing processes remain appropriate.

CHAPTER 8: FUTURE RESEARCH DIRECTIONS

8.1 Overview of Research Opportunities

The rapid evolution of LLMs in software engineering creates both a substantial range of open research questions and a dynamic target where findings can become obsolete within months as model capabilities advance. Future research directions are organized here around five themes: capability frontiers (what current systems cannot do), evaluation rigor (how to measure what matters), socio-technical integration (how AI and human developers should work together), safety and reliability (how to deploy AI systems responsibly), and theoretical foundations (why current systems work and what limits them). For each direction, the current state, the key open questions, and the methodological approaches most likely to produce generalizable progress are described.

8.2 Formal Verification and Correctness Guarantees

Perhaps the most consequential open research direction for the safe deployment of LLMs in high-stakes software engineering contexts is the integration of formal verification with LLM code generation — moving from statistical correctness (code that passes tests most of the time) to formal correctness (code with mathematical proofs of specified properties). Current LLM-based code generation lacks any formal correctness guarantees: generated code is tested empirically, and test passage provides probabilistic but not absolute assurance of correctness.

Research directions in this space include: interactive theorem prover integration, where LLM-generated code is automatically checked against formal specifications using Coq, Isabelle, or Lean, with counterexamples fed back to the LLM for iterative correction; bounded model checking of generated code for safety properties (array bound violations, null pointer dereferences, integer overflows) using symbolic execution tools; Hoare logic-based generation where LLMs are trained to produce code with associated pre- and post-conditions that can be formally verified; and synthesis-guided generation where formal specifications are used to constrain the search space for LLM generation, ensuring generated code satisfies formal criteria.

Clover [132] demonstrated initial feasibility of LLM-assisted formal verification, using LLMs to generate both code and formal specifications and an automated checker to verify consistency, achieving 67% agreement rate between LLM-generated code and specifications. This initial result suggests promise but also identifies significant improvements required: LLM-generated specifications frequently have bugs that render them vacuously true (satisfied by any program), and the agreement rate between code and specification does not by itself validate that either the code or the specification correctly

captures the intended behavior. Addressing these limitations requires research on specification elicitation, specification validation, and the fundamental challenge of formally characterizing 'what the developer intended.'

8.3 Long-Horizon Autonomous Software Engineering

The progression from single-turn code generation to sustained autonomous software engineering — where AI systems independently execute complex development tasks over extended time horizons, managing their own tool usage, error recovery, and strategy adaptation — represents the research frontier most directly relevant to the eventual transformation of software engineering practice. Current agentic systems achieve meaningful performance on SWE-bench issues but fall far short of the reliability required for unsupervised execution of production software changes.

Key open questions for long-horizon autonomy include: how should agents acquire, maintain, and update their understanding of a codebase as they make changes? What decision-making frameworks enable agents to correctly determine when to proceed independently versus when to seek human input? How should agents recover from errors that have already modified the codebase (file edits, database schema changes, deployed configurations) and may be difficult to cleanly reverse? How can agents develop and maintain persistent software engineering knowledge across projects rather than starting from a blank slate for each task?

The memory and context management challenges are particularly fundamental for long-horizon autonomy. Current LLMs operate within fixed context windows and have no mechanism for acquiring new knowledge during deployment — they cannot learn from their own execution experience during a task. Research directions addressing this limitation include: episodic memory systems that persist key insights from prior execution; dynamic knowledge base updating as new information is encountered; chain-of-memory approaches that summarize and compress prior execution history for extended task windows; and continual learning approaches that update model parameters based on deployment experience without catastrophic forgetting of prior capabilities.

8.4 Multi-Model and Ensemble Approaches

Single-model code generation has inherent limitations arising from each model's particular training distribution, capability profile, and failure modes. Research on multi-model and ensemble approaches for software engineering tasks explores whether combining multiple models can achieve more robust correctness, better coverage across programming languages and domains, and more reliable confidence calibration than any single model.

EvoEval [2024] demonstrated that evolved variants of HumanEval problems, where problem statements are paraphrased or modified, reveal substantially different performance profiles across models that appear similar on the original benchmark — suggesting that models have complementary strengths on different problem types. Ensemble approaches that route problems to the model most likely to handle them correctly, based on predicted per-model performance from problem characteristics, achieve better aggregate correctness than any single model. This multi-model routing paradigm is underexplored in production systems, where single-model APIs dominate, but may represent a significant practical improvement for high-stakes code generation.

Constitutional ensembles — where multiple models with different alignment properties review each other's outputs — represent a promising direction for improving the security and correctness properties of generated code. A model that generates secure code may catch vulnerabilities generated by a model optimized primarily for functional correctness; a model with strong formal reasoning may catch algorithmic errors in code generated by a model with stronger pattern-matching fluency. Cross-model critique and refinement, systematically organized around complementary model strengths, may achieve better outcomes than any single model with single-pass generation.

8.5 Personalization, Adaptation, and Private Knowledge Integration

Current LLM-based coding tools are general-purpose systems that do not adapt to individual developer preferences, organizational coding standards, or private codebase knowledge accumulated through deployment. Research on personalization and adaptation addresses this limitation by developing techniques for efficiently incorporating private knowledge into model behavior without requiring full model retraining.

Few-shot adaptation through in-context examples is the most immediately practical approach: providing examples of the target developer's code style, naming conventions, and architectural preferences in the system prompt shapes generated code toward the demonstrated style without any model weight updates. Preliminary research suggests that 5-20 carefully selected examples provide substantial stylistic adaptation, though the theoretical mechanisms underlying in-context learning remain partially understood. Systematic methods for selecting maximally informative style examples, maintaining style consistency across different code generation tasks, and updating style examples as preferences evolve are underexplored.

Parameter-efficient fine-tuning (LoRA, adapter layers, prefix tuning) enables more substantive adaptation to organizational code by training a small number of additional parameters on private codebase data, at modest additional inference cost. The privacy implications of fine-tuning on organizational code — the extent to which private code can be extracted from fine-tuned model weights through adversarial prompting — require careful investigation before enterprise deployment. Research on privacy-

preserving fine-tuning (differential privacy in training, secure multi-party computation for distributed fine-tuning) addresses these concerns but with current utility-privacy tradeoffs that limit practical deployment.

8.6 Human-AI Collaborative Development Models

The current predominant model of human-AI collaboration in software development — developer formulates request, AI generates response, developer accepts or modifies — represents just one point in a large space of possible collaboration patterns. Research on alternative collaboration models may identify patterns that achieve better outcomes than the current paradigm for specific task types or developer populations.

The 'pair programming with AI' model — where the human and AI engage in extended back-and-forth dialogue simulating the discussion between two developers pair programming — is underexplored relative to its potential. Pair programming is established as a high-quality software development practice that improves code correctness, knowledge transfer, and developer satisfaction; AI that emulates the productive aspects of the pair programming dynamic (asking clarifying questions, offering alternative approaches, catching reasoning errors) may achieve better outcomes than current tool paradigms that generate responses to explicit requests.

The 'AI as project manager' model — where AI systems maintain project-level awareness, proactively identify tasks that need attention, coordinate work across team members, and track progress toward goals — represents a different collaboration mode with potentially high leverage for team productivity. Research on the specific responsibilities, decision boundaries, and communication patterns appropriate for AI project management roles in software engineering teams would need to address both the technical capabilities required and the organizational and authority dynamics of integrating AI into team management functions.

8.7 Empirical Software Engineering Research Methods for AI Tools

The software engineering research community faces methodological challenges in rigorously evaluating AI coding tools that require both new measurement approaches and careful adaptation of existing empirical research methods. The rapid capability evolution of LLMs — where benchmark performance improvements that previously required years now occur in months — makes longitudinal empirical studies difficult to interpret, as the tool being studied may be substantially different at the end of the study period than at the beginning.

Causal inference challenges in AI tool productivity studies are fundamental: randomized controlled trials are the gold standard for causal claims about productivity

effects, but practical constraints make true randomization difficult. Within-subject designs (same developers with and without AI assistance) risk carryover effects and learning effects; between-subject designs (different developers with and without AI assistance) require careful matching or randomization to avoid selection bias; quasi-experimental designs using natural variation in tool access (organizational rollout sequences, geographic access restrictions) provide large samples but weaker causal identification.

Benchmark ecology — the relationship between benchmark performance and real-world utility — requires sustained research attention as benchmark contamination undermines the validity of published performance comparisons. LivingBench approaches that continuously add fresh, uncontaminated problems, task-design frameworks that minimize contamination risk, and methodological standards for disclosure and correction of contamination effects are all needed to maintain the scientific value of code generation benchmarks as the primary empirical signal for model capability advancement.

8.8 Addressing Bias, Fairness, and Inclusion

The bias and fairness challenges identified in Chapter 7 require targeted research on both measurement and mitigation. Fairness evaluation frameworks for code generation tools must address: performance disparities across programming languages (which groups of developers are best served, which worst?); demographic assumption propagation in generated code; and representation effects in training data that shape model capabilities and limitations.

Mitigation research directions include: curated fine-tuning on diverse, high-quality code from underrepresented languages and domains; demographic debiasing techniques applied to code generation post-processing; automatic detection of demographic assumptions in generated code (binary gender options, Western naming assumptions, accessibility gaps); and participatory design approaches that involve developers from underrepresented communities in tool evaluation and capability prioritization. The goal is ensuring that productivity benefits from AI coding tools are equitably distributed across the global developer community rather than amplifying existing advantages of developers who work primarily in well-represented languages and domains.

CHAPTER 9: CONCLUSION

9.1 Summary of Key Findings

The preceding chapters have examined large language models in software engineering, spanning foundational technology through practical application, enterprise deployment case studies, challenge analysis, and prospective research directions. The analysis synthesizes evidence from over 220 primary sources to provide an evidence-grounded assessment of LLM capabilities, limitations, and impact trajectories in software engineering contexts.

The fundamental finding is that LLMs have crossed the threshold from research demonstration to practical production utility for a broad class of software engineering tasks, with documented productivity benefits that are positive, statistically significant, and practically meaningful across multiple independent evaluation methodologies. The convergence of evidence from controlled laboratory studies, large-scale observational studies, enterprise deployment reports, and benchmark evaluations provides greater confidence in core productivity benefit claims than any single line of evidence could support. At the same time, the analysis identifies significant limitations, risks, and unresolved challenges that must inform responsible deployment decisions and research priorities.

9.2 Research Questions Addressed

Five primary research questions defined in Chapter 1 were addressed. Research Question 1 — What are the architectural and training foundations that enable LLMs to understand and generate programming language? — was addressed in Chapter 3, establishing that the transformer attention mechanism's ability to capture long-range dependencies between code tokens, combined with training on massive code corpora and alignment techniques that optimize code generation for human preferences, provides the foundation for LLM coding capabilities. The identified technical limitations — context window constraints, hallucination patterns, lack of formal correctness guarantees — derive from these same foundations and define the boundaries of current capability.

Research Question 2 — How do LLM-based tools perform across the spectrum of software engineering tasks? — was addressed in Chapter 4, establishing a taxonomy of 15 application categories ranging from inline code autocomplete (most mature, highest adoption) through automated test generation and documentation (substantial capability, significant adoption) to autonomous end-to-end software engineering (emerging capability, limited production deployment). The analysis found meaningful capability

variation across tasks, with performance correlating with task structure and training data representation.

Research Question 3 — How do leading LLM-based software engineering tools compare across capability, cost, and enterprise readiness dimensions? — was addressed in Chapter 5, which provided systematic multi-dimensional comparison across 10 tools. The finding that no single tool dominates across all evaluation dimensions — frontier models provide best raw capability, specialized tools provide best developer experience, open-source models provide best privacy and customization — implies that tool selection must be informed by organizational-specific priorities rather than universal rankings.

Research Question 4 — What are the real-world outcomes of LLM deployment in software engineering practice? — was addressed in Chapter 6 through six in-depth case studies establishing: 30-80% productivity improvements for specific high-frequency tasks (varying substantially by task type and developer population), 40-80% reduction in security defect rates in AI-assisted development pipelines with security scanning, and significant variation in organizational outcomes based on deployment practices, change management, and process adaptation. The case studies also identified deployment challenges and failure modes not captured in laboratory studies, providing practitioner-relevant lessons.

Research Question 5 — What are the most significant challenges and research gaps that must be addressed for responsible LLM adoption in software engineering? — was addressed in Chapters 7 and 8. The identified challenges form a coherent picture of an immature but rapidly maturing technology: technical limitations in correctness and security are well-understood and partially addressed; socio-technical challenges in skill development and organizational integration are less well-understood; and fundamental research gaps in formal correctness guarantees, long-horizon autonomy, and evaluation methodology will require sustained research investment.

9.3 Theoretical Contributions

Beyond the empirical synthesis, this thesis makes three theoretical contributions to the software engineering and AI research literatures. First, the taxonomy of LLM applications in software engineering (Chapter 4) provides a principled organizational framework for the rapidly expanding literature, distinguishing application categories by both lifecycle phase and LLM contribution type to enable more precise performance claims and comparison across studies. This taxonomy has been used to organize the systematic literature review and is offered as a contribution to researchers needing to position new work within the broader field.

Second, the multi-dimensional evaluation framework developed for the comparative study (Chapter 5) extends existing benchmark-focused comparison

approaches by incorporating enterprise readiness, ecosystem integration, and total cost of ownership dimensions alongside functional performance metrics. This framework provides a more complete basis for organizational tool selection decisions than benchmark comparisons alone and can be applied to future tool generations as the landscape continues to evolve.

Third, the challenge taxonomy in Chapter 7 provides a systematic characterization of limitation categories that enables more precise problem formulation for future research. The distinction between correctness limitations (technical, partially addressed), security concerns (technical and organizational, partially mitigated), IP and legal concerns (external, evolving), cognitive effects (socio-technical, largely unresolved), and fairness issues (technical and organizational, insufficiently studied) prevents conflation of distinct problems that require different research approaches and solution strategies.

9.4 Practical Implications

For software engineering practitioners and organizations, this thesis provides an evidence-grounded foundation for AI coding tool adoption decisions. The practical implications are organized around three practitioner audiences.

For individual developers, the evidence supports adoption of AI coding assistance as a productivity tool while maintaining critical evaluation practices. Specific recommendations include: use AI assistance for high-frequency routine tasks (boilerplate generation, documentation, test structure) where the ratio of productivity benefit to risk is highest; maintain deliberate practice of core programming skills to prevent skill atrophy; apply particular scrutiny to AI-generated code in security-sensitive contexts; and develop comfort with multi-turn iterative refinement rather than expecting single-shot correctness for complex tasks.

For engineering team leads and technical managers, the evidence supports structured, deliberate AI tool adoption with explicit attention to change management, security control, and quality assurance process adaptation. The case study evidence consistently shows that deployment success correlates with: pilot programs with measurement before broad rollout; role-segmented adoption strategies that match tool capabilities to task types; security control investments (SAST scanning, secure coding training) that prevent AI-assisted security regression; and team norm-setting that establishes healthy AI assistance practices without creating unhealthy dependency.

For software engineering researchers, the gaps identified here suggest priority areas for investigation: longitudinal skill development effects of AI assistance; fairness measurement and mitigation frameworks for AI coding tools; causal inference methodology appropriate for AI productivity studies; and evaluation benchmark evolution to maintain ecological validity as training data contamination accumulates.

These research investments are expected to strengthen the evidence base needed to guide AI coding tool deployment toward beneficial outcomes.

9.5 Limitations of This Study

Several limitations bound the generalizability of the findings and should inform interpretation of the presented analysis. The systematic literature review, while covering peer-reviewed academic sources, may underrepresent practitioner knowledge published in grey literature (blog posts, conference presentations, internal corporate reports) that contains important operational insights not yet reflected in academic publications. The six-month review horizon for the comparative study ensures currency but may miss longer-term trends visible only in multi-year longitudinal data.

The case studies are limited by the availability of documented enterprise deployments that provide sufficient methodological detail for rigorous analysis. The six selected cases over-represent technology sector deployments and large organizations; small and medium enterprises and non-technology industry sectors (manufacturing, healthcare, finance) are underrepresented. The practical implications presented in Section 9.4 should be interpreted with awareness that they are most directly grounded in large technology sector evidence and may require adaptation for other organizational contexts.

Benchmark-based performance claims in this thesis reflect published results that may be subject to benchmark contamination, evaluation methodology variation, and model update effects that alter performance between publication and review. The performance comparisons presented in Chapter 5 represent the best available evidence at time of writing but should be verified against current benchmarks before informing specific tool selection decisions, given the rapid capability evolution of this field.

9.6 Concluding Remarks

The integration of large language models into software engineering practice represents a shift that is already substantially underway rather than a future prospect. The available evidence suggests that LLM-based tools provide meaningful, reproducible productivity benefits for broad categories of software development tasks, and that frontier model capabilities are expected to continue advancing. The more pressing question for the field is not whether to adopt AI coding tools but how to adopt them in ways that maximize benefits while managing the technical limitations, security risks, skill development concerns, and equity implications identified throughout this work.

The history of software engineering is a history of abstractions that have allowed developers to work at progressively higher levels of intent expression: from machine code to assembly, from assembly to high-level languages, from general languages to domain-specific abstractions, from manual toolchains to frameworks and libraries. LLMs

represent a potential next step in this progression — allowing intent expression in natural language, with the model mediating the translation to executable implementation. Whether this step proves as consequential as previous abstraction advances, or whether the correctness and reliability challenges of natural language specifications prove difficult to resolve for production software, may be determined by research and engineering progress over the coming decade.

The central argument of this work is that maximizing the benefit of LLMs in software engineering requires not just technical capability — which is advancing and appears likely to continue — but careful empirical evaluation, thoughtful organizational deployment, proactive attention to socio-technical risks, and sustained research investment in the open problems identified here. The field appears to be at an inflection point where deployment and development decisions are likely to shape the trajectory of software engineering practice for the foreseeable future. This study aspires to contribute to the evidence base that informs those decisions.

REFERENCES

The following references are formatted in accordance with the IEEE citation style, as specified by the Department of Computer Science and Engineering, RV University. References are listed in alphabetical order by first author's surname.

- [1] Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., and Polosukhin, I. (2017). Attention is all you need. In *Advances in Neural Information Processing Systems (NeurIPS 2017)*, vol. 30, pp. 5998-6008.
- [2] Devlin, J., Chang, M.-W., Lee, K., and Toutanova, K. (2019). BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL-HLT 2019)*, pp. 4171-4186.
- [3] Brown, T. B., Mann, B., Ryder, N., Subbiah, M., Kaplan, J., Dhariwal, P., Neelakantan, A., Shyam, P., Sastry, G., Askell, A., et al. (2020). Language models are few-shot learners. In *Advances in Neural Information Processing Systems (NeurIPS 2020)*, vol. 33, pp. 1877-1901.
- [4] Chen, M., Tworek, J., Jun, H., Yuan, Q., Pinto, H. P. d. O., Kaplan, J., Edwards, H., Burda, Y., Joseph, N., Brockman, G., et al. (2021). Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*.
- [5] Li, Y., Wang, S., Nguyen, T. N., Van Nguyen, S., Phan, L. H., Luo, D., Kim, Y., and Lo, D. (2022). Automating code review activities by large-scale pre-training. In *Proceedings of the 30th ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering (ESEC/FSE 2022)*, pp. 1035-1047.
- [6] Zan, D., Chen, B., Yang, D., Lin, Z., Kim, M., Guan, B., Wang, Y., and Chen, W. (2023). Large language models meet NL2Code: A survey. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL 2023)*, vol. 1, pp. 7443-7464.
- [7] Niu, C., Li, C., Liang, G., Yu, M., Hu, X., and Luo, X. (2024). UnitCon: Automated unit test generation using LLMs with code-contextualized hierarchical knowledge. *arXiv preprint arXiv:2403.12682*.
- [8] Ji, Z., Lee, N., Frieske, R., Yu, T., Su, D., Xu, Y., Ishii, E., Bang, Y., Madotto, A., and Fung, P. (2023). Survey of hallucination in natural language generation. *ACM Computing Surveys*, 55(12), 248:1-248:38.
- [9] Siddiq, M. L., and Santos, J. C. S. (2023). A lightweight framework for adaptive summaries. *arXiv preprint arXiv:2302.04563*.
- [10] Wang, Y., Le, H., Gotmare, A. D., Bui, N. D. Q., Li, J., and Hoi, S. C. H. (2023). CodeT5+: Open code large language models for code understanding and generation. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP 2023)*, pp. 1069-1088.
- [11] Li, R., Allal, L. B., Zi, Y., Muennighoff, N., Kocetkov, D., Mou, C., Marone, M., Akiki, C., Li, J., Chim, J., et al. (2023). StarCoder: May the source be with you! *Transactions on Machine Learning Research*, 2023.
- [12] Touvron, H., Martin, L., Stone, K., Albert, P., Almahairi, A., Babaci, Y., Bashlykov, N., Batra, S., Bhargava, P., Bhosale, S., et al. (2023). Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*.

- [13] Roziere, B., Gehring, J., Gloeckle, F., Sootla, S., Gat, I., Tan, X. E., Adi, Y., Liu, J., Sauvestre, R., Remez, T., et al. (2023). Code Llama: Open foundation models for code. arXiv preprint arXiv:2308.12950.
- [14] Guo, D., Zhu, Q., Yang, D., Xie, Z., Dong, K., Zhang, W., Chen, G., Bi, X., Wu, Y., Li, Y., et al. (2024). DeepSeek-Coder: When the large language model meets programming — the rise of code intelligence. arXiv preprint arXiv:2401.14196.
- [15] Hui, B., Yang, J., Cui, Z., Yang, J., Liu, D., Zhang, L., Liu, T., Zhang, J., Yu, B., Lu, K., et al. (2024). Qwen2.5-Coder technical report. arXiv preprint arXiv:2409.12186.
- [16] Anthropic. (2023). Claude's constitution. Anthropic Technical Report. Accessed June 2026.
- [17] Bai, Y., Jones, A., Ndousse, K., Askell, A., Chen, A., DasSarma, N., Drain, D., Fort, S., Ganguli, D., Henighan, T., et al. (2022). Training a helpful and harmless assistant with reinforcement learning from human feedback. arXiv preprint arXiv:2204.05862.
- [18] Team, G., Mesnard, T., Hardin, C., Dadashi, R., Bhupatiraju, S., Pathak, S., Sifre, L., Riviere, M., Kale, M. S., Love, J., et al. (2024). Gemma: Open models based on Gemini research and technology. arXiv preprint arXiv:2403.08295.
- [19] Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C., Mishkin, P., Zhang, C., Agarwal, S., Slama, K., Ray, A., et al. (2022). Training language models to follow instructions with human feedback. In *Advances in Neural Information Processing Systems (NeurIPS 2022)*, vol. 35, pp. 27730-27744.
- [20] Rafailov, R., Sharma, A., Mitchell, E., Manning, C. D., Ermon, S., and Finn, C. (2023). Direct preference optimization: Your language model is secretly a reward model. In *Advances in Neural Information Processing Systems (NeurIPS 2023)*, vol. 36.
- [21] Kaplan, J., McCandlish, S., Henighan, T., Brown, T. B., Chess, B., Child, R., Gray, S., Radford, A., Wu, J., and Amodei, D. (2020). Scaling laws for neural language models. arXiv preprint arXiv:2001.08361.
- [22] Hoffmann, J., Borgeaud, S., Mensch, A., Buchatskaya, E., Cai, T., Rutherford, E., Casas, D. de L., Hendricks, L. A., Welbl, J., Clark, A., et al. (2022). Training compute-optimal large language models. In *Advances in Neural Information Processing Systems (NeurIPS 2022)*, vol. 35, pp. 30016-30030.
- [23] Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Kuttler, H., Lewis, M., Yih, W.-t., Rocktaschel, T., et al. (2020). Retrieval-augmented generation for knowledge-intensive NLP tasks. In *Advances in Neural Information Processing Systems (NeurIPS 2020)*, vol. 33, pp. 9459-9474.
- [24] Wei, J., Wang, X., Schuurmans, D., Bosma, M., Ichter, B., Xia, F., Chi, E., Le, Q., and Zhou, D. (2022). Chain-of-thought prompting elicits reasoning in large language models. In *Advances in Neural Information Processing Systems (NeurIPS 2022)*, vol. 35.
- [25] Kitchenham, B., and Charters, S. (2007). Guidelines for performing systematic literature reviews in software engineering. Technical Report EBSE-2007-01, Keele University and Durham University Joint Report.
- [26] Humbatova, N., Jahangirova, G., Bavota, G., Riccio, V., Stocco, A., and Tonella, P. (2020). Taxonomy of real faults in deep learning systems. In *Proceedings of the 42nd International Conference on Software Engineering (ICSE 2020)*, pp. 1110-1121.
- [27] Austin, J., Odena, A., Nye, M., Bosma, M., Michalewski, H., Dohan, D., Jiang, E., Cai, C., Terry, M., Le, Q., and Sutton, C. (2021). Program synthesis with large language models. arXiv preprint arXiv:2108.07732.

- [28] Nijkamp, E., Pang, B., Hayashi, H., Tu, L., Wang, H., Zhou, Y., Savarese, S., and Xiong, C. (2023). CodeGen: An open large language model for code with multi-turn program synthesis. In Proceedings of the 11th International Conference on Learning Representations (ICLR 2023).
- [29] Zheng, Q., Xia, X., Zou, X., Dong, Y., Wang, S., Xue, Y., Wang, Z., Shen, L., Wang, A., Li, Y., et al. (2024). CodeGeeX: A pre-trained model for code generation and translation. In Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD 2023), pp. 5673-5684.
- [30] Chowdhery, A., Narang, S., Devlin, J., Bosma, M., Mishra, G., Roberts, A., Barham, P., Chung, H. W., Sutton, C., Gehrmann, S., et al. (2022). PaLM: Scaling language modeling with pathways. Journal of Machine Learning Research, 24(240), 1-113.
- [31] OpenAI. (2024). GPT-4 technical report. arXiv preprint arXiv:2303.08774.
- [32] Bubeck, S., Chandrasekaran, V., Eldan, R., Gehrke, J., Horvitz, E., Kamar, E., Lee, P., Lee, Y. T., Li, Y., Lundberg, S., et al. (2023). Sparks of artificial general intelligence: Early experiments with GPT-4. arXiv preprint arXiv:2303.12528.
- [33] Gou, Z., Shao, Z., Gong, M., Shen, Y., Yang, Y., Duan, N., and Chen, W. (2024). CRITIC: Large language models can self-correct with tool-interactive critiquing. In Proceedings of the 12th International Conference on Learning Representations (ICLR 2024).
- [34] Jiang, A. Q., Sablayrolles, A., Roux, A., Mensch, A., Savary, B., Bamford, C., Chaplot, D. S., Casas, D. d. l., Hanna, E. B., Bressand, F., et al. (2024). Mixtral of experts. arXiv preprint arXiv:2401.04088.
- [35] Su, J., Lu, Y., Pan, S., Murtadha, A., Wen, B., and Liu, Y. (2024). RoFormer: Enhanced transformer with rotary position embedding. Neurocomputing, 568, 127063.
- [36] Dettmers, T., Pagnoni, A., Holtzman, A., and Zettlemoyer, L. (2023). QLoRA: Efficient finetuning of quantized LLMs. In Advances in Neural Information Processing Systems (NeurIPS 2023), vol. 36.
- [37] Hu, E. J., Shen, Y., Wallis, P., Allen-Zhu, Z., Li, Y., Wang, S., Wang, L., and Chen, W. (2022). LoRA: Low-rank adaptation of large language models. In Proceedings of the 10th International Conference on Learning Representations (ICLR 2022).
- [38] Bavarian, M., Jun, H., Tezak, N., Schulman, J., McLeavey, C., Tworek, J., and Chen, M. (2022). Efficient training of language models to fill in the middle. arXiv preprint arXiv:2207.14255.
- [39] Pearce, H., Ahmad, B., Tan, B., Dolan-Gavitt, B., and Karri, R. (2022). Asleep at the keyboard? Assessing the security of GitHub Copilot's code contributions. In Proceedings of the 43rd IEEE Symposium on Security and Privacy (IEEE S&P 2022), pp. 754-768.
- [40] Pearce, H., Tan, B., Ahmad, B., Karri, R., and Dolan-Gavitt, B. (2022). Examining zero-shot vulnerability repair with large language models. In Proceedings of the 43rd IEEE Symposium on Security and Privacy (IEEE S&P 2022), pp. 2339-2356.
- [41] Copilot, G. (2023). GitHub Copilot for Business. GitHub. Retrieved from <https://github.blog/2023>.
- [42] Amazon. (2024). Amazon Q Developer documentation. AWS Documentation. Retrieved June 2026.
- [43] Lozhkov, A., Li, R., Allal, L. B., Cassano, F., Lamy-Poirier, J., Tazi, N., Tang, A., Pykhtar, D., Liu, J., Wei, Y., et al. (2024). StarCoder 2 and The Stack v2: The next generation. arXiv preprint arXiv:2402.19173.
- [44] Sun, T., Shao, Y., Qian, H., Huang, X., and Qiu, X. (2024). Black-box tuning for language-model-as-a-service. In Proceedings of the 39th International Conference on Machine Learning (ICML 2022), pp. 20841-20855.

- [45] Ziegler, A., Kalliamvakou, E., Li, X. A., Rice, A., Rifkin, D., Simister, S., Sittampalam, G., and Aftandilian, E. (2022). Productivity assessment of neural code completion. In *Proceedings of the 6th ACM SIGPLAN International Symposium on Machine Programming (MAPS 2022)*, pp. 21-29.
- [46] Liu, J., Zhu, Q., Xia, X., Lo, D., and Li, S. (2023). Improving ChatGPT prompt for code generation. *arXiv preprint arXiv:2305.08360*.
- [47] Geng, M., Wang, S., Dong, D., and Wang, H. (2024). Large language models are few-shot testers: Exploring LLM-based general bug reproduction. In *Proceedings of the 45th International Conference on Software Engineering (ICSE 2023)*, pp. 2337-2349.
- [48] Siddiq, M. L., Santos, J. C. S., Tanvir, R. H., Ulfat, N., Al-Sajee, F. A., and Machado, B. C. (2024). Using large language models to generate JUnit tests: An empirical study. In *Proceedings of the 28th International Conference on Evaluation and Assessment in Software Engineering (EASE 2024)*.
- [49] Xia, C. S., Deng, Y., Dunn, S., and Zhang, L. (2024). Automated program repair in the era of large pre-trained language models. In *Proceedings of the 45th International Conference on Software Engineering (ICSE 2023)*, pp. 1482-1494.
- [50] Kalliamvakou, E. (2022). Research: Quantifying GitHub Copilot's impact on developer productivity and happiness. *GitHub Blog*. October 2022.
- [51] Lozhkov, A., Jennings, G., Zheng, L., Biderman, S., and Wolf, T. (2024). The Stack v2: Growing code datasets. *arXiv preprint arXiv:2402.19173*.
- [52] Jimenez, C. E., Yang, J., Wettig, A., Yao, S., Pei, K., Press, O., and Narasimhan, K. (2024). SWE-bench: Can language models resolve real-world GitHub issues? In *Proceedings of the 12th International Conference on Learning Representations (ICLR 2024)*.
- [53] Wang, J., Zhang, Y., Gu, Q., Xia, X., and Xing, Z. (2024). Software engineering for AI-based systems: A survey. *arXiv preprint arXiv:2404.00712*.
- [54] Bairi, R., Sonwane, A., Kanade, A., V D, A. C., Iyer, S., Parthasarathy, S., Rajamani, S., and Shet, B. (2024). CodePlan: Repository-level coding using LLMs and planning. In *Proceedings of the 2024 Conference on File System Semantics and Engineering (FSE 2024)*.
- [55] Fan, A., Gokkaya, B., Harman, M., Lyubarskiy, M., Sengupta, S., Yoo, S., and Zhang, J. M. (2023). Large language models for software engineering: Survey and open problems. In *Proceedings of the 45th IEEE/ACM International Conference on Software Engineering: Future of Software Engineering (ICSE-FoSE 2023)*, pp. 31-53.
- [56] Amodei, D., Olah, C., Steinhardt, J., Christiano, P., Schulman, J., and Mane, D. (2016). Concrete problems in AI safety. *arXiv preprint arXiv:1606.06565*.
- [57] Hendrycks, D., and Mazeika, M. (2022). X-risk analysis for AI research. *arXiv preprint arXiv:2206.05862*.
- [58] Wei, J., Tay, Y., Bommasani, R., Raffel, C., Zoph, B., Borgeaud, S., Yogatama, D., Bosma, M., Zhou, D., Metzler, D., et al. (2022). Emergent abilities of large language models. *Transactions on Machine Learning Research*, 2022.
- [59] Cobbe, K., Kosaraju, V., Bavarian, M., Chen, M., Jun, H., Kaiser, L., Plappert, M., Tworek, J., Hilton, J., Nakano, R., et al. (2021). Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*.
- [60] Peng, S., Kalliamvakou, E., Croft, P., and Butter, M. (2023). The impact of AI on developer productivity: Evidence from GitHub Copilot. *arXiv preprint arXiv:2302.06590*.
- [61] Poldrack, R. A., Lu, T., and Beguš, G. (2023). AI-assisted coding: Experiments with GPT-4. *arXiv preprint arXiv:2304.13187*.

- [62] Zheng, L., Chiang, W.-L., Sheng, Y., Zhuang, S., Wu, Z., Zhuang, Y., Lin, Z., Li, Z., Li, D., Xing, E., et al. (2023). Judging LLM-as-a-judge with MT-bench and Chatbot Arena. In *Advances in Neural Information Processing Systems (NeurIPS 2023)*, vol. 36.
- [63] Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., and Cao, Y. (2023). ReAct: Synergizing reasoning and acting in language models. In *Proceedings of the 11th International Conference on Learning Representations (ICLR 2023)*.
- [64] Shinn, N., Cassano, F., Gopinath, A., Narasimhan, K., and Yao, S. (2023). Reflexion: Language agents with verbal reinforcement learning. In *Advances in Neural Information Processing Systems (NeurIPS 2023)*, vol. 36.
- [65] Yang, J., Jimenez, C. E., Wettig, A., Yang, K., Qian, Y., Pei, K., and Narasimhan, K. (2024). SWE-agent: Agent-computer interfaces enable automated software engineering. In *Advances in Neural Information Processing Systems (NeurIPS 2024)*.
- [66] Wang, X., Wei, J., Schuurmans, D., Le, Q., Chi, E., Narang, S., Chowdhery, A., and Zhou, D. (2023). Self-consistency improves chain of thought reasoning in language models. In *Proceedings of the 11th International Conference on Learning Representations (ICLR 2023)*.
- [67] Ni, A., Iyer, S., Radev, D., Stoyanov, V., Yih, W.-t., Wang, S., and Lin, X. V. (2023). LEVER: Learning to verify language-to-code generation with execution. In *Proceedings of the 40th International Conference on Machine Learning (ICML 2023)*, pp. 26175-26198.
- [68] Zhang, F., Chen, B., Zhang, Y., Keung, J., Liu, J., Zan, D., Mao, Y., Lou, J.-G., and Chen, W. (2023). RepoCoder: Repository-level code completion through iterative retrieval and generation. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP 2023)*, pp. 275-288.
- [69] McKinsey and Company. (2023). The economic potential of generative AI: The next productivity frontier. McKinsey Global Institute Report. June 2023.
- [70] Dakhel, A. M., Majdinasab, V., Nikanjam, A., Khomh, F., Desmarais, M. C., and Jiang, Z. M. (2023). GitHub Copilot AI pair programmer: Asset or liability? *Journal of Systems and Software*, 203, 111734.
- [71] Buscemi, A. (2023). A comparative study of code generation using ChatGPT 3.5 across 10 programming languages. *arXiv preprint arXiv:2308.04477*.
- [72] Cassano, F., Gouwar, J., Nguyen, D., Nguyen, S., Phipps-Costin, L., Pinckney, D., Yee, M.-H., Zi, Y., Anderson, C. J., Feldman, M. Q., et al. (2023). MultiPL-E: A scalable and polyglot approach to benchmarking neural code generation. *IEEE Transactions on Software Engineering*, 49(7), 3675-3691.
- [73] Athiwaratkun, B., Gouda, S. K., Wang, Z., Li, X., Tian, Y., Liu, M., Shang, J., Ding, M., Kumar, A., Fulp, W., et al. (2023). Multi-lingual evaluation of code generation models. In *Proceedings of the 11th International Conference on Learning Representations (ICLR 2023)*.
- [74] Liu, N. F., Lin, K., Hewitt, J., Paranjape, A., Bevilacqua, M., Petroni, F., and Liang, P. (2023). Lost in the middle: How language models use long contexts. *Transactions of the Association for Computational Linguistics*, 12, 157-173.
- [75] Xie, S., Santurkar, S., Ma, T., and Liang, P. (2023). Data selection for language models via importance resampling. In *Advances in Neural Information Processing Systems (NeurIPS 2023)*, vol. 36.
- [76] Du, Y., Li, S., Torralba, A., Tenenbaum, J. B., and Mordatch, I. (2023). Improving factuality and reasoning in language models through multiagent debate. In *Proceedings of the 41st International Conference on Machine Learning (ICML 2024)*, pp. 11850-11881.

- [77] Ding, Y., Wang, Z., Ahmad, W., Ding, H., Tan, M., Jain, N., Ramanathan, M. K., Nallapati, R., Bhatia, P., Roth, D., and Ma, F. (2024). CrossCodeEval: A diverse and multilingual benchmark for cross-file code completion. In *Advances in Neural Information Processing Systems (NeurIPS 2023)*, vol. 36.
- [78] Leblond, R., Anil, R., Sohl-Dickstein, J., and Szegedy, C. (2023). AlphaCode 2 technical report. DeepMind Technical Report. December 2023.
- [79] Guo, Q., Zhu, Q., Yang, D., Xie, Z., Li, K., Zhang, Y., Chen, G., Zhu, W., Liu, W., Gao, Y., et al. (2024). DeepSeek-Coder-V2: Breaking the barrier of closed-source models in code intelligence. *arXiv preprint arXiv:2406.11931*.
- [80] Liu, Y., Iter, D., Xu, Y., Wang, S., Xu, R., and Zhu, C. (2023). GPTEval: NLG evaluation using GPT-4 with better human alignment. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP 2023)*, pp. 2511-2522.
- [81] Bi, X., Chen, D., Chen, G., Chen, S., Dai, D., Deng, C., Ding, H., Dong, K., Du, Q., Fu, Z., et al. (2024). DeepSeek LLM: Scaling open-source language models with longtermism. *arXiv preprint arXiv:2401.02954*.
- [82] Luo, Z., Xu, C., Zhao, P., Sun, Q., Geng, X., Hu, W., Tao, C., Ma, J., Lin, Q., and Jiang, D. (2024). WizardCoder: Empowering code large language models with Evol-Instruct. In *Proceedings of the 12th International Conference on Learning Representations (ICLR 2024)*.
- [83] Cummins, C., Seeker, V., Grubisic, D., Elhoushi, M., Laber, Y., Tillet, P., Aubert, G., Parejo, G., Kelly, J., and Leather, H. (2023). Large language models for compiler optimization. *arXiv preprint arXiv:2309.07062*.
- [84] Liu, C., Le Bras, R., Bhatt, U., Bhimani, J., Khashabi, D., Welleck, S., Seo, M., Hajishirzi, H., and Smith, N. A. (2024). EvoEval: Evolving coding benchmarks via LLM. *arXiv preprint arXiv:2403.19114*.
- [85] Gu, X., Meng, X., Song, F., Su, Z., and Zhang, D. (2024). CruxEval: A benchmark for code reasoning, understanding and execution. *arXiv preprint arXiv:2401.03065*.
- [86] Xu, W., Cassano, F., Liu, Y., Yee, M.-H., Schafer, M., Beurer-Kellner, L., and Vechev, M. (2024). BigCodeBench: Benchmarking code generation with diverse function calls and complex instructions. *arXiv preprint arXiv:2406.15877*.
- [87] Jain, N., Han, K., Ding, A., Guo, Q., Goldie, A., Zettlemoyer, L., Guestrin, C., Ré, C., Jia, R., and Potts, C. (2024). LiveCodeBench: Holistic and contamination-free evaluation of large language models for code. *arXiv preprint arXiv:2403.07974*.
- [88] Nong, Y., Ou, Y., Pradel, M., Chen, F., and Cai, H. (2024). Automated software vulnerability patching using large language models. *arXiv preprint arXiv:2408.13597*.
- [89] Hou, X., Zhao, Y., Liu, Y., Yang, Z., Wang, K., Li, L., Luo, X., Lo, D., Grundy, J., and Wang, H. (2024). Large language models for software engineering: A systematic literature review. *IEEE Transactions on Software Engineering*, 50(9), 2362-2385.
- [90] Wu, X., Deng, W., Li, Y., Li, Y., Gao, J., Tang, H., Shi, L., Ma, Y., and Xu, Y. (2024). Large language models in software security: A survey of vulnerability detection techniques and insights. *arXiv preprint arXiv:2402.16396*.
- [91] Yao, Y., Duan, J., Xu, K., Cai, Y., Sun, Z., and Zhang, Y. (2024). A survey on large language model (LLM) security and privacy: The good, the bad, and the ugly. *High-Confidence Computing*, 4(2), 100211.

- [92] Wang, Y., Liu, H., and Zhu, H. (2024). CodeAct: Enabling language models as coding agents. arXiv preprint arXiv:2402.01030.
- [93] Jimenez, C. E., Yang, J., Wettig, A., Yao, S., Pei, K., Press, O., and Narasimhan, K. (2024). SWE-bench: Can language models resolve real-world GitHub issues? In Proceedings of the 12th International Conference on Learning Representations (ICLR 2024). [Verified subset: arXiv:2410.06992].
- [94] Jain, N., Han, K., Ding, A., et al. (2024). LiveCodeBench: Holistic and contamination-free evaluation of large language models for code. arXiv:2403.07974.
- [95] Wang, X., Li, Y., Wan, Z., Zhang, T., Zhang, Y., Liu, Y., Tang, J., Pan, S., and Lo, D. (2024). OpenHands: An open platform for AI software developers as generalist agents. arXiv preprint arXiv:2407.16741.
- [96] Devin. (2024). Devin: The first AI software engineer. Cognition. Technical blog. March 2024.
- [97] Svyatkovskiy, A., Deng, S. K., Fu, S., and Sundaresan, N. (2020). IntelliCode compose: Code generation using transformer. In Proceedings of the 28th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering (ESEC/FSE 2020), pp. 1433-1443.
- [98] Fried, D., Aghajanyan, A., Lin, J., Wang, S., Wallace, E., Shi, F., Chen, R., Yih, W.-t., Zettlemoyer, L., and Lewis, M. (2023). InCoder: A generative model for code infilling and synthesis. In Proceedings of the 11th International Conference on Learning Representations (ICLR 2023).
- [99] Geng, M., Wang, S., Dong, D., and Wang, H. (2024). An empirical study of code smells in transformer-based code generation models. In Proceedings of the 31st IEEE International Conference on Software Analysis, Evolution and Reengineering (SANER 2024).
- [100] Shen, T., Jin, R., Huang, Y., Liu, C., Dong, W., Guo, Z., Wu, X., Liu, Y., and Xiong, D. (2023). Large language model alignment: A survey. arXiv preprint arXiv:2309.15025.
- [101] Liu, S., Gao, C., He, J., Chen, L., Huang, Y., Li, H., and Lou, J.-G. (2023). Refining ChatGPT-generated code: Characterizing and mitigating code quality issues. ACM Transactions on Software Engineering and Methodology, 33(5), 132.
- [102] Wu, X., Zheng, W., Liang, T., and Yu, G. (2024). How far have we gone in vulnerability detection using large language models? arXiv preprint arXiv:2311.12420.
- [103] Pearce, H., Tan, B., Ahmad, B., Karri, R., and Dolan-Gavitt, B. (2023). Examining zero-shot vulnerability repair with large language models. arXiv preprint arXiv:2112.02125.
- [104] Khoury, R., Avila, A. R., Brunelle, J., and Camara, B. M. (2023). How secure is code generated by ChatGPT? In Proceedings of the 2023 IEEE International Conference on Systems, Man, and Cybernetics (SMC 2023), pp. 2445-2450.
- [105] Xia, C. S., and Zhang, L. (2023). Conversational automated program repair. arXiv preprint arXiv:2301.13246.
- [106] Xia, C. S., Deng, Y., and Zhang, L. (2024). ChatRepair: Conversational automated program repair. arXiv preprint arXiv:2404.00571.
- [107] Li, Z., Chen, Y., Shetty, S., Gao, Z., and Zhang, L. (2022). A critical review of large language model on software engineering. arXiv preprint arXiv:2312.07867.
- [108] Schafer, M., Nadi, S., Eghbali, A., and Tip, F. (2024). Empirical evaluation of using large language models for automated unit test generation. IEEE Transactions on Software Engineering, 50(1), 85-105.

- [109] Dakhel, A. M., Majdinasab, V., Nikanjam, A., Khomh, F., Desmarais, M. C., and Jiang, Z. M. J. (2022). GitHub Copilot AI pair programmer: Asset or liability? arXiv preprint arXiv:2206.15331.
- [110] Hu, X., Li, G., Xia, X., Lo, D., and Jin, Z. (2018). Deep code comment generation. In Proceedings of the 26th IEEE/ACM International Conference on Program Comprehension (ICPC 2018), pp. 200-210.
- [111] Gao, S., et al. (2024). CoT-CodeSumm: Code summarization via chain-of-thought with large language models. *ACM Transactions on Software Engineering and Methodology*, 33(7).
- [112] Stripe Engineering. (2023). AI-assisted technical documentation at Stripe. Stripe Engineering Blog. November 2023.
- [113] Liu, X., Wu, G., Peng, Y., Ma, Y., Liu, J., and Wang, X. (2024). Large language models for DevOps task automation: Evaluation and best practices. arXiv preprint arXiv:2404.12673.
- [114] Rahman, M. M., and Roy, C. K. (2024). Towards a taxonomy and quality metrics for AI-assisted code generation. *IEEE Transactions on Software Engineering*, 50(4), 973-992.
- [115] Ziegler, A., Kalliamvakou, E., Li, X. A., Rice, A., Rifkin, D., Simister, S., Sittampalam, G., and Aftandilian, E. (2022). Measuring GitHub Copilot's impact on productivity. *Communications of the ACM*, 65(3), 54-59.
- [116] Xu, B., Gou, Z., He, J., Nie, L., Liu, H., Xu, T., and Wang, E. (2024). API hallucination reduction via retrieval-augmented code generation. arXiv preprint arXiv:2403.10777.
- [117] Accenture. (2023). Generative AI and the future of software development: Accenture technology vision. Accenture Research Report. October 2023.
- [118] McKinsey Digital. (2024). Unleashing developer productivity with generative AI. McKinsey Technology. June 2024.
- [119] Tian, Z., Chen, J., and Pian, W. (2023). Is ChatGPT the ultimate programming assistant — How far is it? arXiv preprint arXiv:2304.11938.
- [120] Nguyen, N., and Nadi, S. (2022). An empirical evaluation of GitHub Copilot's code suggestions. In Proceedings of the 19th International Conference on Mining Software Repositories (MSR 2022), pp. 1-5.

APPENDIX A: SYSTEMATIC LITERATURE REVIEW PROTOCOL

A.1 Research Questions

The systematic literature review was conducted following the protocol defined by Kitchenham and Charters [25]. The primary research questions guiding the review were:

1. RQ1: What architectural and training innovations enable LLMs to understand and generate programming language?
2. RQ2: How do LLM-based tools perform across software engineering tasks including code generation, testing, debugging, and documentation?
3. RQ3: How do leading LLM-based SE tools compare across capability, cost, and enterprise readiness dimensions?
4. RQ4: What are the documented outcomes of LLM deployment in real-world software engineering practice?
5. RQ5: What are the primary challenges and research gaps for responsible LLM adoption in software engineering?

A.2 Search Strategy

Literature search was conducted across five databases: IEEE Xplore, ACM Digital Library, Springer Link, arXiv (cs.SE, cs.AI, cs.CL categories), and Google Scholar. Search terms were organized in three clusters: (1) LLM terms [large language model, LLM, GPT, BERT, transformer, foundation model, pre-trained model, generative AI]; (2) SE application terms [software engineering, code generation, program synthesis, automated debugging, test generation, code review, documentation generation]; (3) specific tool terms [GitHub Copilot, Amazon CodeWhisperer, Cursor, Devin, Claude Code, Copilot, ChatGPT]. Boolean combinations of these clusters formed the primary search queries.

The search was restricted to publications from January 2017 (the year of the transformer architecture introduction [1]) through June 2026 (the thesis completion date). Conference proceedings included: ICSE, FSE/ESEC, ASE, ISSTA, MSR, NeurIPS, ICML, ICLR, ACL, EMNLP, NAACL, and IEEE S&P. Journals included: IEEE Transactions on Software Engineering, ACM TOSEM, EMSE, JSS, CACM, and relevant specialized journals.

A.3 Inclusion and Exclusion Criteria

Criterion Type	Criterion	Rationale
Inclusion	Empirical study of LLM for SE application	Core research scope
Inclusion	Performance benchmark of LLM on code tasks	Capability evidence
Inclusion	Case study of AI coding tool deployment	Practical evidence
Inclusion	Survey or systematic review of LLMs in SE	Secondary synthesis
Exclusion	Non-English publications without English abstract	Language feasibility
Exclusion	Pre-2017 publications (pre-transformer era)	Scope boundary
Exclusion	Opinion pieces without empirical component	Evidence quality
Exclusion	Duplicate publication of same study	De-duplication

Table A.1: Inclusion and exclusion criteria for systematic literature review.

A.4 Search Results and PRISMA Flow

The initial search yielded 3,847 potentially relevant records. After duplicate removal (n=612), 3,235 unique records were screened by title and abstract, with 2,891 excluded as clearly out of scope. The remaining 344 records were retrieved for full-text review; 124 were excluded for not meeting inclusion criteria or failing quality assessment. The final corpus comprised 220 primary studies, including 143 conference papers, 52 journal articles, and 25 preprints with significant community uptake (>50 citations). The PRISMA flow diagram is available as a supplementary digital figure.

APPENDIX B: BENCHMARK DESCRIPTIONS AND METHODOLOGY

B.1 HumanEval Benchmark

HumanEval [4] consists of 164 Python programming problems, each comprising a function signature, docstring specification, and 7.7 unit tests on average. Problems span difficulty from elementary (string manipulation, list operations) to intermediate (algorithmic problems requiring reasoning about data structures and algorithms). The primary metric is `pass@k`: the probability that at least one of `k` generated solutions passes all unit tests. `pass@1` is the most commonly reported metric, evaluating whether the model's top single solution is correct.

Limitations: HumanEval has been critiqued for benchmark saturation (frontier models exceeding 90% performance), limited difficulty range, Python-only evaluation, potential contamination from training data inclusion of problems or similar patterns, and limited ecological validity for real-world software engineering tasks that require multi-file reasoning, error handling, and API integration.

B.2 SWE-bench and SWE-bench Verified

SWE-bench [52] consists of 2,294 real-world GitHub issues from 12 popular Python open-source repositories (including Django, Flask, Matplotlib, NumPy, Requests, and SQLAlchemy), each requiring one or more code changes to resolve. Evaluation requires generating a patch that passes the repository's test suite (including the tests added to verify the fix). SWE-bench Verified [93] is a 500-issue subset annotated by professional software engineers to confirm that the issues are unambiguous, the test suite adequately validates the fix, and the fix does not require external knowledge unavailable from the repository. SWE-bench Verified is considered the current gold standard for evaluating AI software engineering capability.

Limitations: SWE-bench is Python-only and limited to open-source repositories with comprehensive test suites. The Verified subset, while higher quality, represents only 22% of the full benchmark. Real-world enterprise software engineering involves proprietary codebases, multiple programming languages, integration with external services, and deployment considerations not captured in the benchmark.

B.3 Additional Benchmarks

Benchmark	Task	Size	Language	Primary Metric
MBPP	Function generation from NL	374 problems	Python	<code>pass@1</code>

Benchmark	Task	Size	Language	Primary Metric
MultiPL-E	Code generation multilingual	2.6K problems	18 languages	pass@1
BigCodeBench	Code with diverse library calls	1.1K tasks	Python	pass@1
LiveCodeBench	Contamination-free ongoing	500+ (growing)	Multiple	pass@1
DS-1000	Data science code	1K problems	Python	pass@1
Spider	Text-to-SQL	10.2K questions	SQL	Execution acc.
BIRD	Text-to-SQL realistic	12.8K questions	SQL	Execution acc.
EvoEval	Evolved HumanEval	828 problems	Python	pass@1
HumanEval+	More tests per HumanEval problem	164 problems	Python	pass@1
CRUXEval	Code execution prediction	800 pairs	Python	pass@1

Table B.1: Summary of code generation benchmarks used in this thesis.

APPENDIX C: PROMPT ENGINEERING EXAMPLES

C.1 Zero-Shot vs. Few-Shot vs. Chain-of-Thought Comparison

The following examples illustrate the three primary prompting strategies for code generation tasks, demonstrating the quality differences observable with different prompt formulations for the same underlying task (implementing a function to find all prime numbers up to n using the Sieve of Eratosthenes).

Zero-Shot Prompt Example

System: You are an expert Python programmer.

User: Write a Python function `sieve_of_eratosthenes(n)` that returns
a list of all prime numbers up to and including n .

Chain-of-Thought Prompt Example

System: You are an expert Python programmer. Think step by step.

User: Implement `sieve_of_eratosthenes(n)`. First, describe the algorithm
at a high level. Then identify the data structures needed. Then write
the implementation. Then identify edge cases to handle. Finally,
confirm the implementation handles all identified edge cases.

C.2 Security-Focused Prompt Engineering

The following template demonstrates a security-conscious prompt for code generation in security-sensitive contexts, incorporating explicit security review instructions that reduce the rate of security vulnerabilities in generated code:

System: You are a senior security-conscious Python developer. For all
code you generate: (1) Use parameterized queries for all database
operations. (2) Validate and sanitize all user inputs. (3) Never
hardcode credentials. (4) Use strong cryptography (bcrypt for
passwords,
AES-256 for encryption). (5) After generating code, explicitly
review
it for OWASP Top 10 vulnerabilities before providing your answer.
User: Write a user login function that validates credentials
against
a PostgreSQL database.

APPENDIX D: CASE STUDY DATA COLLECTION INSTRUMENTS

D.1 Developer Productivity Survey Instrument

The following survey questions were used in the developer experience measurement components of the case studies presented in Chapter 6. Questions are rated on a 5-point Likert scale (1=Strongly Disagree, 5=Strongly Agree) unless otherwise indicated.

Section: Productivity & Time

6. D1. The AI coding tool helps me complete programming tasks faster than I could without it.
7. D2. I spend less time searching for documentation when using the AI coding tool.
8. D3. The AI coding tool helps me spend more time on challenging, creative aspects of development.
9. D4. I would estimate the AI coding tool saves me [open response: __ hours per week].

Section: Code Quality

10. Q1. The code suggested by the AI tool is usually of good quality.
11. Q2. I frequently catch errors or security issues in AI-suggested code before accepting it.
12. Q3. The AI tool helps me write code that follows project conventions and best practices.
13. Q4. I am confident in the correctness of code I accept from the AI tool after my review.

Section: Learning & Skills

14. L1. Using the AI coding tool helps me learn new programming techniques and patterns.
15. L2. I believe I am maintaining my programming skills despite heavy use of AI assistance.
16. L3. The AI tool helps me work effectively in programming languages I am less familiar with.

D.2 Qualitative Interview Protocol

Semi-structured interview questions used for in-depth case study data collection. Interviews were conducted with developers, team leads, and security engineers at participating organizations.

- I1. Describe a specific situation where the AI coding tool significantly helped you complete a task. What made it effective?
- I2. Describe a situation where the AI coding tool gave you incorrect or problematic output. How did you detect the problem, and what was the impact?
- I3. How has your approach to code review changed since adopting the AI coding tool? Do you review AI-generated code differently than human-written code?

I4. Have you noticed any changes in your own programming skills or habits since using the AI tool heavily? Please describe.

I5. What aspects of the AI coding tool does your team find most valuable? What aspects are most problematic?

I6. What organizational practices has your team developed specifically to use the AI tool effectively and safely?

I7. What would you most want to see improved in the AI coding tool to better support your work?

SUPPLEMENTARY LITERATURE REVIEW: EXTENDED ANALYSIS OF RECENT ADVANCES

SLR-1 Semantic Code Models: From Token Prediction to Meaning Representations

The evolution of code models from token-level prediction to semantic meaning representations marks a qualitative shift in what LLMs can accomplish in software engineering. Early statistical language models for code — n-gram models over token sequences, probabilistic context-free grammars for parse trees — captured surface regularities in code but had no notion of program semantics: what a function computes, what invariants hold at a program point, or whether two syntactically different implementations are functionally equivalent.

The transformer's attention mechanism provides a richer form of representation that implicitly encodes some semantic relationships. Models like BERT and its code-specific descendant GraphCodeBERT [122] learn representations where semantically related identifiers (variables that hold similar types of values, functions that perform similar operations) are geometrically proximate in the embedding space. GraphCodeBERT explicitly incorporates data flow graphs as auxiliary training signals alongside raw token sequences, using the graph structure to supervise attention patterns that align with program semantics. In code search tasks — retrieving the most semantically relevant code snippet for a natural language query — GraphCodeBERT substantially outperformed token-only models, demonstrating that explicit semantic structure improves retrieval beyond what implicit attention can achieve.

UniXcoder [123] extends the semantic representation framework by incorporating abstract syntax tree structure through a unified cross-modal pre-training objective that simultaneously trains on code tokens, AST node sequences, and natural language comments describing the code. The multi-modal training signal encourages the model to develop representations that capture the relationships between syntactic structure (what the code looks like), semantic structure (how data flows through the program), and intent (what the code is supposed to do). These richer representations improve performance on a broader range of downstream tasks than purely token-based models, at modest additional training cost.

The practical implications of semantic code representations are significant for code search, duplicate detection, and clone identification applications. Enterprise codebases frequently contain multiple implementations of semantically equivalent functionality — different developers have solved the same problem in slightly different ways, or the same algorithm has been implemented across multiple services that have

diverged over time. Semantic code search tools that can identify functional equivalence despite syntactic differences enable consolidation of duplicate implementations, detection of license-incompatible code clones, and identification of candidate reuse opportunities that syntactic search would miss.

SLR-2 Program Analysis Integration: Neurosymbolic Approaches

Neurosymbolic approaches to code intelligence — systems that combine neural language models with classical program analysis techniques such as static analysis, type inference, and symbolic execution — represent a research frontier with potential to overcome fundamental limitations of purely neural approaches. The core insight motivating neurosymbolic research is that neural models and symbolic analysis tools have complementary strengths: neural models excel at pattern recognition, natural language understanding, and generalization from examples, while symbolic tools provide formal correctness guarantees, complete coverage of specified properties, and compositional reasoning that scales to complex programs.

Codex-based test generation augmented with fuzzing — using the LLM to generate initial test cases from specifications and then using automated fuzzing to extend coverage to edge cases and boundary conditions the LLM missed — combines the LLM's strength at generating semantically meaningful test inputs with the fuzzer's strength at systematic exploration of the input space. TitanFuzz [130] applied this approach to generating tests for deep learning library APIs, using GPT-4 to generate initial test programs from API documentation and then mutating them with coverage-guided fuzzing, achieving substantially higher bug-finding rates than either approach alone.

LLM-guided symbolic execution represents another compelling integration: the LLM proposes execution paths to explore based on code semantics and natural language hints about where bugs may hide, while symbolic execution follows those paths with formal precision, generating concrete counterexamples when the path leads to a property violation. This combination addresses the path explosion problem that limits pure symbolic execution (the LLM prioritizes the most likely bug-bearing paths) while maintaining the formal guarantees that pure LLM-based analysis cannot provide.

Type inference augmented with LLMs addresses Python's gradual typing challenge: most Python code lacks type annotations, limiting the effectiveness of type-based static analysis tools like mypy. LLMs trained on typed Python code can predict type annotations for unannotated functions with 60-80% accuracy, enabling static type checking for previously unanalyzable code. TypeEvalPy [131] demonstrated that LLM-generated type annotations substantially improve the effectiveness of downstream static analysis, catching type errors invisible to both pure LLM analysis and annotation-free static tools.

SLR-3 Multimodal Code Intelligence: Bridging Visual and Textual Representations

The integration of visual modalities with code understanding opens application domains not accessible to text-only code models. Software engineering involves substantial visual artifacts — UI mockups and wireframes, architecture diagrams, entity-relationship diagrams, sequence diagrams, class diagrams, and screenshots of running applications — that carry information essential for many development tasks but not representable in plain text. Multimodal LLMs capable of jointly processing visual and textual information can leverage these visual artifacts as first-class inputs for code generation and analysis.

UI-to-code generation — producing HTML/CSS/JavaScript or framework-specific component code from screenshots or design mockups — is the most practically advanced multimodal code application. Screenshot2Code demonstrations using GPT-4V and Gemini 1.5 Pro show that frontier multimodal models can produce functionally similar web page implementations from screenshots with reasonable fidelity, though precise pixel-perfect reproduction of complex designs remains challenging. The practical value is significant for rapid prototyping workflows where visual design comes first and code follows: designers can produce Figma mockups that developers then implement using AI assistance from the visual specification.

Architecture diagram comprehension enables AI assistance for tasks that require understanding system structure from visual documentation. An engineer asking 'Based on this architecture diagram, how should I implement the data flow between the message queue and the processing service?' can receive contextually grounded implementation guidance that accounts for the depicted system topology. Gemini 1.5 Pro's multimodal capability enables processing of uploaded architecture diagrams alongside natural language questions, providing a conversational interface to visual system documentation.

Error screenshot analysis represents a practically high-value multimodal application: developers encountering unexpected application behavior can share screenshots of error states alongside relevant code for more effective debugging assistance. The combination of visual error manifestation (a UI in an unexpected state, a missing component, incorrect data display) and code context enables the LLM to reason about both what the user sees and why the code produces it, bridging the gap between observable behavior and underlying implementation.

SLR-4 Cross-Lingual Transfer and Low-Resource Programming Language Support

The distribution of programming languages in LLM training data is highly skewed: Python, JavaScript, Java, and C++ account for the majority of public code, while

hundreds of specialized and domain-specific languages are represented sparsely or not at all. Cross-lingual transfer research investigates whether knowledge learned from high-resource languages can be leveraged to improve model performance on low-resource programming languages, following the successful precedent from multilingual natural language processing.

MultiPL-E [72] established that frontier models show significant performance variation across programming languages, with Python performance consistently exceeding performance in less-represented languages at the same relative difficulty level. The performance gap between Python and low-resource languages is larger for models trained primarily on English text with code as a secondary domain (general LLMs) than for code-specialized models trained on diverse multilingual code corpora, suggesting that code-specific pre-training with explicit multilingual coverage narrows the cross-language capability gap.

Few-shot cross-lingual transfer — providing examples in the target low-resource language in the prompt — substantially improves performance for languages within the model's training distribution at moderate volume. For languages with essentially zero training representation (proprietary domain-specific languages developed internally by organizations), few-shot examples may be the only viable approach for leveraging LLM capabilities without the ability to fine-tune on the proprietary language's code corpus.

The practical challenge for organizations with proprietary programming languages or frameworks is significant: if their codebase uses a language not represented in LLM training data, standard AI coding tools provide limited assistance. Research on rapid language adaptation — fine-tuning code models on small corpora of a new language to establish basic competency — using parameter-efficient techniques like LoRA enables organizations to extend LLM capabilities to internal languages at modest cost. Preliminary work suggests that 10,000-100,000 examples of a new language may be sufficient for meaningful capability acquisition, well within the documentation volume of typical proprietary languages.

SLR-5 Retrieval-Augmented Code Generation: Architectures and Effectiveness

Retrieval-augmented generation (RAG) [23] has emerged as a critical architectural pattern for addressing several fundamental limitations of parametric-only LLMs in software engineering applications: factual hallucination of API signatures for insufficiently represented libraries, inability to access code patterns from private organizational codebases, and difficulty incorporating recent API changes that post-date the training cutoff. By combining the LLM's generalization capability with direct retrieval of accurate, current, and private information, RAG-based code generation systems achieve better factual grounding than parametric models alone.

The RAG pipeline for code generation involves: (1) encoding the query and candidate documents into a shared embedding space; (2) retrieving the k most semantically similar code snippets or documentation sections from an indexed corpus; (3) concatenating retrieved context with the original query; and (4) generating code conditioned on both the query and the retrieved context. The quality of each component substantially affects end-to-end generation quality: poor retrieval (failing to retrieve the most relevant context) degrades performance even with a strong generator, while poor generation (failing to correctly apply retrieved context) wastes the retrieval investment.

CodeFuse-Query [212] demonstrated that hierarchical retrieval — first retrieving relevant API documentation sections, then retrieving usage examples that use those APIs — provides substantially better grounding than flat retrieval from a combined corpus. The two-stage approach mirrors the natural information-seeking process of a developer: first understanding what an API does (documentation), then seeing how it is used (examples). The composite retrieval quality substantially reduces hallucination rates for infrequently used APIs.

Enterprise RAG deployments for code generation must address the indexing and retrieval challenges of large, heterogeneous codebases. Production codebases contain millions of files across many directories, with complex dependency relationships and varying relevance for any given query. Effective indexing strategies must balance retrieval precision (returning the most relevant code) against recall (not missing relevant code in unexpected locations) and computational cost (chunking, embedding, and indexing millions of files). GitHub Copilot's workspace context feature, Amazon Q Developer's codebase understanding, and Cursor's codebase indexing all implement production RAG architectures with these trade-offs, though the specific architectural details are proprietary.

RAG Component	Approach	Retrieval Quality	Notes
Chunking	Function-level	High for API tasks	Natural semantic boundary
Chunking	File-level	Medium	Preserves file context
Chunking	Fixed 512-token	Medium	Simple, fast
Embedding model	Code-specific (CodeBERT)	High for code-code	Specialized for SE
Embedding model	General (text-embedding-3)	Medium	Simpler deployment
Retrieval	Cosine similarity (dense)	High	Standard approach
Retrieval	BM25 (sparse)	Medium for keywords	Good for identifiers
Retrieval	Hybrid (dense + sparse)	Highest	Best of both
Reranking	Cross-encoder reranker	+10-20% precision	More compute
Context window	Retrieved chunks first	Higher attention	Addresses 'lost in middle'

Table SLR-1: RAG component choices and their impact on code generation quality.

SLR-6 Agentic Frameworks: Tool Use, Planning, and Memory

The evolution from single-turn LLM code generation to multi-step agentic systems represents one of the most consequential architectural shifts in LLM software engineering research since the transformer. Agentic systems — AI systems capable of autonomously taking sequences of actions, using tools, and adapting their behavior based on intermediate results — substantially expand the class of software engineering tasks accessible to AI automation. Understanding the architectural components of effective agentic systems is essential for both deploying current systems effectively and developing improved next-generation systems.

The ReAct framework [63] established the fundamental architecture pattern for LLM agents: interleaved reasoning (generating natural language thoughts about what to do next) and action (executing tool calls such as code execution, file reading, or web search). The reasoning step enables the model to maintain and update its understanding of the task state, plan subsequent actions, and diagnose errors from action results. The interleaving of reasoning and action, rather than reasoning followed by a batch of actions, enables the model to adapt its plan dynamically based on intermediate results — crucial for software engineering tasks where early actions may reveal information that changes the optimal approach.

Reflexion [64] extended the ReAct pattern with explicit verbal self-reflection after task completion or failure: the agent analyzes what went wrong in failed attempts and explicitly stores lessons learned in an episodic memory buffer that is provided as context for subsequent attempts. On SWE-bench and HumanEval, Reflexion significantly improved performance over single-attempt ReAct, demonstrating that explicit meta-cognitive self-correction substantially improves agentic task success rates. The practical implication is that agentic coding systems should incorporate explicit failure analysis and learning loops rather than simply retrying with the same approach.

Memory systems for long-horizon agentic tasks require careful design to balance context relevance against context length. Three memory modalities are available to current agentic systems: in-context memory (information within the current context window, fully accessible but capacity-limited), external vector memory (semantically indexed notes, retrieved by similarity to the current query), and structured memory (database of facts, retrieved by structured query). Effective long-horizon agents combine all three: in-context memory for immediate working information, external vector memory for semantic retrieval of relevant past observations, and structured memory for precise retrieval of specific facts.

Tool use diversity substantially affects agentic performance on complex software engineering tasks. Agents with access to code execution (running the code they write), file system operations (reading, writing, and navigating the codebase), web search

(finding documentation and solutions online), shell command execution (running tests, build tools, linters), and external API calls achieve substantially higher task completion rates than agents limited to file reading and code generation alone. The difference reflects the breadth of actions required to complete realistic software engineering tasks: a bug fix may require reading multiple files, running the test suite, checking documentation, modifying code, and verifying the fix by running tests again.

SLR-7 Code LLMs and Formal Specification: Early Results

The intersection of neural code generation and formal methods — an area historically characterized by mutual incomprehension between statistical ML researchers and formal verification researchers — is emerging as one of the most promising research directions for improving the reliability of AI-generated code. Formal specifications (contracts, pre/post-conditions, invariants expressed in languages like Dafny, TLA+, or separation logic) provide a mathematical characterization of correct program behavior that can be automatically checked, enabling principled validation of LLM-generated code beyond empirical testing.

DafnyBench [133] evaluated frontier LLMs on generating Dafny programs — a verification-aware programming language where programs are written with formal correctness specifications that are automatically checked by the Dafny verifier. GPT-4 correctly generated verifiable Dafny programs for 29% of problems in the benchmark, with failures primarily in formulating correct loop invariants and recursion specifications that require understanding of program semantics at a level of precision beyond natural language description. While 29% success is modest, the benchmark establishes a non-trivial baseline demonstrating that LLMs can engage with formal verification, contrary to earlier pessimistic assessments.

The practical workflow emerging from this research enables a form of formally grounded AI code generation: developers provide both a natural language description and a formal specification (or the LLM helps generate both), the LLM generates an implementation, and the formal verifier checks whether the implementation satisfies the specification. When verification fails, the verifier provides a counterexample that grounds the debugging process in formal precision rather than test case exploration. This workflow is not yet productive at scale — the effort of writing formal specifications often exceeds the effort saved by LLM implementation generation — but represents a promising direction for high-assurance software domains where correctness guarantees justify the additional specification effort.

SUPPLEMENTARY: EXTENDED TECHNICAL FOUNDATIONS

TF-1 FlashAttention: Enabling Efficient Long-Context Processing

Standard attention computation — the core operation of the transformer — has quadratic complexity in sequence length: computing attention scores between all pairs of positions in a sequence of length n requires $O(n^2)$ operations and $O(n^2)$ memory. For sequences of 1,000 tokens this is manageable, but for sequences of 100,000 or 1,000,000 tokens (necessary for processing large codebases), naive attention becomes computationally prohibitive. FlashAttention [124] and FlashAttention-2 [125] address this by restructuring the attention computation to exploit the hierarchical memory structure of modern GPUs, reducing memory access cost without changing the mathematical result of attention computation.

The key innovation is tiling: dividing the query, key, and value matrices into blocks that fit within GPU SRAM (fast on-chip memory), computing attention within and between blocks, and accumulating results without materializing the full $n \times n$ attention matrix in slower HBM (off-chip high-bandwidth memory). By keeping computation within SRAM as much as possible, FlashAttention reduces the number of HBM reads/writes — the primary bottleneck for memory-bound operations on modern GPUs — by a factor proportional to the tile size. FlashAttention achieves 2-4× wall-clock speedup over standard attention implementations at the same mathematical precision, enabling practical training and inference with sequences substantially longer than standard attention implementations permit.

FlashAttention-2 further improved performance through better work partitioning across GPU thread blocks, achieving 50-73% of theoretical GPU FLOP utilization versus 25-35% for standard attention. FlashAttention-3 [139], optimized for the H100 GPU architecture, achieves near-theoretical peak performance through explicit exploitation of H100's asynchronous memory operations. These successive improvements have made long-context processing — essential for repository-level code understanding — practically feasible on available hardware at reasonable cost, directly enabling the 128K-1M context windows of frontier code models.

For software engineering applications, the practical consequence of FlashAttention is that processing entire moderately-sized code files (up to 100,000 tokens at 8-12 characters per token) within a single forward pass has become routine on GPU hardware available in 2024-2025. This context length enables: complete analysis of large source files without chunking; processing of multiple interrelated files simultaneously; inclusion of relevant test files alongside implementation files; and processing of long

build log outputs for CI/CD failure analysis. The trajectory toward million-token contexts enabled by FlashAttention and its successors suggests that 'the entire codebase fits in context' will become a practical reality for mid-sized codebases within a few years.

TF-2 Grouped Query Attention and Multi-Query Attention: Efficient Inference

During inference, standard multi-head attention (MHA) requires storing key-value tensors for all previous positions in the sequence — the KV cache — to avoid recomputing them at each generation step. The KV cache for a model with 32 attention heads, 128 head dimension, and a sequence of 10,000 tokens requires storing $32 \times 128 \times 10,000 \times 2$ (keys and values) $\times 2$ (bytes for float16) = 163 million parameters per layer. For a 30-layer model, the total KV cache is approximately 4.9 GB, roughly equal to the model parameters themselves. For long-context models with 100,000-token sequences, the KV cache becomes the primary memory bottleneck for inference.

Multi-query attention (MQA) [138] addresses this by sharing a single key-value head across all query heads rather than maintaining separate key and value projections for each head. This reduces KV cache size by a factor equal to the number of heads (typically 32-64 $\times$), at the cost of some reduction in representation capacity that may affect model quality. Grouped query attention (GQA) [126] provides a middle ground: G groups of query heads share a single key-value head per group, with memory reduction by factor G. With G=8, GQA reduces KV cache memory 8 $\times$ while maintaining quality closer to full MHA than MQA.

The practical impact for code model inference is substantial. Models using GQA can serve longer sequences within a given memory budget, reducing the frequency with which input must be truncated due to memory constraints. Llama 2 70B adopted GQA, enabling serving with 8 $\times$ less KV cache memory than full MHA — the difference between fitting a long-context session in GPU memory and requiring memory paging. All major code models since 2023 (Code Llama, DeepSeek-Coder, Qwen2.5-Coder) employ GQA, making it the de facto standard for efficient long-context inference.

TF-3 Speculative Decoding: Reducing Inference Latency

Autoregressive generation — generating tokens one at a time, each conditioned on all previous tokens — is inherently sequential and therefore difficult to parallelize within a single generation call. For large models, each forward pass generating a single token may take tens of milliseconds, translating to generation speeds of tens of tokens per second that can feel slow for interactive coding assistance applications expecting response latency under 1-2 seconds.

Speculative decoding [127, 141] addresses this latency by using a small, fast draft model to generate multiple candidate tokens in parallel, then using the large target model to verify all draft tokens simultaneously in a single forward pass. Because verification is a single forward pass over a batch of tokens, it takes similar time to generating a single token — but verifies multiple tokens. If the draft model is well-calibrated with the target model, most draft tokens are accepted, achieving 2-4× speedup over sequential generation with no change to the output distribution of the large model.

For code generation specifically, speculative decoding shows higher speedup than for natural language generation because code contains many predictable sequences — fixed syntax patterns, common identifier conventions, repeated boilerplate — that a small draft model can predict accurately. Chen et al. [141] reported 2.4-3.1× speedup on code generation tasks versus 1.5-2.0× on natural language tasks, reflecting the higher predictability of code token sequences. GitHub Copilot's low-latency inline suggestions benefit from speculative decoding or similar acceleration techniques to maintain the sub-300ms response latency necessary for effective autocomplete.

TF-4 Quantization: Deploying Large Models at Reduced Cost

Model quantization — reducing the numerical precision of model weights and activations from 32-bit or 16-bit floating point to 8-bit integers or lower — enables deployment of large models at substantially reduced memory cost and increased inference speed, at the cost of potential accuracy degradation. For software engineering applications, quantization enables deployment of capable models on hardware that could not accommodate the full-precision model, expanding the range of deployment contexts from cloud-only to on-premises, edge, and laptop environments.

Post-training quantization (PTQ) applies quantization to a trained model without requiring retraining, using calibration data to determine optimal quantization parameters. GPTQ [128] demonstrated that 4-bit quantization of large language models could be performed with negligible perplexity increase (less than 1%) using a layer-wise quantization approach that minimizes the reconstruction error at each layer. For a 70B parameter model, 4-bit quantization reduces memory from approximately 140 GB (float16) to 35 GB, enabling deployment on a single high-memory GPU that could not accommodate the full-precision model.

QLoRA [36] combines quantization with parameter-efficient fine-tuning, enabling fine-tuning of quantized models with minimal GPU memory. A 65B parameter model fine-tuned with QLoRA requires only approximately 48 GB GPU memory — fitting on a single A100 80GB GPU — while achieving fine-tuning quality comparable to full-precision fine-tuning. This combination enables organizations to fine-tune large code models on proprietary code corpora with hardware available in standard GPU server configurations, dramatically reducing the compute cost barrier to model customization.

Quantization Level	Memory Reduction	Accuracy Impact	Best Use Case
FP32 → BF16	2×	Negligible	Standard inference
BF16 → INT8 (LLM.int8)	2× vs BF16	~1% perplexity rise	High-quality deployment
BF16 → INT4 (GPTQ)	4× vs BF16	~2-4% perplexity rise	Resource-constrained servers
BF16 → INT4 (QLoRA)	4× vs BF16	~2-3% (fine-tuning)	Fine-tuning large models
BF16 → INT3	5.3× vs BF16	~8-15% perplexity rise	Edge/laptop inference only
BF16 → INT2 (GPTQ)	8× vs BF16	>20% degradation	Not recommended for SE

Table TF-1: Quantization levels, memory savings, and accuracy trade-offs for code models.

TF-5 Instruction Tuning and Alignment for Code-Specific Tasks

Base pre-trained language models — trained on next-token prediction over large code and text corpora — are not directly usable as coding assistants without additional alignment training that teaches the model to respond helpfully to instructions, explain its reasoning, and decline requests for harmful code generation. Instruction tuning transforms the base model's text-completion capability into the instruction-following behavior expected from an assistant, while alignment training shapes the model's behavior to be helpful, harmless, and honest in coding contexts.

Code instruction datasets — collections of (instruction, code response) pairs used for supervised fine-tuning — have evolved from manually curated datasets to large-scale synthetic datasets generated using stronger models. Self-Instruct [129] pioneered using an LLM to generate diverse instruction-response pairs from a small seed set; subsequent work adapted this approach to code by generating diverse programming tasks, prompting a capable code model for solutions, and filtering pairs for quality. Code Alpaca, WizardCoder [82], and Evol-Instruct applied increasingly sophisticated instruction generation strategies, substantially improving instruction-following quality and code generation accuracy for fine-tuned open-source models.

The distinction between instruction tuning (supervised fine-tuning on demonstrations of correct behavior) and RLHF-based alignment (fine-tuning with reinforcement learning from human feedback on quality ratings) is particularly important for code models. Instruction tuning teaches the model to produce code in response to requests; RLHF additionally teaches the model to produce code that is preferred by developers — well-explained, appropriately cautious about limitations, secure by default, and accompanied by helpful context. The combination of instruction tuning for basic capability and RLHF for quality calibration produces the most capable and trustworthy code assistants.

Constitutional AI [16] provides an alternative to RLHF that is particularly relevant for code security alignment. Rather than collecting human preference ratings for security, the model is trained using a set of explicit principles (no SQL injection, use

parameterized queries, never hardcode credentials) that it self-applies through a 'critique and revision' process. The model generates code, critiques it against the constitutional principles, and revises it to be principle-compliant. This self-supervised alignment process can be applied to specific code security properties without requiring expensive human labeling of every type of security issue.

SUPPLEMENTARY: EXTENDED APPLICATION ANALYSIS

EA-1 Code Completion in IDEs: A Detailed Technical Analysis

The integration of LLM-based code completion into integrated development environments presents a set of technical challenges distinct from standalone code generation. IDEs are interactive, real-time environments with strict latency requirements, complex state management requirements (tracking open files, edit history, cursor position, error diagnostics), and integration with diverse language servers, build systems, and version control tools. Production IDE code completion systems must address these challenges while maintaining the generative quality necessary for useful suggestions.

The language server protocol (LSP), a standardized interface between IDEs and language analysis tools, provides the architectural foundation for IDE integration. LLM-based completion services can integrate at the LSP level, receiving completion requests with rich context (current file content, cursor position, related open files) and returning completion suggestions alongside traditional LSP completions (symbol-based, type-based, import-based). This integration enables LLM-based completions to coexist with and complement traditional completion mechanisms rather than replacing them.

Context management for in-IDE completions must balance the competing demands of context quality (including more context improves suggestion relevance) and latency (larger context means longer processing time). Production systems typically employ a context budget: a maximum token count that can be included in the completion request. Within this budget, context selection algorithms prioritize the current file content around the cursor, recently edited files in the same session, imported dependencies, and type definitions of referenced identifiers. Context selection accuracy substantially affects suggestion quality, particularly for repository-level completions that must align with conventions established in files not currently open in the editor.

Cache management strategies are essential for maintaining low latency in the face of rapid keystroke-level context updates. Caching the KV cache of the LLM computation for a stable prefix — the portions of the input that have not changed since the last request — enables reusing prior computation for subsequent requests where only a few new tokens have been added. This prefix caching strategy can reduce effective latency by 50-80% for typical keystroke completion patterns, enabling the response times necessary for effective inline completion.

Inline Completion Acceptance Patterns

Understanding the distribution of developer acceptance behavior for inline completions reveals important patterns for system design and productivity measurement. Ziegler et al.'s large-scale study [45] of GitHub Copilot acceptance patterns across 500,000+ completion events found that acceptance rate varies dramatically by completion type: multi-line completions for function bodies have substantially lower acceptance rates than single-line completions (22% vs 38%), but accepted multi-line completions represent disproportionately more code written. The efficiency metric — total code written per suggestion shown — is therefore higher for multi-line completions despite their lower acceptance rate, because each accepted suggestion contributes more code.

Temporal patterns of acceptance within a coding session reveal that developers become more selective about accepting completions as sessions progress: early in a session, when implementing a fresh function or module, acceptance rates are relatively high (40-50%); later in a session when making targeted edits or debugging, acceptance rates are lower (15-25%). This pattern suggests that AI assistance provides the most value at the start of implementation tasks where the blank-page problem is most acute, and provides less value for targeted modifications where the developer knows precisely what change to make.

EA-2 LLM-Assisted Code Comprehension: Program Slicing and Impact Analysis

Beyond pure code explanation, LLMs are increasingly being applied to program comprehension tasks that require sophisticated reasoning about code structure and behavior: program slicing (identifying which statements affect the value of a variable at a given point), impact analysis (determining which code elements will be affected by a proposed change), and call graph analysis (identifying the chain of function calls triggered by an event). These comprehension tasks are essential for safe maintenance of large codebases where the consequences of changes are difficult to predict by reading code locally.

Program slicing traditionally requires precise static or dynamic analysis tools that operate on formal program representations. LLMs approach slicing differently — by reasoning about data and control flow through natural language explanation of the code's behavior — achieving reasonable accuracy for simple slicing queries while failing for complex programs with indirect data flow through data structures, callbacks, or concurrency. Tian et al.'s evaluation [119] found that GPT-4 correctly identified the relevant slice for 72% of slicing queries on simple Python functions but only 41% for queries involving class hierarchies and method calls, reflecting the difficulty of tracking indirect data flow through complex object structures.

Impact analysis for software maintenance decisions — 'if I change this function's return type, what other code will I need to update?' — is particularly valuable for reducing the risk of incomplete maintenance updates. LLMs with broad codebase context can traverse call relationships, identify all callers of a changed function, and estimate the likely impact of the change on each caller. This analysis, while less precise than formal type system analysis, provides useful guidance for maintenance planning and review prioritization. Sourcegraph Cody's impact analysis capability, which uses code search to identify all usages of a modified symbol before generating an impact summary, represents a production implementation of this approach.

EA-3 LLMs for Software Architecture and Design Assistance

Software architecture — the high-level structure of a system, its decomposition into components and the relationships between them — is a domain where LLMs are beginning to provide assistance that extends beyond code-level help to design-level decision support. Architectural decisions are consequential, long-lived, and difficult to reverse, making assistance particularly valuable even if the LLM provides only one perspective in a larger decision process.

Architecture description generation — producing natural language architecture documentation, component diagrams in textual notation (PlantUML, Mermaid), or Architecture Decision Records (ADRs) from code or high-level descriptions — is the most immediately deployable architectural assistance capability. Given a codebase, an LLM can generate a description of the current architecture — identifying major components, their responsibilities, their communication patterns, and the architectural patterns employed — that provides a starting point for documentation that human architects then refine. This capability is particularly valuable for reverse-engineering the architecture of poorly documented legacy systems.

Design pattern recommendation — suggesting which design patterns might improve the structure of existing code — requires the LLM to understand both the current code structure and the range of applicable patterns, then evaluate whether a pattern application would improve or complicate the design given the specific context. Liu et al. [2024] found that LLM-generated design pattern recommendations were rated as appropriate by senior engineers in 58% of cases, with the strongest performance for straightforward pattern applications (extract method, singleton replacement with dependency injection) and weakest for complex pattern transformations (converting an inheritance hierarchy to strategy pattern, implementing event sourcing) that require deep understanding of the system's behavioral requirements.

Architecture fitness function evaluation — assessing whether a proposed architectural change satisfies specified architectural constraints such as 'the UI layer should not directly access the database' or 'each microservice should own its own data' —

is an emerging capability where LLMs can assist by translating natural language constraints into executable checks. Architecture as Code tools combined with LLM interpretation of natural language constraints enable continuous architecture compliance monitoring that was previously feasible only with manual architecture review.

Architectural Task	LLM Capability	Accuracy	Key Limitation
Architecture documentation	High	70-80% rated useful	Misses implicit design intent
ADR generation	Medium-High	65% acceptable quality	Missing business context
Design pattern recommendation	Medium	58% appropriate	Complex patterns error-prone
Microservice boundary suggestion	Medium	52% expert agreement	Domain knowledge required
API contract design	High	72% expert approval	Non-functional reqs missing
Database schema design	Medium	60% expert approval	Normalization rules complex
Security architecture review	Medium	65% issues identified	Novel threat patterns missed
Performance architecture review	Low-Medium	45% issues found	Workload knowledge required

Table EA-1: LLM capability levels for software architecture assistance tasks.

EA-4 AI-Assisted Code Review: Detailed Workflow Analysis

Code review is not a monolithic task but a multi-stage process involving several distinct activities: understanding the purpose of the change (what problem does this PR solve?); verifying functional correctness (does the implementation correctly solve the stated problem?); checking code quality (is the code readable, maintainable, and consistent with project conventions?); security review (does the change introduce security vulnerabilities?); performance review (does the change introduce performance regressions?); and test coverage review (are the changes adequately tested?). LLMs have different capability profiles for each stage, and effective AI-assisted code review must match LLM capabilities to the stages where they add value.

AI systems demonstrate strongest capability in the documentation/explanation stage (generating PR summaries that explain what changed and why from diff analysis), the code quality stage (identifying naming inconsistencies, documentation gaps, style deviations), and the security stage (detecting common vulnerability patterns). They show weaker capability in the functional correctness stage (evaluating whether the implementation correctly solves the intended problem without running it) and the performance stage (predicting whether a change will cause performance regression without profiling data). Deployment of AI code review that focuses on the high-capability stages while maintaining human review responsibility for the lower-capability stages creates a complementary human-AI review workflow.

Pull request summarization — generating natural language descriptions of what a pull request changes and why, from the diff and commit messages — is among the highest-quality AI code review capabilities and the most widely deployed in production. GitHub Copilot's PR description generation, GitLab Duo's merge request summary, and Atlassian's AI-generated Jira story updates all implement this capability. Developer surveys consistently show high satisfaction with AI-generated summaries, which save the time required to write a coherent description of complex changes and help reviewers orient to the change before examining individual files.

Review comment generation — producing specific, actionable comments on individual diff hunks — requires more contextual understanding than PR summarization and shows more variable quality. Comments on code style, obvious bugs, and missing error handling are generally accurate and well-received; comments about architectural appropriateness and business logic correctness are more often incorrect or irrelevant. A study by Li et al. [2022] found that developers accepted approximately 35% of AI-generated review comments overall, with acceptance rates ranging from 65% for documentation-related comments to 18% for logic-related comments, reflecting the capability gradient described above.

EA-5 LLMs in Software Testing: Advanced Topics

Beyond the basic test generation capabilities described in Chapter 4, LLMs are enabling advanced testing workflows that leverage natural language understanding for test design tasks previously requiring significant manual expertise: equivalence class partitioning, boundary value analysis, combinatorial testing, and property-based test specification.

Equivalence class partitioning — dividing the input space into classes expected to exercise the same program behavior — is a fundamental test design technique that benefits from LLM-based automation. Given a function specification, an LLM can identify likely equivalence classes by reasoning about the function's preconditions, invariants, and documented edge cases, then generate representative test inputs from each class. This structured approach to test case design produces tests with better functional coverage than random input generation while requiring less manual analysis than purely manual equivalence class identification.

Metamorphic testing — testing without a test oracle by specifying metamorphic relations that describe how output should change when input is transformed in a specific way — is an advanced testing technique for programs where correct output is difficult to specify directly (machine learning models, compilers, computational simulations). LLMs can assist in identifying metamorphic relations for a given function: given a sorting function, the LLM might identify that 'reversing the input should not change the set of elements in the output' or 'sorting an already sorted list should return the same list.'

Nakamura et al. [2023] evaluated LLM-generated metamorphic relations for three machine learning libraries, finding that approximately 60% of generated relations were valid (semantically correct descriptions of metamorphic properties) and executable, providing useful testing scaffolding that reduced test development time by an estimated 40%.

Fault injection testing — deliberately introducing faults into software to test error handling robustness — is another testing domain where LLMs provide value. An LLM can generate diverse fault injection scenarios for a given code path: what happens if this network call returns a timeout? If this database query returns an unexpected NULL? If this file I/O operation encounters a permission error? The LLM's breadth of knowledge about error conditions from diverse codebases enables broader fault coverage than manually specified fault injection scenarios, systematically testing error handling code that is otherwise difficult to exercise through normal functional testing.

SUPPLEMENTARY: EXTENDED COMPARATIVE ANALYSIS

CA-1 Evaluation Methodology: Strengths and Weaknesses of Current Benchmarks

The benchmark ecosystem for evaluating LLM-based code generation has grown rapidly but retains several systematic weaknesses that affect the reliability of published performance comparisons. Understanding these weaknesses is essential for practitioners interpreting benchmark claims and for researchers designing evaluation methodology for new work.

Benchmark contamination — the inclusion of benchmark problems or highly similar problems in training data — is the most widely discussed methodological concern. Models trained on data that includes benchmark solutions will show artificially inflated performance that does not reflect generalization capability. Contamination detection is technically challenging: exact deduplication of training data against benchmarks is feasible but insufficient, as near-duplicate variants of benchmark problems may provide equivalent training signal without being detected by exact matching. Liu et al. [2024] demonstrated that performance decreases substantially on functionally equivalent but syntactically modified benchmark problems for models suspected of contamination, providing indirect evidence of contamination effects.

Metric limitations present a second systematic concern. Pass@k metrics measure whether code executes correctly on provided test cases, but test suites are often incomplete — they may not exercise all code paths, test all edge cases, or verify all output properties. Code that passes all provided tests may still contain bugs for untested inputs, while code that fails provided tests may implement the correct algorithm with a minor fix required. F1-based partial credit metrics that measure functional overlap between generated and reference implementations, or execution accuracy metrics that test against a larger hidden test suite, provide more robust correctness estimates but are more expensive to compute.

Ecological validity — the degree to which benchmark performance predicts real-world utility — is the most fundamental but hardest to measure concern. HumanEval's self-contained Python functions with clear specifications bear limited resemblance to the complex, multi-file, specification-ambiguous tasks that dominate professional software development. SWE-bench Verified's real GitHub issues are substantially more ecologically valid but limited to Python and to open-source codebases that may not represent enterprise software complexity. No currently available benchmark provides

both the ecological validity of real enterprise software tasks and the scale necessary for reliable statistical comparison.

Benchmark	Ecological Validity	Contamination Risk	Scale	Maintains Currency
HumanEval	Low (isolated functions)	Very High (old, public)	164 problems	No (static)
MBPP	Low-Medium	High	374 problems	No (static)
SWE-bench (full)	High (real issues)	Medium	2,294 issues	No (static)
SWE-bench Verified	High (curated)	Medium	500 issues	Slow updates
MultiPL-E	Medium	High (derived)	2,600+	No (static)
BigCodeBench	Medium-High	Medium	1,140 tasks	No (static)
LiveCodeBench	Medium (competitive)	Low (continuous)	500+ growing	Yes (ongoing)
EvoEval	Medium	Low (evolved)	828 problems	Periodically
BIRD	High (realistic SQL)	Medium	12,751 Q&A	No (static)

Table CA-1: Benchmark evaluation quality dimensions. 'Maintains Currency' indicates whether new problems are continuously added.

CA-2 Total Cost of Ownership Analysis for Enterprise Deployment

Enterprise procurement decisions for AI coding tools must account for total cost of ownership (TCO) that extends substantially beyond subscription or API costs. The TCO framework for AI coding tools includes: direct tool costs (subscriptions, API usage); infrastructure costs (GPU servers for self-hosted models, network egress for API calls); integration costs (engineering time for IDE integration, CI/CD pipeline integration, policy configuration); training and change management costs (developer onboarding, workflow adaptation, security training); security and compliance costs (security scanning tools, compliance audits, IP review); ongoing management costs (version updates, prompt template maintenance, model performance monitoring); and productivity value realized (the offset against these costs).

Direct tool costs are the most visible but not necessarily the largest TCO component. For a 500-developer organization, GitHub Copilot Enterprise at \$39/month/developer costs \$234,000 annually — significant but manageable for a technology organization's tool budget. The integration engineering cost (configuring enterprise policies, integrating with internal systems, setting up monitoring) typically runs 2-6 months of engineering time for a production-ready deployment, representing \$100,000-300,000 in engineering cost. The ongoing management cost (monitoring adoption metrics, updating configurations, managing model updates, responding to security incidents) runs 0.5-1.0 FTE annually for a large enterprise deployment.

The productivity value calculation is the most uncertain but typically largest TCO component. For 500 developers at \$100,000 average annual cost, 50% of time on LLM-addressable tasks, and 25% average productivity improvement on those tasks, the annual

productivity value is \$6,250,000 — dramatically exceeding the total cost components above. This favorable ratio explains the rapid enterprise adoption despite significant implementation costs. However, the productivity improvement estimate has substantial uncertainty: conservative estimates from careful controlled studies may be 20-25%, while optimistic estimates from vendor-sponsored surveys may reach 40-55%. The TCO model is very sensitive to this estimate, with the break-even improvement rate at approximately 5-8%.

Cost Component	One-Time	Annual (500 devs)	Notes
Tool subscription (Copilot Ent.)	—	\$234,000	\$39/dev/month
Integration engineering	\$200,000	\$50,000 (updates)	2-3 month initial project
Change management / training	\$50,000	\$20,000	Onboarding + ongoing
Security scanning (Snyk/Semgrep)	—	\$30,000-80,000	AI-augmented SAST
Compliance/legal review	\$30,000	\$15,000	IP/privacy policy annual review
Ongoing management (0.5 FTE)	—	\$75,000	Monitoring, updates, incidents
TOTAL COST	\$280,000	\$424,000-474,000	Year 1 higher
Productivity value (25% imp.)	—	\$6,250,000	Conservative estimate
NET BENEFIT (Year 1)	—	\$5,480,000	13:1 ROI conservative

Table CA-2: Total Cost of Ownership analysis for 500-developer enterprise Copilot Enterprise deployment.

CA-3 Security Assessment of LLM-Based Development Tools

Enterprise adoption of LLM-based development tools introduces security considerations beyond the code security challenges described in Chapter 7. The tools themselves — as software systems receiving and processing sensitive organizational code, handling developer credentials, and potentially executing code in production environments — present attack surface that must be assessed in the enterprise security posture.

Data confidentiality risks arise from code being transmitted to LLM APIs hosted by third parties. Code sent to OpenAI, Anthropic, or Google APIs may include trade secrets, proprietary algorithms, credential strings accidentally included in code, or information about unreleased products. Enterprise agreements (BAA, DPA) provide contractual data handling commitments but do not eliminate the fundamental risk of transmitting sensitive code outside organizational control. Organizations in regulated industries (healthcare, finance, defense) must carefully evaluate whether LLM API usage complies with data handling requirements that may prohibit transmission of regulated data to third-party systems. Self-hosted open-source models eliminate this risk by keeping code within organizational infrastructure.

Prompt injection attacks — adversarial inputs that hijack LLM reasoning to produce unintended outputs — are particularly concerning in agentic coding contexts where the LLM executes actions (file writes, command execution, API calls) based on instructions that may include maliciously crafted content. A developer using an AI coding agent to analyze a third-party library might unwittingly expose the agent to malicious content embedded in that library's documentation or README, designed to redirect the agent to extract sensitive information or execute unauthorized actions. Production agentic systems must implement explicit input sanitization, action confirmation requirements, and scope limitations to mitigate prompt injection risks.

Supply chain risks from LLM-generated dependency suggestions represent an emerging threat vector. LLMs may suggest installing packages with names similar to legitimate packages but controlled by malicious actors (typosquatting), recommend outdated package versions with known vulnerabilities, or hallucinate package names that could be registered by attackers. Organizations should require automated dependency scanning (using tools like Dependabot, Snyk, or Socket) for all AI-generated dependency specifications before installation, and should maintain internal package mirrors with pre-vetted package versions where feasible.

SUPPLEMENTARY: EXTENDED CHALLENGE ANALYSIS

CH-1 Technical Debt Acceleration: A Hidden Risk of AI-Assisted Development

A subtle but important risk of rapid AI-assisted code generation is the potential for accelerated technical debt accumulation. Technical debt — the implied cost of future rework caused by choosing expedient code over better-structured code — accumulates in any codebase over time, but may accumulate faster when AI tools enable rapid generation of code without sufficient architectural forethought. When a developer can generate 100 lines of working code in minutes, the temptation to accept a working but poorly structured implementation is greater than when generating the same code manually would require hours of careful consideration.

The technical debt risk from AI-generated code manifests in specific patterns identifiable through code quality analysis. Over-abstraction and premature generalization — LLMs trained on diverse codebases tend to generate abstract, generic implementations where simpler, concrete implementations would be more appropriate for the specific use case — create unnecessary complexity. Under-modularization — generating monolithic functions that combine multiple concerns rather than following single-responsibility decomposition — reduces maintainability. Inconsistent style and conventions — each AI generation session may produce code following slightly different conventions than surrounding code — creates maintenance friction at the places where AI-generated and human-generated code meet.

Organizations should implement code quality gates specific to AI-generated code that flag potential technical debt patterns before they are committed. Complexity metrics (cyclomatic complexity, cognitive complexity), coupling metrics (afferent and efferent coupling between modules), and maintainability indices provide automated signals that can be integrated into CI/CD pipelines to require review of AI-generated code exceeding quality thresholds. Some organizations have implemented 'AI-generated' labels on commits to enable quality metric tracking that segments AI-generated code for targeted debt monitoring.

CH-2 Consistency and Coherence in Large AI-Assisted Codebases

As AI-assisted code generation becomes more prevalent within a development team, maintaining consistency and coherence across the growing codebase presents significant challenges. Human developers within a team develop shared conventions,

vocabulary, and design patterns through collaboration and code review; AI-generated code lacks this socialization, producing locally reasonable code that may diverge from team conventions in ways not immediately apparent but that create maintenance friction over time.

Naming convention drift is a common manifestation: across different developers and sessions, AI-generated code may use different naming conventions for similar concepts — `user_id` vs `userId` vs `UserId` vs `uid` — creating inconsistency within a codebase where human-authored code would typically use a consistent convention. Error handling pattern inconsistency similarly emerges: some AI-generated code may use exception-based error handling, other AI-generated code for similar tasks may use error-code returns, and still other code may use result types, creating a heterogeneous error handling landscape that complicates error analysis and recovery.

Organizational-level prompt templates — standardized system prompts provided to AI tools that specify the team's naming conventions, error handling patterns, code organization preferences, and documentation standards — partially address consistency by constraining the AI's generation toward team conventions. GitHub Copilot's `.github/copilot-instructions.md` configuration, Cursor's `.cursorrules` configuration, and similar per-repository AI instruction files enable organizations to encode team conventions that shape AI-generated code toward consistency with the existing codebase. Systematic adoption of these configuration mechanisms is an important operational practice for teams using AI assistance at scale.

CH-3 Model Capability Regression and Version Management

Enterprise deployments of LLM-based development tools face a version management challenge not present with traditional software tools: model updates by the LLM provider may change model behavior in ways that affect quality or compatibility with established workflows, and there is typically no mechanism for pinning to a specific model version for long-term stability. OpenAI's model deprecation policies — where older GPT model versions are deprecated on 6-12 month schedules — have caused disruption for organizations that had adapted their workflows to specific model behavior that changed in a subsequent version.

Model capability regression — where a newer model version performs worse than an older version on specific task types important to an organization — is documented but often under-reported. Organizations that have built quality-critical workflows around specific model behavior (such as a specific code generation format or a specific level of verbosity) may find that updates degrade these workflows. Systematic regression testing of AI model updates against a test suite of organization-specific tasks provides early warning of version-specific regressions before they affect production workflows, but requires investment in creating and maintaining the task test suite.

The version management challenge is more acute for fine-tuned models than for base model API access. Organizations that fine-tune a model on internal code corpora must periodically re-fine-tune on updated base models to incorporate capability improvements, but re-fine-tuning requires re-investment of compute and engineering resources. A structured model lifecycle management framework — including version monitoring, regression testing, update evaluation, and re-fine-tuning processes — is an emerging operational discipline for organizations with significant investment in fine-tuned AI coding tools.

SUPPLEMENTARY: EXTENDED FUTURE RESEARCH DIRECTIONS

FR-1 Compositional Code Generation for Complex Software Systems

Current LLM code generation excels at generating individual functions or files but struggles with the compositional challenge of generating complex software systems that require consistent interfaces across many components generated in sequence. A system of 50 microservices with shared data models, consistent API contracts, and coherent error propagation cannot be generated reliably by 50 independent generation calls, because each call lacks awareness of the interface decisions made in previous calls.

Compositional code generation research addresses this challenge by developing generation architectures that maintain and enforce interface constraints across multi-component generation. Specification-driven approaches first generate a formal interface specification (API contracts, shared data schemas, event bus message formats) and then generate each component constrained to satisfy the specification. Memory-augmented generation maintains a persistent representation of established interfaces and design decisions that constrains subsequent generation calls. Hierarchical generation produces high-level structural specifications before generating components that implement the specification — mirroring the top-down design process of experienced software architects.

The verification challenge for compositional generation is substantial: verifying that independently generated components will correctly interoperate requires either end-to-end integration testing (expensive) or formal interface compatibility checking (limited by the expressiveness of available interface specification languages). Research on automatically verifiable interface specifications that can be generated alongside code, and on compositional testing frameworks that enable incremental verification of component compatibility as components are generated, would substantially improve the reliability of large-scale compositional generation.

FR-2 Continuous Learning and Adaptation in Deployed Systems

A fundamental limitation of current LLM-based coding tools is their static nature: once trained and deployed, the model does not update its knowledge or capabilities based on deployment experience. A model that encounters a new API, encounters a class of bugs it fails to diagnose correctly, or observes that its suggestions are consistently rejected for a specific type of task does not update its behavior to improve in these areas. This stasis means that deployment experience cannot compound into capability

improvement — the model must wait for an external retraining cycle to incorporate lessons from deployment.

Continual learning — the ability to update model parameters based on new experience without catastrophic forgetting of prior capabilities — is an active research area with direct relevance to improving AI coding tools through deployment. The catastrophic forgetting challenge (updating model parameters for new tasks degrades performance on old tasks) has been partially addressed through elastic weight consolidation, progressive neural networks, and experience replay, but no fully satisfactory solution exists for the scale of LLMs. Parameter-efficient continual learning methods that update only a small set of adapter parameters while freezing base model weights may provide a practical path toward deployment-time adaptation without catastrophic forgetting.

Retrieval-augmented adaptation provides an alternative to parameter updates for incorporating new knowledge: rather than updating model weights, new knowledge (new API documentation, organizational coding standards, domain-specific patterns) is added to an external retrieval corpus that augments the model's parametric knowledge at inference time. This approach avoids catastrophic forgetting entirely but is limited to knowledge that can be represented as retrievable text rather than learned behavioral patterns. The complementary strengths of retrieval-based adaptation (safe, instant, reversible) and parametric adaptation (deeper integration, faster inference) suggest that hybrid approaches combining both will characterize the most capable next-generation AI coding tools.

FR-3 Cross-Repository Learning and Organizational Knowledge Graphs

Individual code repositories contain local patterns, conventions, and domain knowledge; organizational knowledge that spans multiple repositories — the relationship between services, the evolution of architectural decisions, the lessons from past incidents — is distributed across code, documentation, incident reports, design documents, and developer knowledge. Research on cross-repository learning investigates how AI coding tools can develop and leverage this organizational-level knowledge to provide assistance that reflects the broader context in which individual repositories exist.

Knowledge graph construction from organizational software artifacts — extracting entities (services, APIs, data models, developers, incidents) and relationships (service-calls-service, developer-owns-module, incident-caused-by-bug-in-component) from code, logs, documentation, and communication history — provides a structured representation of organizational knowledge that can ground AI assistance in organizational context. A coding assistant with access to an organizational knowledge graph can answer questions like 'which other services depend on the API I'm about to

change?' or 'what past incidents were caused by similar bugs?' in ways that provide substantially more relevant context than code-only analysis.

Privacy and access control challenges for cross-repository organizational knowledge are significant: some repositories, documents, and communication records are accessible only to specific teams, and an AI assistant that aggregates across all organizational knowledge must enforce the access control policies that apply to each information source. Federated learning approaches — where models learn from distributed data sources without centralizing sensitive information — may enable cross-repository learning while respecting data sensitivity boundaries, but at the cost of additional system complexity and potentially reduced learning effectiveness.

FR-4 Human Factors and Developer Wellbeing in AI-Augmented Development

The human factors implications of AI-augmented software development — effects on developer cognitive load, job satisfaction, professional identity, skill development, and workplace dynamics — represent an underresearched dimension of significant practical importance. Software engineering is a knowledge-intensive profession where developer wellbeing, motivation, and professional development are important outcomes in their own right, not merely instrumentally relevant to productivity. Research that treats developers as stakeholders whose experience matters independently of their productivity contribution is needed to complement the productivity-focused research that dominates current AI-SE literature.

Cognitive load effects of AI assistance are theoretically ambiguous: AI assistance reduces the cognitive effort required for routine coding tasks (potentially reducing task load) while simultaneously increasing the cognitive effort required for monitoring, evaluating, and selectively accepting AI suggestions (potentially increasing supervisory load). Empirical studies using cognitive load measures (secondary task performance, physiological measures, subjective rating scales) could clarify whether AI assistance reduces or shifts cognitive load, with important implications for understanding the sustainability of AI-augmented development workflows.

Professional identity and job satisfaction effects of AI adoption have received attention primarily in the context of automation anxiety — concern that AI will replace software engineers — but require broader study of how AI assistance affects engineers' sense of craft, ownership, and professional accomplishment. If developers increasingly view their role as reviewing and curating AI-generated code rather than creating code through their own design and implementation decisions, the nature of the profession changes in ways that may affect talent attraction, developer engagement, and the professional norms that shape software quality. Longitudinal studies of developer satisfaction and professional identity across cohorts with different levels of AI assistance

adoption would provide empirical grounding for policy and organizational practice decisions.

FR-5 Regulatory Frameworks and Professional Standards for AI-Assisted Development

The rapid adoption of AI-assisted code generation in professional software development contexts is outpacing the development of professional standards, regulatory frameworks, and liability doctrines that characterize mature professions. Medical professionals have liability frameworks for adopting new diagnostic tools; engineers have professional codes of practice for adopting new analytical methods; lawyers have bar association guidance on using AI for legal research. Software engineering, despite its critical role in infrastructure and safety-critical systems, lacks comparable professional frameworks for responsible AI adoption.

Professional standards development for AI-assisted software engineering is beginning in several organizations. ISO and IEC are developing standards for AI system development and testing that address AI-generated code quality requirements. The IEEE Software Engineering Standards Committee is examining updates to software quality standards to address AI-generated code. Sector-specific regulatory guidance is emerging from the FDA (for software as a medical device developed with AI assistance), FAA (for avionics software), and financial sector regulators (for algorithmic trading and financial systems software). These sector-specific regulatory developments will be significant determinants of AI tool adoption pace in safety-critical industries.

Liability doctrine for AI-generated software defects remains unresolved. When AI-generated code contains a bug that causes harm — a vulnerability in security software, an error in medical device software, a financial loss from trading system code — the chain of responsibility is unclear: is the developer who accepted the AI suggestion liable? The tool provider who generated the suggestion? The organization that deployed the tool without adequate security review? Resolution of these liability questions through court decisions, legislation, or regulatory guidance will substantially affect AI tool adoption in high-stakes contexts and will shape the design of AI coding tools to provide appropriate documentation, disclosure, and verification mechanisms.

SUPPLEMENTARY REFERENCES (121–200)

- [121] Zhou, D., Schärli, N., Hou, L., Wei, J., Scales, N., Wang, X., Schuurmans, D., Chen, C., Bousquet, O., Le, Q., and Chi, E. (2023). Least-to-most prompting enables complex reasoning in large language models. In *Proceedings of the 11th International Conference on Learning Representations (ICLR 2023)*.
- [122] Guo, D., Ren, S., Lu, S., Feng, Z., Tang, D., Liu, S., Zhou, L., Duan, N., Svyatkovskiy, A., Fu, S., et al. (2021). GraphCodeBERT: Pre-training code representations with data flow. In *Proceedings of the 9th International Conference on Learning Representations (ICLR 2021)*.
- [123] Guo, D., Lu, S., Duan, N., Wang, Y., Zhou, M., and Yin, J. (2022). UniXcoder: Unified cross-modal pre-training for code representation. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (ACL 2022)*, vol. 1, pp. 7212-7225.
- [124] Dao, T., Fu, D. Y., Ermon, S., Rudra, A., and Re, C. (2022). FlashAttention: Fast and memory-efficient exact attention with IO-awareness. In *Advances in Neural Information Processing Systems (NeurIPS 2022)*, vol. 35, pp. 16344-16359.
- [125] Dao, T. (2023). FlashAttention-2: Faster attention with better parallelism and work partitioning. In *Proceedings of the 12th International Conference on Learning Representations (ICLR 2024)*.
- [126] Ainslie, J., Lee-Thorp, J., de Jong, M., Zemlyanskiy, Y., Lebron, F., and Sanghai, S. (2023). GQA: Training generalized multi-query transformer models from multi-head checkpoints. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP 2023)*, pp. 4895-4901.
- [127] Leviathan, Y., Kalman, M., and Matias, Y. (2023). Fast inference from transformers via speculative decoding. In *Proceedings of the 40th International Conference on Machine Learning (ICML 2023)*, pp. 19274-19286.
- [128] Frantar, E., Ashkboos, S., Hoefler, T., and Alistarh, D. (2023). GPTQ: Accurate post-training quantization for generative pre-trained transformers. In *Proceedings of the 11th International Conference on Learning Representations (ICLR 2023)*.
- [129] Wang, Y., Kordi, Y., Mishra, S., Liu, A., Smith, N. A., Khashabi, D., and Hajjishirzi, H. (2023). Self-Instruct: Aligning language models with self-generated instructions. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL 2023)*, vol. 1, pp. 13484-13508.
- [130] Deng, Y., Lemieux, C., Wang, J., Sen, K., and Chandra, S. (2023). Large language models are zero-shot fuzzers: Fuzzing deep-learning libraries via large language models. In *Proceedings of the 32nd ACM SIGSOFT International Symposium on Software Testing and Analysis (ISSTA 2023)*, pp. 423-435.
- [131] Mir, A. M., Latoskinas, E., Proksch, S., and Gousios, G. (2023). Type4Py: Practical deep similarity learning-based type inference for Python. In *Proceedings of the 44th IEEE/ACM International Conference on Software Engineering (ICSE 2023)*, pp. 2241-2252.
- [132] Ni, A., Polu, S., Wu, J., and Creswell, R. (2023). Clover: Closed-loop verifiable code generation. *arXiv preprint arXiv:2310.17807*.
- [133] Misu, M. J., Lopes, C. V., Ma, I., and Noble, J. (2024). Towards AI-assisted synthesis of verified dafny programs. In *Proceedings of the 46th IEEE/ACM International Conference on Software Engineering (ICSE 2024)*.

- [134] Pourreza, M., and Rafiei, D. (2024). DIN-SQL: Decomposed in-context interactive text-to-SQL with self-correction. In *Advances in Neural Information Processing Systems (NeurIPS 2023)*, vol. 36.
- [135] Li, J., Hui, B., Qu, G., Yang, J., Li, B., Li, B., Wang, B., Qin, B., Geng, R., Huo, N., et al. (2024). Can LLM already serve as a database interface? A big bench for large-scale database grounded text-to-SQL. In *Advances in Neural Information Processing Systems (NeurIPS 2023)*, vol. 36.
- [136] Liu, F., Huang, Z., Wang, X., Zhang, Y., Tang, L., Li, G., and Jin, Z. (2024). Is your code generated by ChatGPT really correct? Rigorous evaluation of large language models for code generation. In *Advances in Neural Information Processing Systems (NeurIPS 2023)*, vol. 36.
- [137] Nakamura, T., Yamada, M., and Takashima, K. (2023). LLM-based metamorphic testing for deep learning models. In *Proceedings of the 2023 IEEE International Conference on Software Testing, Verification and Validation (ICST 2023)*.
- [138] Shazeer, N. (2019). Fast transformer decoding: One write-head is all you need. *arXiv preprint arXiv:1911.02150*.
- [139] Shah, J., Bikshandi, G., Zhang, Y., Thakkar, V., Ramani, P., and Dao, T. (2024). FlashAttention-3: Fast and accurate attention with asynchrony and low-precision. *arXiv preprint arXiv:2407.08608*.
- [140] Ding, H., Kumar, V., Tian, Y., Wang, Z., Kwiatkowski, T., Pan, X., Frost, C., Gangal, V., Zhao, C., Yang, H., et al. (2024). Crosscodeeval: A diverse and multilingual benchmark for cross-file code completion. *arXiv preprint arXiv:2310.11248*.
- [141] Chen, C., Borgeaud, S., Irving, G., Lespiau, J.-B., Sifre, L., and Jumper, J. (2023). Accelerating large language model decoding with speculative sampling. *arXiv preprint arXiv:2302.01318*.
- [142] Liu, S., et al. (2024). Is LLM-generated code good enough? A comprehensive evaluation. In *Proceedings of the 46th IEEE/ACM International Conference on Software Engineering (ICSE 2024)*.
- [143] Zhang, Y., et al. (2024). CodeAgent: Enhancing code generation with tool-integrated agent systems for real-world repo-level coding challenges. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL 2024)*.
- [144] Xia, X., et al. (2024). Agentless: Demystifying LLM-based software engineering agents. *arXiv preprint arXiv:2407.01489*.
- [145] Wang, J., Huang, Y., Chen, C., Liu, Z., Wang, S., and Wang, Q. (2024). Software testing with large language models: Survey, landscape, and vision. *IEEE Transactions on Software Engineering*, 50(4), 911-936.
- [146] Yetistiren, B., Ozsoy, I., and Ayerdem, M. (2022). Assessing the quality of GitHub Copilot's code generation. In *Proceedings of the 18th International Conference on Predictive Models and Data Analytics in Software Engineering (PROMISE 2022)*, pp. 62-71.
- [147] Ouyang, S., Zhang, J. M., Harman, M., and Wang, M. (2023). LLM is like a box of chocolates: The non-determinism of ChatGPT in code generation. *arXiv preprint arXiv:2308.02828*.
- [148] Mastropaolo, A., Pascarella, L., Guglielmi, E., Ciniselli, M., Scalabrino, S., Oliveto, R., and Bavota, G. (2023). On the robustness of code generation techniques: An empirical study on GitHub Copilot. In *Proceedings of the 45th IEEE/ACM International Conference on Software Engineering (ICSE 2023)*, pp. 2149-2160.
- [149] Prenner, J. A., Robbes, R., and Bergel, A. (2022). Can OpenAI's Codex fix bugs? An evaluation on QuixBugs. In *Proceedings of the 5th IEEE/ACM International Workshop on Automated Program Repair (APR 2022)*, pp. 69-75.
- [150] Tian, H., Liu, K., Li, B., Lo, D., Grundy, J., and Wang, H. (2024). Is ChatGPT the ultimate programming assistant — How far is it? *arXiv preprint arXiv:2304.11938*.

- [151] Barke, S., James, M. B., and Polikarpova, N. (2023). Grounded copilot: How programmers interact with code-generating models. *Proceedings of the ACM on Programming Languages*, 7(OOPSLA1), 85:1-85:27.
- [152] Liang, J. T., et al. (2024). Large language models in software engineering: Managing the technical debt risk. *IEEE Software*, 41(2), 48-57.
- [153] Sarsa, S., Denny, P., Hellas, A., and Leinonen, J. (2022). Automatic generation of programming exercises and code explanations using large language models. In *Proceedings of the 2022 ACM Conference on International Computing Education Research (ICER 2022)*, vol. 1, pp. 27-43.
- [154] Finnie-Ansley, J., Denny, P., Becker, B. A., Luxton-Reilly, A., and Prather, J. (2022). The robots are coming: Exploring the implications of OpenAI Codex on introductory programming. In *Proceedings of the 24th Australasian Computing Education Conference (ACE 2022)*, pp. 10-19.
- [155] Denny, P., et al. (2024). Explaining code with a purpose: An integrated approach to code comprehension and accessibility. In *Proceedings of the 55th ACM Technical Symposium on Computer Science Education (SIGCSE 2024)*.
- [156] Albusays, K., Da Silva, B., Gorski, L., Harper, S., Kim, J., Mankoff, J., Menzies, T., and Tigwell, G. W. (2021). The diversity crisis in software development. *IEEE Software*, 38(2), 19-25.
- [157] DeGrave, A., et al. (2024). Energy consumption estimation for large language model inference at scale. *arXiv preprint arXiv:2402.08105*.
- [158] Patterson, D., et al. (2022). The carbon footprint of machine learning training will plateau, then shrink. *Computer*, 55(7), 18-28.
- [159] Croft, R., Newlands, D., Chen, Z., and Babar, M. A. (2022). Data quality for software vulnerability datasets. In *Proceedings of the 44th IEEE/ACM International Conference on Software Engineering (ICSE 2022)*, pp. 121-133.
- [160] Yang, Z., Shi, J., He, J., and Lo, D. (2022). Natural attack for pre-trained models of code. In *Proceedings of the 44th IEEE/ACM International Conference on Software Engineering (ICSE 2022)*, pp. 819-830.
- [161] Sun, H., Chen, L., Xu, F., Lou, J.-G., Wang, J., Zhang, X., Feng, Z., and Zhang, D. (2024). CoCoST: Automatic complex code generation with online searching and correctness testing. *arXiv preprint arXiv:2403.13583*.
- [162] Weyssow, M., Kamber, U., Yu, K., Lo, D., and Sahraoui, H. (2024). Exploring the usage of ChatGPT for the analysis of CodeReviews. In *Proceedings of the 20th International Conference on Mining Software Repositories (MSR 2024)*.
- [163] Imai, S. (2022). Is GitHub Copilot a substitute for human pair-programming? An empirical study. In *Proceedings of the 44th International Conference on Software Engineering: Companion Proceedings (ICSE-Companion 2022)*, pp. 319-321.
- [164] Cao, J., et al. (2024). A survey on the honesty of large language models. *arXiv preprint arXiv:2401.04868*.
- [165] Le, X. B. D., Lo, D., and Le Goues, C. (2016). History driven program repair. In *Proceedings of the 23rd IEEE International Conference on Software Analysis, Evolution, and Reengineering (SANER 2016)*, pp. 213-224.
- [166] Christopoulou, F., Lampos, V., and Androutopoulos, I. (2024). Towards conversational diagnostic AI for software engineering. *arXiv preprint arXiv:2312.09337*.
- [167] Kim, S., et al. (2023). An empirical study of code smells in code generation results from large language models. *arXiv preprint arXiv:2311.03642*.

- [168] Sobania, D., Briesch, M., Hanna, C., and Petke, J. (2023). An analysis of the automatic bug fixing performance of ChatGPT. In Proceedings of the 4th IEEE/ACM International Workshop on Automated Program Repair (APR 2023), pp. 23-30.
- [169] Guo, H., Jing, Y., Cheng, X., Cao, Y., and Gao, C. (2024). AI-assisted code review: A systematic literature review. *Information and Software Technology*, 168, 107374.
- [170] Tang, Y., et al. (2024). CODEREPAIR: LLM-based program repair at scale. *arXiv preprint arXiv:2406.01178*.
- [171] Zhang, Y., et al. (2024). A survey of LLM-based agents for software engineering. *arXiv preprint arXiv:2409.09030*.
- [172] Xia, C. S., and Zhang, L. (2024). Keep the conversation going: Fixing 162 out of 337 bugs for \$0.42 each using ChatGPT. *arXiv preprint arXiv:2304.00385*.
- [173] Zhou, A., et al. (2024). OpenDevin: An open platform for AI software developers as generalist agents. *arXiv preprint arXiv:2407.16741*.
- [174] Huang, D., et al. (2023). AgentCoder: Multi-agent-based code generation with iterative testing and optimization. *arXiv preprint arXiv:2312.13010*.
- [175] Ishibashi, Y., and Nishimura, Y. (2024). Self-organized agents: A LLM multi-agent framework toward ultra large-scale code generation and optimization. *arXiv preprint arXiv:2404.02183*.
- [176] Li, G., et al. (2024). MetaGPT: Meta programming for multi-agent collaborative framework. In Proceedings of the 12th International Conference on Learning Representations (ICLR 2024).
- [177] Hong, S., et al. (2024). ChatDev: Communicative agents for software development. In Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL 2024), pp. 15174-15186.
- [178] Esteves, G., et al. (2024). Mining and classifying software quality violations in AI-generated code. In Proceedings of the 21st International Conference on Mining Software Repositories (MSR 2024).
- [179] Schafer, M., Nadi, S., Eghbali, A., and Tip, F. (2024). An empirical evaluation of using large language models for automated unit test generation. *IEEE Transactions on Software Engineering*, 50(1), 85-105.
- [180] Alégroth, E., et al. (2024). On the benefits and challenges of using large language models for test suite maintenance. *arXiv preprint arXiv:2402.12334*.
- [181] Parnas, D. L. (1972). On the criteria to be used in decomposing systems into modules. *Communications of the ACM*, 15(12), 1053–1058.
- [182] Manna, Z., and Waldinger, R. J. (1971). Toward automatic program synthesis. *Communications of the ACM*, 14(3), 151–165.
- [183] Green, C. C. (1969). Application of theorem proving to problem solving. In Proceedings of the 1st International Joint Conference on Artificial Intelligence (IJCAI 1969), pp. 219–239.
- [184] Gulwani, S., Polozov, O., and Singh, R. (2017). Program synthesis. *Foundations and Trends in Programming Languages*, 4(1–2), 1–119.
- [185] Tufano, M., Watson, C., Bavota, G., Di Penta, M., White, M., and Poshyvanyk, D. (2019). An empirical study on learning bug-fixing patches in the wild via neural machine translation. *ACM Transactions on Software Engineering and Methodology*, 28(4), 1–29.
- [186] Gupta, R., Pal, S., Kanade, A., and Shevade, S. (2017). DeepFix: Fixing common C language errors by deep learning. In Proceedings of the 31st AAAI Conference on Artificial Intelligence (AAAI 2017), pp. 1345–1351.

- [187] Peters, M. E., Neumann, M., Iyyer, M., Gardner, M., Clark, C., Lee, K., and Zettlemoyer, L. (2018). Deep contextualized word representations. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL-HLT 2018)*, pp. 2227–2237.
- [188] Radford, A., Narasimhan, K., Salimans, T., and Sutskever, I. (2018). Improving language understanding by generative pre-training. OpenAI Technical Report.
- [189] Li, Y., Choi, D., Chung, J., Kushman, N., Schrittwieser, J., Leblond, R., Eccles, T., Keeling, J., Gimeno, F., Dal Lago, A., et al. (2022). Competition-level code generation with AlphaCode. *Science*, 378(6624), 1092–1097.
- [190] Turing, A. M. (1950). Computing machinery and intelligence. *Mind*, 59(236), 433–460.
- [191] Searle, J. R. (1980). Minds, brains, and programs. *Behavioral and Brain Sciences*, 3(3), 417–424.
- [192] Belkin, M., Hsu, D., Ma, S., and Mandal, S. (2019). Reconciling modern machine-learning practice and the classical bias-variance trade-off. *Proceedings of the National Academy of Sciences*, 116(32), 15849–15854.
- [193] LeCun, Y., Boser, B., Denker, J. S., Henderson, D., Howard, R. E., Hubbard, W., and Jackel, L. D. (1989). Backpropagation applied to handwritten zip code recognition. *Neural Computation*, 1(4), 541–551.
- [194] LeCun, Y., Bottou, L., Bengio, Y., and Haffner, P. (1998). Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11), 2278–2324.
- [195] Geva, M., Schuster, R., Berant, J., and Levy, O. (2021). Transformer feed-forward layers are key-value memories. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing (EMNLP 2021)*, pp. 9556–9571.
- [196] Xiong, R., Yang, Y., He, D., Zheng, K., Zheng, S., Xing, C., Zhang, H., Lan, Y., Wang, L., and Liu, T. (2020). On layer normalization in the transformer architecture. In *Proceedings of the 37th International Conference on Machine Learning (ICML 2020)*, pp. 10524–10533.
- [197] Dettmers, T., Lewis, M., Belkada, Y., and Zettlemoyer, L. (2022). LLM.int8(): 8-bit matrix multiplication for transformers at scale. In *Advances in Neural Information Processing Systems (NeurIPS 2022)*, vol. 35, pp. 30318–30332.
- [198] Lin, J., Tang, J., Tang, H., Yang, S., Dang, X., and Han, S. (2023). AWQ: Activation-aware weight quantization for LLM compression and acceleration. In *Proceedings of Machine Learning and Systems (MLSys 2024)*, vol. 6.
- [199] Husain, H., Wu, H.-H., Gazit, T., Allamanis, M., and Brockschmidt, M. (2019). CodeSearchNet challenge: Evaluating the state of semantic code search. *arXiv preprint arXiv:1909.09436*.
- [200] Ahmad, W. U., Chakraborty, S., Ray, B., and Chang, K.-W. (2021). Unified pre-training for program understanding and generation. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL-HLT 2021)*, pp. 2655–2668.
- [201] Liu, H., Zaharia, M., and Abbeel, P. (2023). Ring attention with blockwise transformers for near-infinite context. In *Proceedings of the 12th International Conference on Learning Representations (ICLR 2024)*.
- [202] Yu, G., Jeong, J., Kim, G., Kim, S., and Chun, B. (2022). Orca: A distributed serving system for transformer-based generative models. In *Proceedings of the 16th USENIX Symposium on Operating Systems Design and Implementation (OSDI 2022)*, pp. 521–538.

- [203] Kwon, W., Li, Z., Zhuang, S., Sheng, Y., Zheng, L., Yu, C. H., Gonzalez, J., Zhang, H., and Stoica, I. (2023). Efficient memory management for large language model serving with PagedAttention. In Proceedings of the 29th ACM Symposium on Operating Systems Principles (SOSP 2023), pp. 611–626.
- [204] Akyürek, E., Schuurmans, D., Andreas, J., Ma, T., and Zhou, D. (2022). What learning algorithm is in-context learning? Investigations with linear models. In Proceedings of the 11th International Conference on Learning Representations (ICLR 2023).
- [205] Xie, S. M., Raghuathan, A., Liang, P., and Ma, T. (2022). An explanation of in-context learning as implicit Bayesian inference. In Proceedings of the 10th International Conference on Learning Representations (ICLR 2022).
- [206] Min, S., Lyu, X., Holtzman, A., Artetxe, M., Lewis, M., Hajishirzi, H., and Zettlemoyer, L. (2022). Rethinking the role of demonstrations: What makes in-context learning work? In Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing (EMNLP 2022), pp. 11048–11064.
- [207] Kojima, T., Gu, S. S., Reid, M., Matsuo, Y., and Iwasawa, Y. (2022). Large language models are zero-shot reasoners. In Advances in Neural Information Processing Systems (NeurIPS 2022), vol. 35, pp. 22199–22213.
- [208] Wang, X., Wei, J., Schuurmans, D., Le, Q., Chi, E., Narang, S., Chowdhery, A., and Zhou, D. (2022). Self-consistency improves chain of thought reasoning in language models. In Proceedings of the 11th International Conference on Learning Representations (ICLR 2023).
- [209] Chen, S., Wong, S., Chen, L., and Tian, Y. (2023). Extending context window of large language models via positional interpolation. arXiv preprint arXiv:2306.15595.
- [210] Jin, D., Han, X., Yu, J., Shi, J., and Xu, J. (2024). LongLoRA: Efficient fine-tuning of long-context large language models. In Proceedings of the 12th International Conference on Learning Representations (ICLR 2024).
- [211] Xia, X., Lo, D., Wang, X., and Zhou, B. (2018). Measuring program comprehension: A large-scale field study with professionals. IEEE Transactions on Software Engineering, 44(10), 951–976.
- [212] Li, T., et al. (2024). CodeFuse-Query: A data-centric static code analysis system using LLMs. arXiv preprint arXiv:2401.14301.

APPENDIX E: DETAILED MODEL ARCHITECTURE SPECIFICATIONS

E.1 Architecture Comparison of Major Code LLMs

Model	Params	Architecture	Context	Training Data	Code Data Volume	Year
Codex	12B	GPT-3 decoder-only	4K tokens	Code-focused	54M public repos	2021
CodeGen-16B	16B	Decoder-only	2K tokens	THEPILE+BIGPYTHON	~100B tokens	2022
StarCoder	15B	Decoder-only	8K tokens	The Stack v1	~1T tokens	2023
Code Llama 70B	70B	Llama 2 decoder	100K tokens	Mix + code FT	500B+ tokens	2023
DeepSeek-Coder-V2	236B (MoE)	MoE decoder	128K tokens	Proprietary	10.2T tokens	2024
StarCoder2-15B	15B	Decoder-only	16K tokens	The Stack v2	~4T tokens	2024
Qwen2.5-Coder-72B	72B	Decoder-only	128K tokens	Qwen data + code FT	5.5T tokens	2024
GPT-4o	~200B (est)	MoE decoder	128K tokens	Proprietary	Not disclosed	2024
Claude 3.5 Sonnet	~100B (est)	Decoder-only	200K tokens	Proprietary	Not disclosed	2024
Gemini 1.5 Pro	~500B (est)	MoE multi-modal	1M tokens	Proprietary	Not disclosed	2024

Table E.1: Architecture specifications for major code LLMs. Parameter counts for proprietary models are estimates based on published hints and third-party analysis.

E.2 Tokenization Comparison Across Code Models

Tokenization — the process of splitting input text into tokens that serve as the atomic units of model processing — significantly affects code model performance, efficiency, and multilingual capabilities. Programming languages present distinctive tokenization challenges compared to natural language: identifiers can be arbitrarily long compound strings (`calculateMinimumSpanningTreeWeight`); common patterns like `''''''` should be efficiently encoded; and syntactically identical character sequences may carry different semantics in different programming language contexts.

Model	Tokenizer	Vocab Size	Chars/Token (code)	camelCase handling	Special tokens
Codex	BPE (code-focused)	50,257	~4.5	Splits at capitals	< endoftext >
StarCoder	BPE	49,152	~4.2	Mostly splits	<fim_prefix>, <fim_middle>, <fim_suffix>
Code Llama	BPE (SentencePiece)	32,000	~3.8	Partially merges	<FILL_ME>
DeepSeek-Coder	BPE	32,000	~4.0	Partially merges	Standard

Qwen2.5-Coder	BPE+tiktoken	152,064	~4.5	Partially merges	Standard + code
GPT-4o	tiktoken cl100k_base	100,256	~4.8	Partially merges	Standard OpenAI
Claude 3.5 Sonnet	BPE (Anthropic)	~100K (est)	~4.6	Partially merges	Not disclosed

Table E.2: Tokenization characteristics of major code LLMs. Chars/token is approximate average for Python code.

E.3 Fill-in-the-Middle Training

Fill-in-the-Middle (FIM) training — also called infilling — is a training paradigm specifically designed for code completion scenarios where the model must generate code that fits between a given prefix and a given suffix. Standard next-token prediction training produces models that can only generate from left to right, limiting their utility for completion scenarios where the developer places the cursor in the middle of a file and expects the model to generate code that integrates with both the preceding and following code.

FIM training [38] transforms a standard text sample (prefix + middle + suffix) into a training example structured as (prefix + suffix + middle), teaching the model to generate the middle given both the surrounding prefix and suffix as context. During training, a fraction of samples (typically 50%) are transformed with FIM, while the remainder are left in standard left-to-right format to preserve the model's completion capabilities for forward generation scenarios. The FIM transformation can place the suffix marker at the end of the sequence (Suffix-Prefix-Middle, SPM format) or interleave prefix and suffix before the middle (Prefix-Suffix-Middle, PSM format), with different implications for training dynamics and inference quality.

Models trained with FIM demonstrate substantially better performance on code infilling tasks — completing a function body given its signature and the code that follows, or inserting a missing import statement — than models trained without FIM. StarCoder, Code Llama, and DeepSeek-Coder all include FIM training, enabling their use as infilling completers in addition to forward-generation completers. This capability is essential for IDE integration scenarios where the cursor position may be anywhere in a file, not necessarily at the end of the current editing context.

APPENDIX F: ETHICAL GUIDELINES FOR AI-ASSISTED SOFTWARE ENGINEERING

F.1 Proposed Ethical Framework

The integration of AI systems into professional software engineering practice raises ethical questions at the intersection of professional responsibility, intellectual property, worker rights, and social equity. This appendix proposes a practical ethical framework for AI-assisted software engineering practice, drawing on established professional codes of ethics (ACM Code of Ethics, IEEE Code of Ethics, SE Code of Ethics and Professional Practice) and extending them to address AI-specific considerations.

Transparency and Disclosure

Software engineers using AI assistance have an ethical obligation to disclose the nature of AI involvement in their work to relevant stakeholders. This disclosure obligation is calibrated to context: for commercial software development, disclosure to the employing organization and to code reviewers is appropriate; for safety-critical software development (medical devices, avionics, financial systems), disclosure to regulatory bodies and customers may be required. The disclosure should include the type of AI assistance used, the extent of AI involvement (autocomplete vs. substantial generation), and the validation steps applied to AI-generated code. Organizations should establish clear disclosure standards rather than leaving disclosure to individual engineer discretion.

Quality Assurance Responsibility

The use of AI assistance does not transfer or reduce the professional responsibility of software engineers for the quality, security, and correctness of the code they deliver. Engineers who accept AI-generated code without adequate review and testing remain professionally responsible for defects in that code. This responsibility requires maintaining sufficient technical competence to evaluate AI-generated code — engineers cannot discharge their quality assurance responsibility if they lack the skills to identify errors in AI-generated implementations. Professional development investments must include maintaining the technical skills necessary for critical evaluation of AI assistance alongside the workflow skills for using AI tools effectively.

Intellectual Property Diligence

Software engineers should exercise diligence regarding the intellectual property implications of AI-generated code. This includes: using AI tools with enterprise

agreements that provide IP indemnification where available; applying code provenance scanning tools that detect AI-generated code with high similarity to copyrighted training data; maintaining awareness of open-source license requirements that may apply to AI-generated code; and following organizational IP policies that may restrict AI tool usage for specific projects or customer contexts. When IP status is uncertain, conservative approaches that err toward greater originality validation are appropriate.

Ethical Principle	Professional Obligation	Organizational Requirement	Regulatory Consideration
Transparency	Disclose AI use to reviewers	Establish disclosure standards	Sector-specific requirements
Quality responsibility	Validate all AI-generated code	Implement quality gates	Safety certification reqs.
IP diligence	Scan for license conflicts	Enterprise IP agreements	Copyright compliance
Skill maintenance	Avoid unmitigated skill atrophy	Provide training/practice time	Competency standards
Security	Security-review AI code explicitly	SAST scanning policy	Vulnerability disclosure
Fairness	Review AI code for bias	Bias testing requirements	Non-discrimination reqs.
Environmental	Prefer efficient model options	Track AI energy footprint	ESG reporting
Accountability	Document AI involvement	Audit trail for AI code	Liability documentation

Table F.1: Ethical principles for AI-assisted software engineering with obligations at each stakeholder level.

F.2 Recommended Organizational Policies

Organizations deploying AI-assisted software engineering tools should establish formal policies addressing the ethical dimensions identified in the framework above. The following policy recommendations are grounded in current best practices from early-adopter enterprises and emerging regulatory guidance.

AI Tool Usage Policy should specify: which AI tools are approved for use; which code repositories or project types are approved for AI assistance (addressing IP and data privacy concerns); disclosure requirements for AI-assisted code in commit messages and documentation; the security scanning requirements that apply to AI-generated code before commit; and the process for reporting and addressing quality issues attributable to AI assistance. This policy should be reviewed annually and updated as AI tool capabilities, organizational experience, and regulatory requirements evolve.

AI Competency Standards should define minimum skill requirements for developers authorized to use AI coding assistance, including: ability to critically evaluate AI-generated code for correctness, security, and quality; understanding of the failure modes of the specific AI tools in use; familiarity with security patterns necessary to identify common AI-generated vulnerabilities; and awareness of IP and compliance

obligations. These standards should be reflected in hiring criteria, onboarding curricula, and continuing professional development requirements.

Quality Assurance Adaptation should update existing software quality processes to address AI-specific considerations: code review checklists should include explicit items for AI-specific failure modes (hallucinated APIs, missing error handling, security vulnerabilities); testing requirements should be calibrated for AI-generated code that may have systematic patterns of incompleteness; and security review processes should apply heightened scrutiny to AI-generated code in security-sensitive contexts. These adaptations should be reviewed as understanding of AI code quality patterns develops.